


CS 33

Machine Programming (4)

Some of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Digression (Again): Where Stuff Is (Roughly)



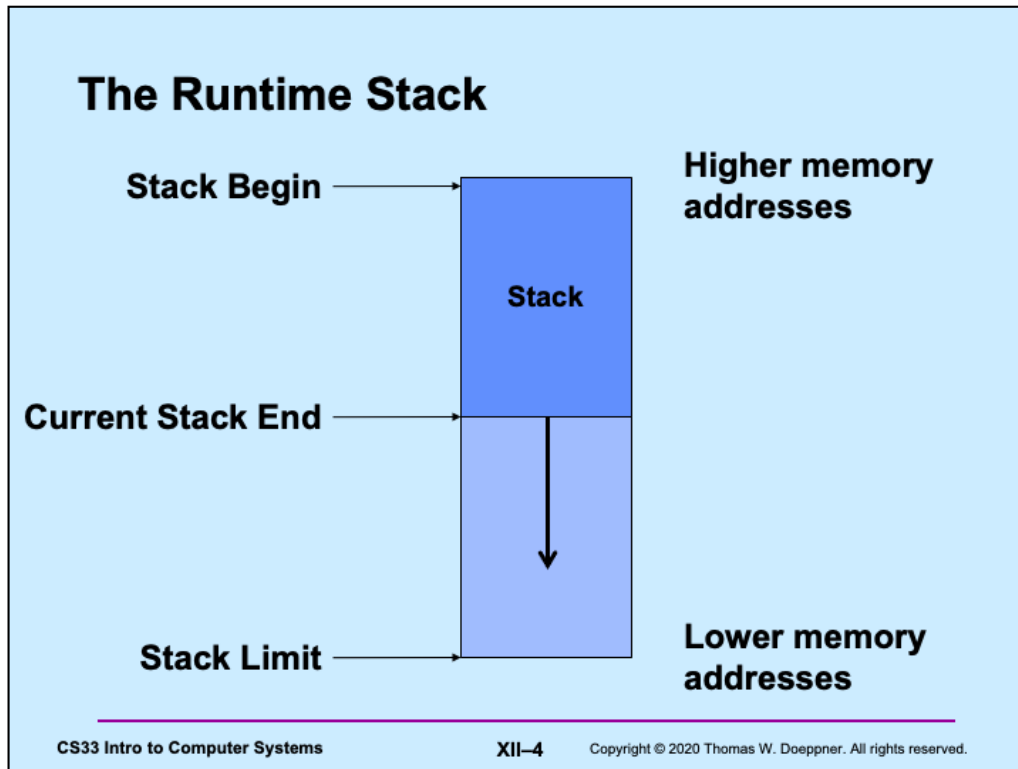
Here we revisit the slide we saw a few weeks ago, this time drawing it with high addresses at the top and low addresses at the bottom. The point is that a large amount of virtual memory is reserved for the stack. In most cases there's plenty of room for the stack and we don't have to worry about exceeding its bounds. However, if we do exceed its bounds (by accessing memory outside of what's been allocated), the program will get a seg fault.

Function Call and Return

- **Function A calls function B**
- **Function B calls function C**

... several million instructions later

- **C returns**
 - how does it know to return to B?
- **B returns**
 - how does it know to return to A?



Stacks, as implemented on the X86 for most operating systems (and, in particular, Linux, OSX, and Windows) grow "downwards", from high memory addresses to low memory addresses. To avoid confusion, we will not use the words "top of stack" or "bottom of stack" but will instead use "stack begin" and "current stack end". The total amount of memory available for the stack is that between the beginning of the stack and the "stack limit". When the stack end reaches the stack limit, we're out of memory for the stack.


Stack Operations



The stack-pointer register (`%rsp`) points to the last byte of the stack. Thus, with little-endian addressing, it points to the least-significant byte of the data item at the end of the stack. Thus, `%rsp` in the slide points to what's perhaps an 8-byte item at the end of the stack.


Push

```
pushl $0x1234
```



Here we call `pushl` to push a 4-byte item onto the end of the stack. First `%rsp` is decremented by 4 bytes, then the item is copied into the 4-byte location now pointed to by `%rsp`.

Pop



Here we pop an item off the stack. The `popl` instruction copies the 4-byte item pointed to by `%rsp` into its argument, then increments `%rsp` by 4.

Call and Return

```
0x1000: call func
0x1004: addq $3, %rax
```

```
0x2000: func:
    ... ..
0x2200: movq $6, %rax
0x2203: ret
```

When a function is called (using the *call* instruction), the (8-byte) address of the instruction just after the *call* (the "return address") is pushed onto the stack. Then when the called function returns (via the *ret* instruction), the 8-byte address at the end of the stack (pointed to by `%rsp`) is copied into the instruction pointer (`%rip`), thus causing control to resume at the instruction following the original call.

Call and Return

0x2000: func:

... ..

0x2200: movq \$6, %rax

0x2203: ret

→ 0x1000: call func
0x1004: addq \$3, %rax

stack growth
↓

0xffff10018

0xffff10010

0xffff10008

0xffff10000 ←

								%rax
00	00	00	00	00	00	10	00	%rip
00	00	00	0f	ff	f1	00	00	%rsp

Here we begin walking through what happens during a call and return.

Initially, %rip (the instruction pointer, shown with a red arrow pointing to the right) points to the call instruction – thus it's the next instruction to be executed. %rsp (the stack pointer, shown with a green arrow pointing to the left) points to the current end of the stack. The actual values contained in the relevant registers are shown at the bottom of the slide (%rax isn't relevant yet, but will be soon!).

Call and Return

```
0x1000: call func
0x1004: addq $3, %rax
```

```
→ 0x2000: func:
    ...
0x2200: movq $6, %rax
0x2203: ret
```

stack growth ↓

00	00	00	00	00	00	10	04

0xffff10018

0xffff10010

0xffff10008

0xffff10000

0xffff0fff8 ←

00	00	00	00	00	00	20	00
00	00	00	0f	ff	f0	ff	f8

%rax

%rip

%rsp

When the *call* instruction is executed, the address of the instruction after the *call* is pushed onto the stack. Thus %rsp is decremented by eight and 0x1004 is copied to the 8-byte location that is now at the end of the stack. The instruction pointer, %rip, now points to the first instruction of *func*.

Call and Return

```
0x1000: call func
0x1004: addq $3, %rax
```

```
0x2000: func:
```

```
... ..
```

```
0x2200: movq $6, %rax
```

```
0x2203: ret
```

stack growth ↓

00	00	00	00	00	00	10	04

```
0xffff10018
```

```
0xffff10010
```

```
0xffff10008
```

```
0xffff10000
```

```
0xffff0fff8 ←
```

00	00	00	00	00	00	00	06	%rax
00	00	00	00	00	00	22	03	%rip
00	00	00	0f	ff	f0	ff	f8	%rsp

Our function `func` puts its return value (6) into `%rax`, then executes the `ret` instruction. At this point, the address of the instruction following the `call` is at the end of the stack.

Call and Return

0x1000: call func
 → 0x1004: addq \$3, %rax

0x2000: func:

 0x2200: movq \$6, %rax
 0x2203: ret

stack growth ↓

00	00	00	00	00	00	10	04

0xffff10018
 0xffff10010
 0xffff10008
 0xffff10000 ←
 0xffff0fff8

00	00	00	00	00	00	00	06	%rax
00	00	00	00	00	00	10	04	%rip
00	00	00	0f	ff	f1	00	00	%rsp

The address at the end of the stack (0x1004) is popped off the stack and into %rip. Thus execution resumes at the instruction following the call and %rsp is incremented by 8. The function's return value is in %rax, for access by its caller.

Arguments and Local Variables

```
int mainfunc() {  
    long array[3] =  
        {2,117,-6};  
    long sum =  
        ASum(array, 3);  
    ...  
    return sum;  
}  
  
long ASum(long *a,  
          unsigned long size) {  
    long i, sum = 0;  
    for (i=0; i<size; i++)  
        sum += a[i];  
    return sum;  
}
```

- **Local variables usually allocated on stack**
- **Arguments to functions pushed onto stack**
- **Local variables may be put in registers (and thus not on stack)**

We explore these two functions in the next set of slides, looking at how arguments and local variables are stored on the stack.

Arguments and Local Variables

```
mainfunc:
    pushq %rbp                # save old %rbp
    movq %rsp, %rbp          # set %rbp to point to stack frame
    subq $32, %rsp           # alloc. space for locals (array and sum)
    movq $2, -32(%rbp)        # initialize array[0]
    movq $117, -24(%rbp)      # initialize array[1]
    movq $-6, -16(%rbp)       # initialize array[2]
    pushq $3                  # push arg 2
    leaq -32(%rbp), %rax      # array address is put in %rax
    pushq %rax               # push arg 1
    call ASum
    addq $16, %rsp            # pop args
    movq %rax, -8(%rbp)       # copy return value to sum
    addq $32, %rsp            # pop locals
    popq %rbp                # pop and restore old %rbp
    ret
```

Here we have compiled code for *mainfunc*. We'll work through this in detail in upcoming slides.

A function's stack frame is that part of the stack that holds its arguments, local variables, etc. In this example code, register `%rbp` points to a known location towards the beginning of the stack frame so that the arguments and local variables are located as offsets from what `%rbp` points to.

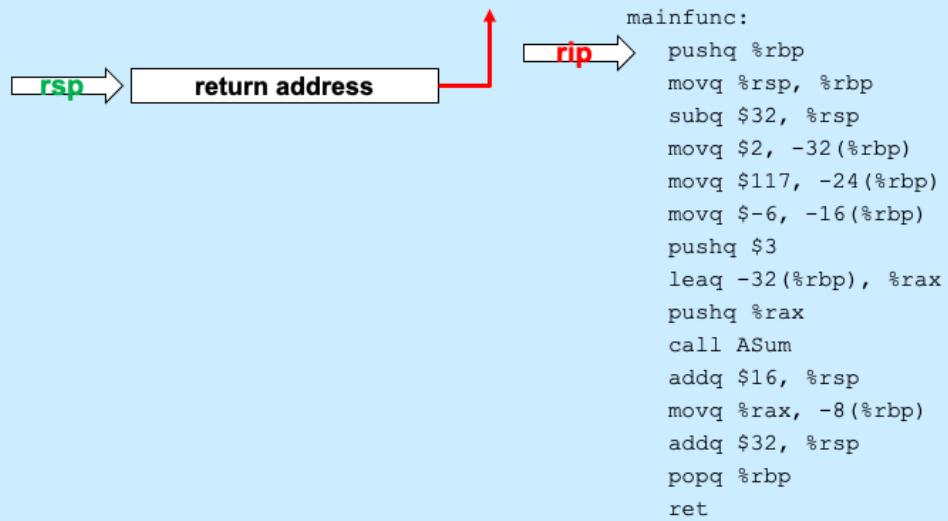
Note, as will be explained, this is not what one would see when compiling it for department computers, on which arguments are passed using registers.

Arguments and Local Variables

```
ASum:
    pushq %rbp                # save old %rbp
    movq %rsp, %rbp          # set %rbp to point to stack frame
    movq $0, %rcx             # i in %rcx
    movq $0, %rax             # sum in %rax
    movq 16(%rbp), %rdx        # copy arg 1 (array) into %rdx
loop:
    cmpq 24(%rbp), %rcx        # i < size?
    jge done
    addq (%rdx,%rcx,8), %rax    # sum += a[i]
    incq %rcx                  # i++
    ja loop
done:
    popq %rbp                 # pop and restore %rbp
    ret
```

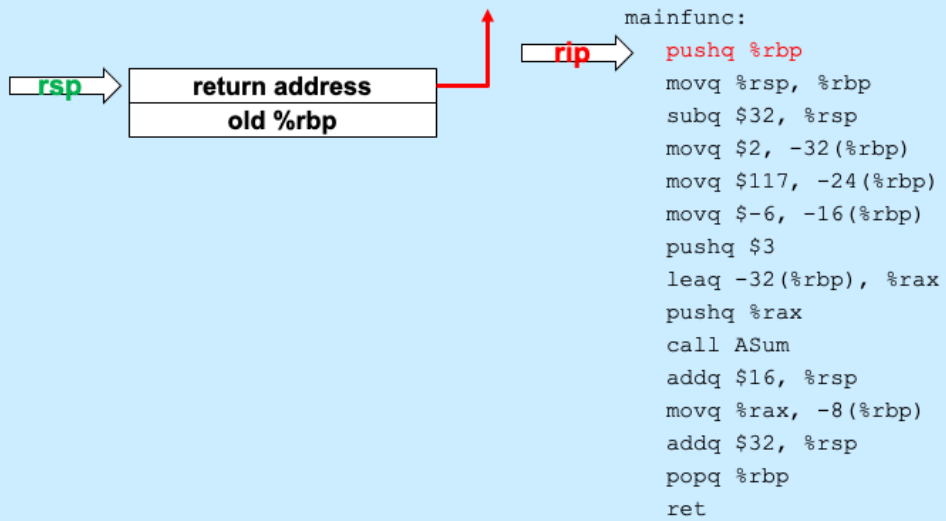
And here is the compiled code for *ASum*. The same caveats as given for the previous slide apply to this one as well.

Enter mainfunc



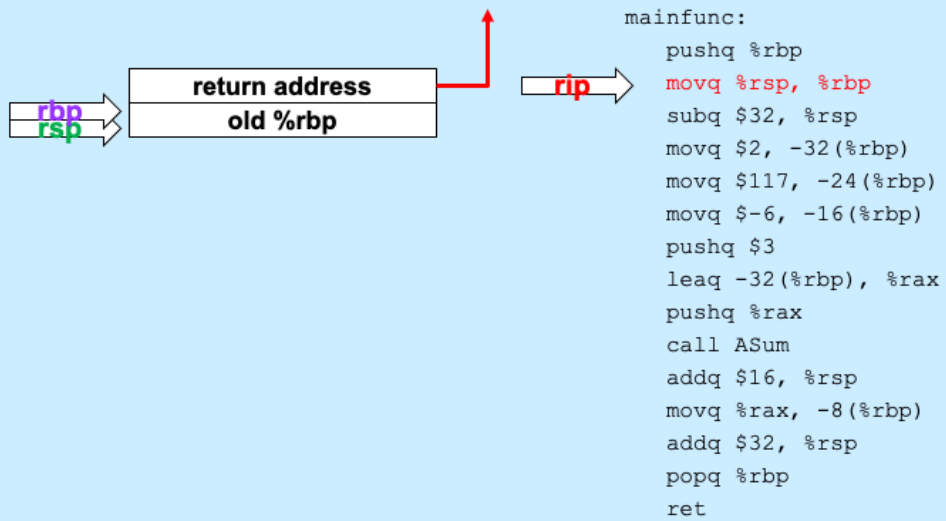
On entry to *mainfunc*, `%rsp` points to the caller's return address.

Enter mainfunc



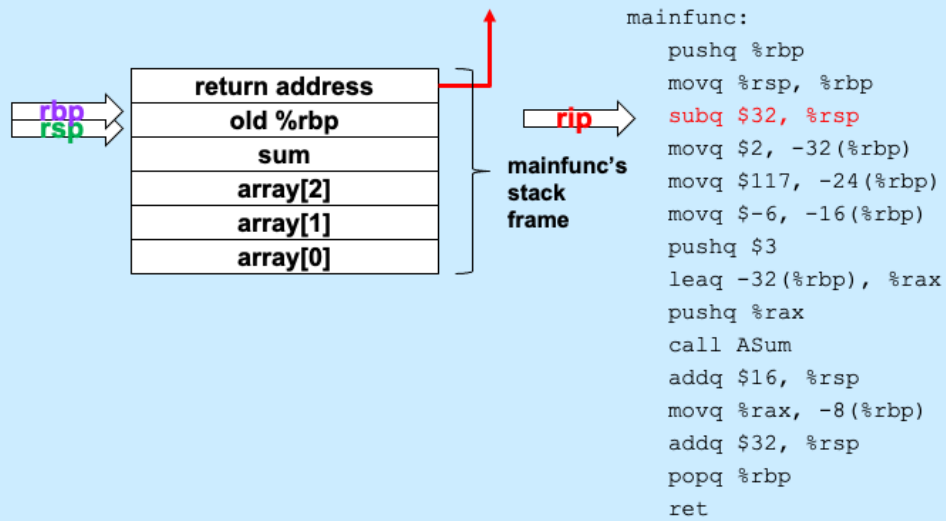
On entry to *mainfunc*, `%rsp` points to the caller's return address.

Setup Frame



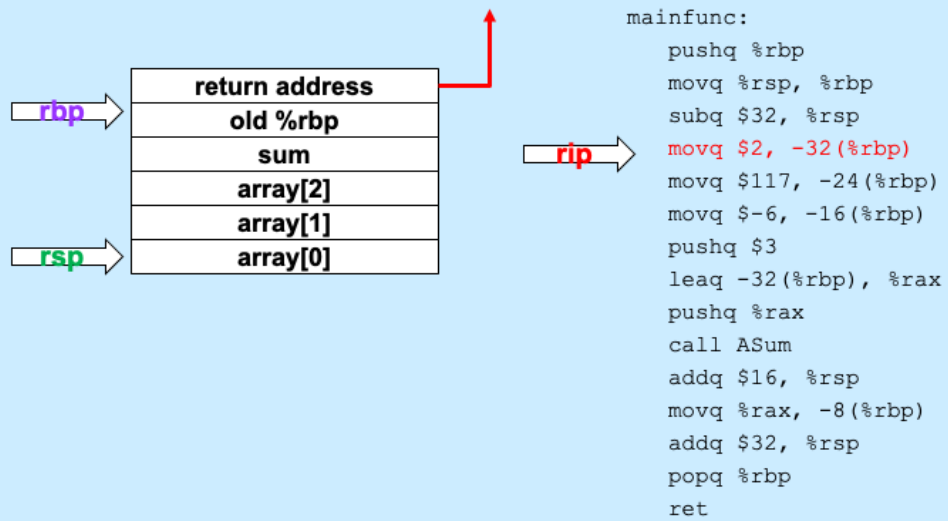
The first thing done by *mainfunc* is to save the caller's `%rbp` by pushing it onto the stack.

Allocate Local Variables



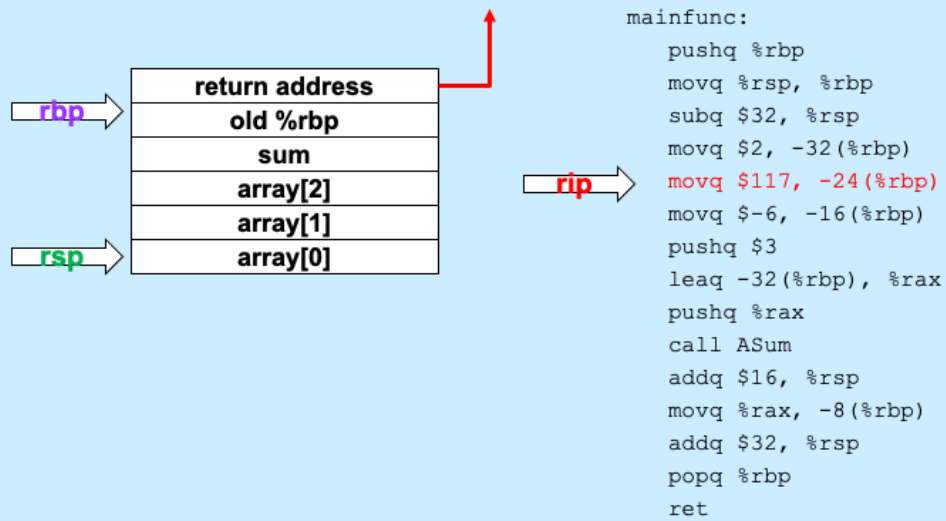
Next, space for *mainfunc*'s local variables is allocated on the stack by decrementing `%rsp` by their total size (32 bytes). At this point we have *mainfunc*'s stack frame in place.

Initialize Local Array

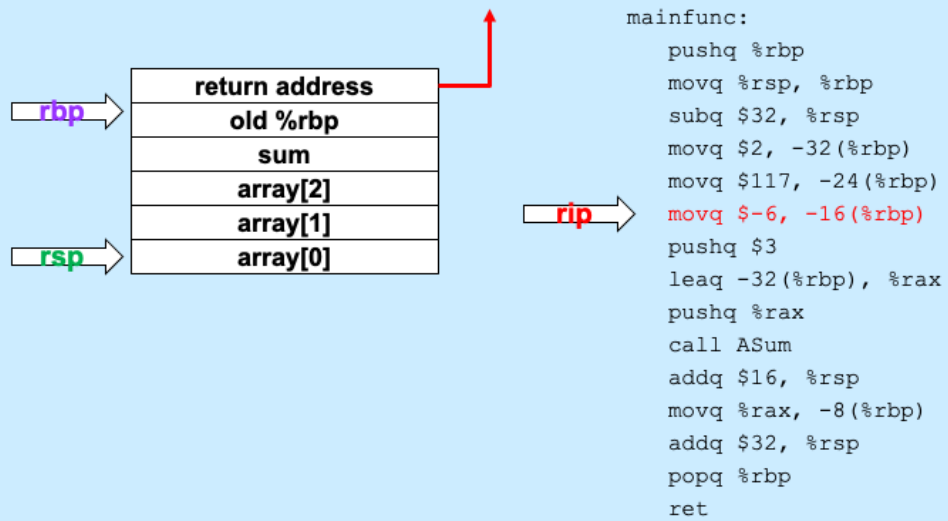


ASum now initializes the stack space containing its local variables.

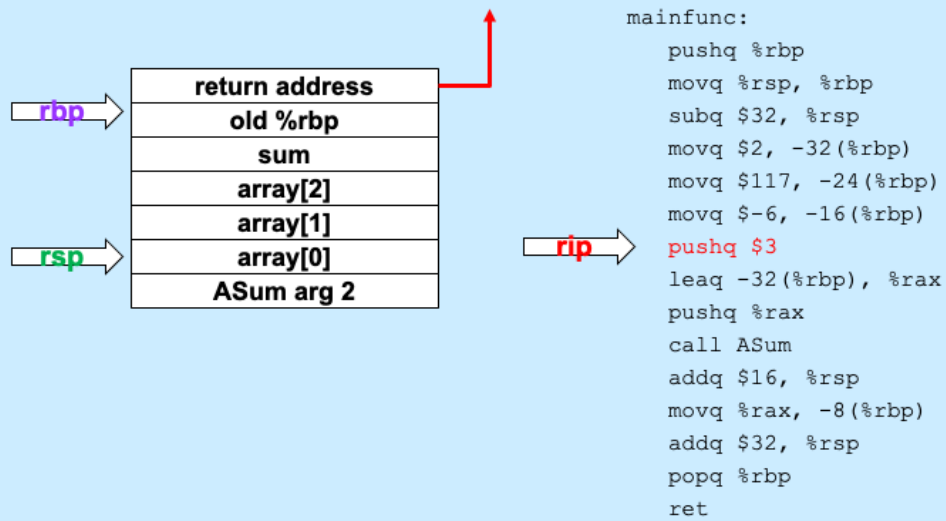
Initialize Local Array



Initialize Local Array

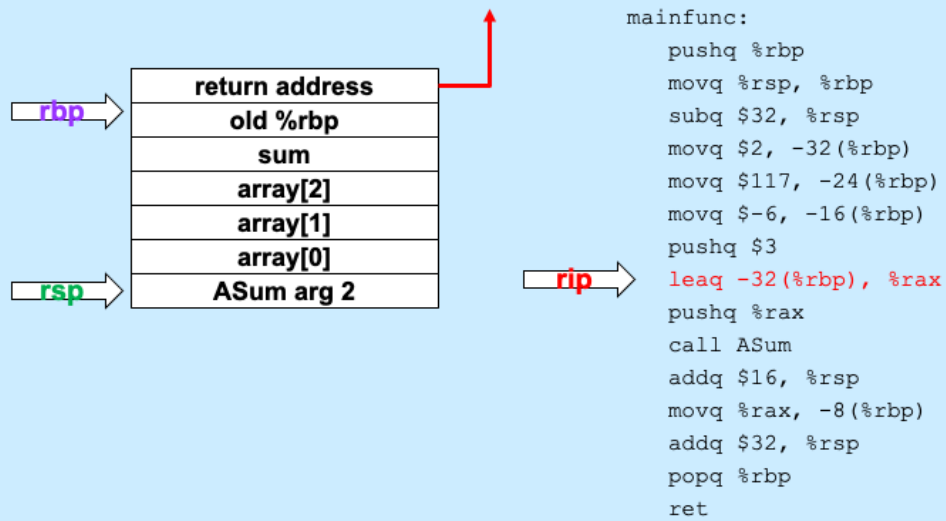


Push Second Argument



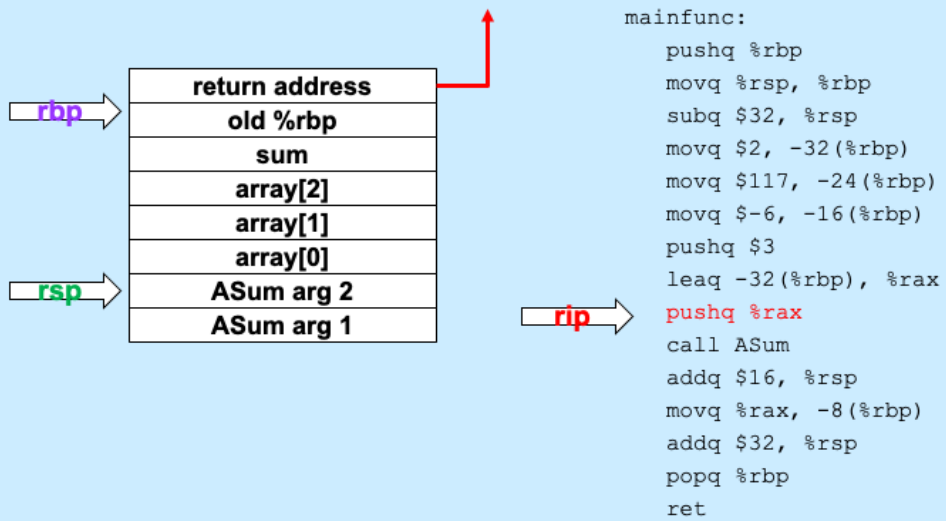
The second argument (3) to *ASum* is pushed onto the stack.

Get Array Address



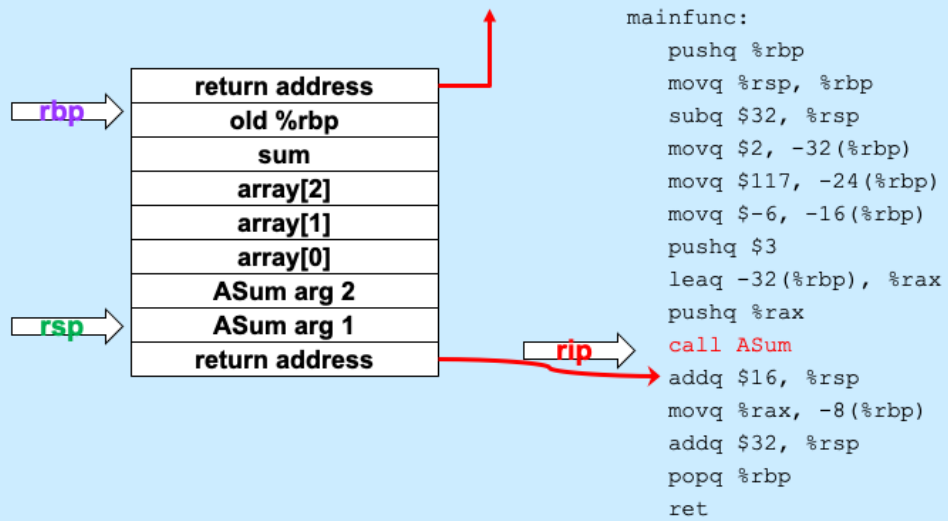
In preparation for pushing the first argument to *ASum* onto the stack, the address of the array is put into `%rax`.

Push First Argument



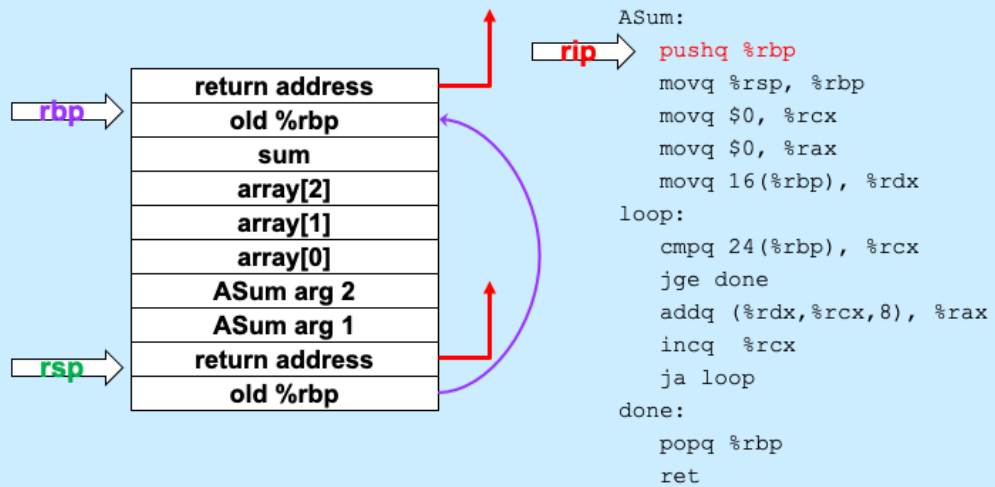
And finally, the address of the array is pushed onto the stack as *ASum*'s first argument.

Call ASum



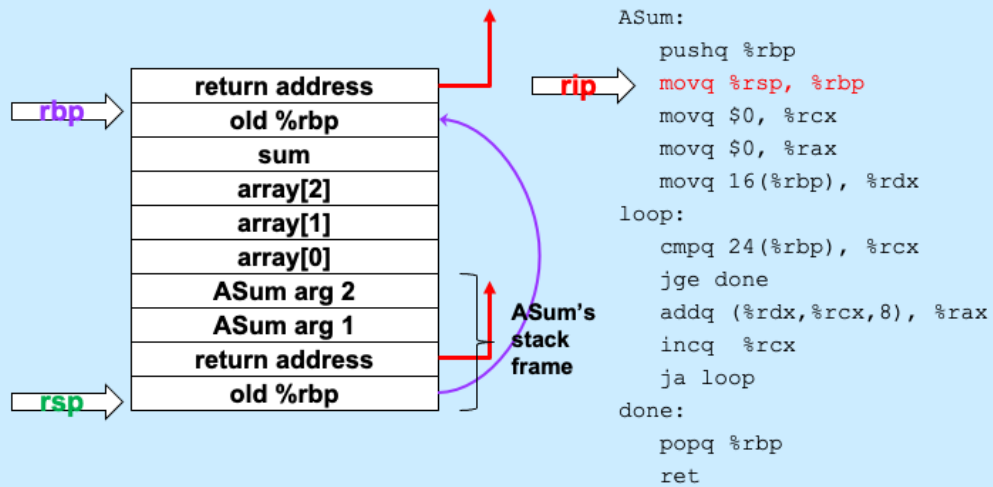
mainfunc now calls *ASum*, pushing its return address onto the stack.

Enter ASum



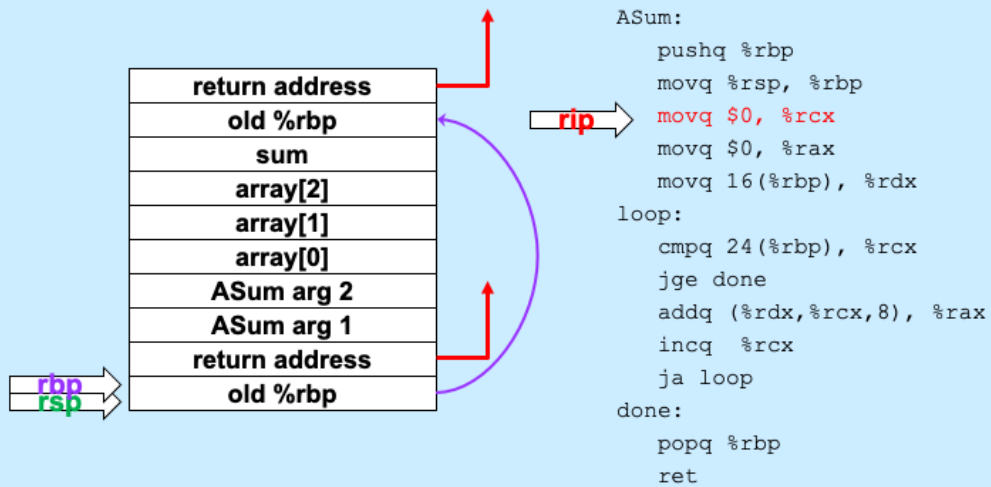
As on entry to *mainfunc*, `%rbp` is saved by pushing it onto the stack.

Setup Frame



`%rbp` is now modified to point into *ASum*'s stack frame.

Execute the Function



ASum's code is now executed, summing the contents of its first argument and storing the result in `%rax`.

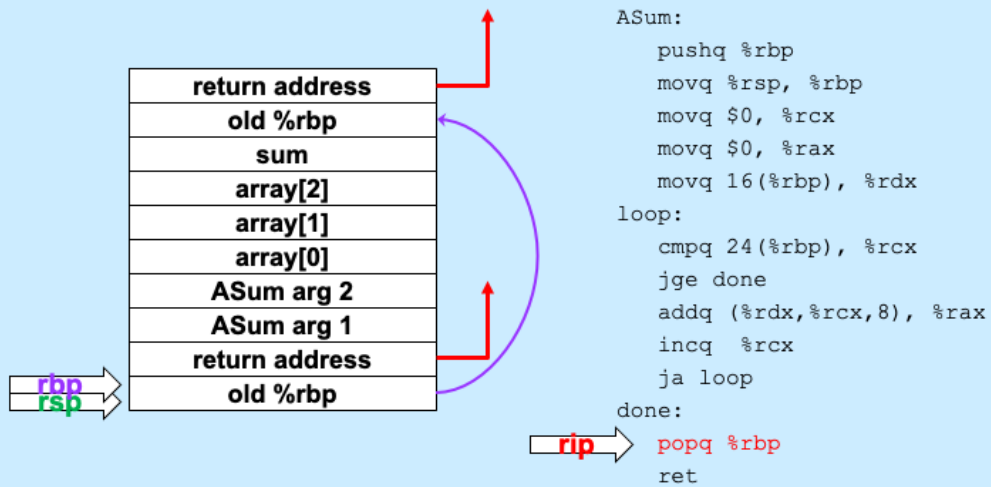
Quiz 1

What's at 24(%rbp)?

- a) a local variable**
- b) the first argument to ASum**
- c) the second argument to ASum**
- d) something else**

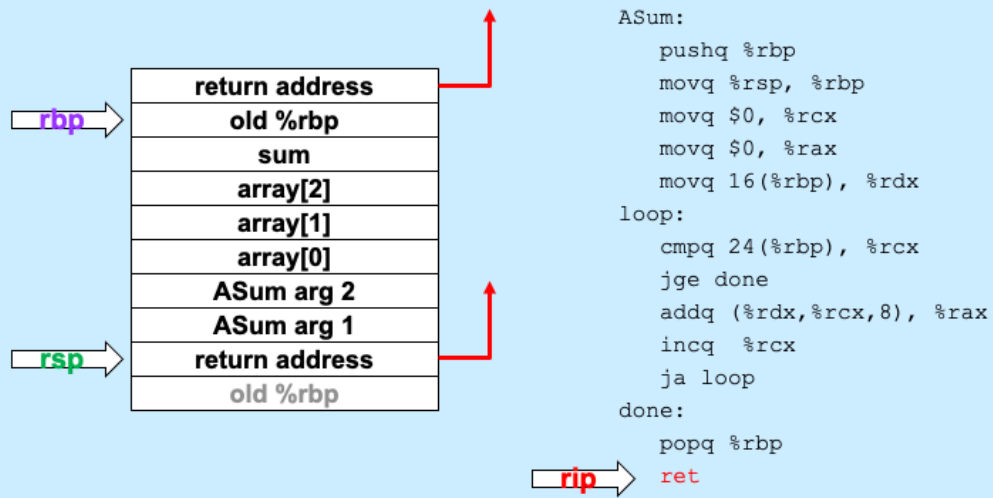
```
ASum:
    pushq %rbp
    movq %rsp, %rbp
    movq $0, %rcx
    movq $0, %rax
    movq 16(%rbp), %rdx
loop:
    cmpq 24(%rbp), %rcx
    jge done
    addq (%rdx,%rcx,8), %rax
    incq %rcx
    ja loop
done:
    popq %rbp
    ret
```

Prepare to Return



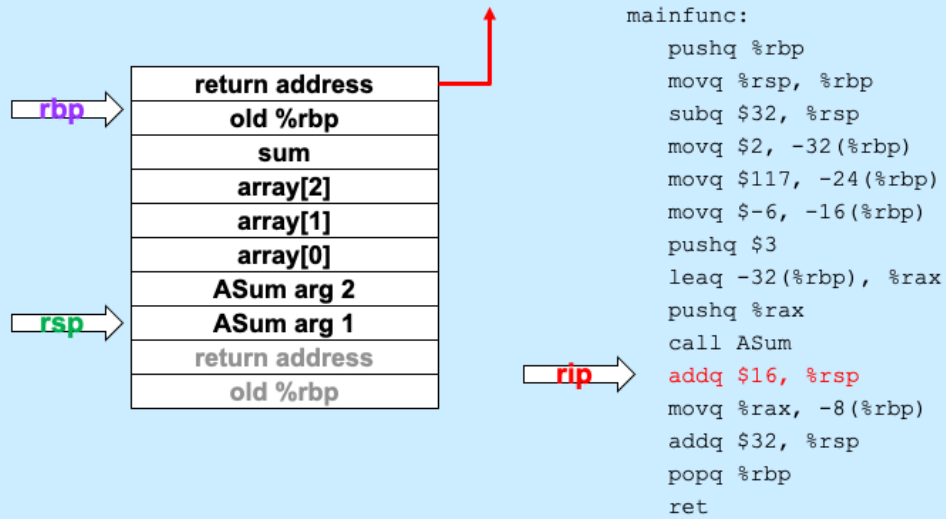
In preparation for returning to its caller, *ASum* restores the previous value of `%rbp` by popping it off the stack.

Return



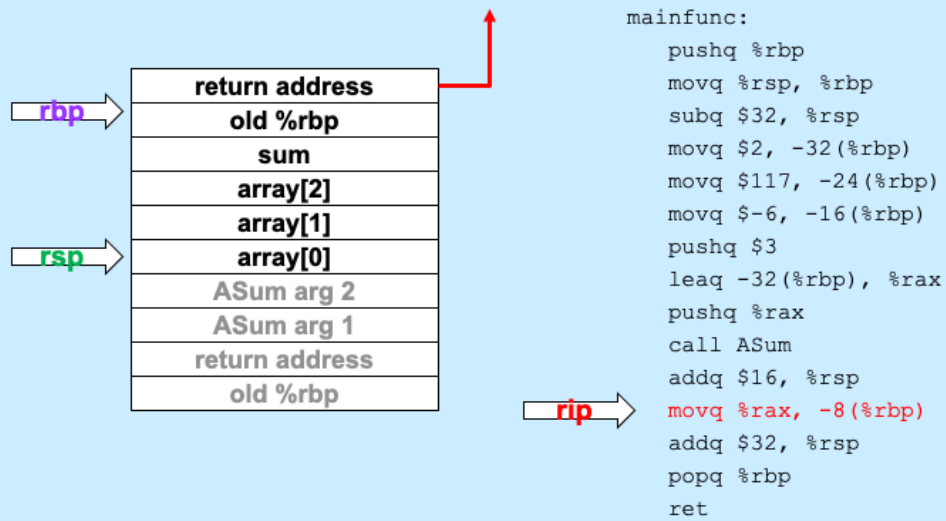
ASum returns by popping the return address off the stack and into `%rip`, so that execution resumes in its caller (*mainfunc*).

Pop Arguments



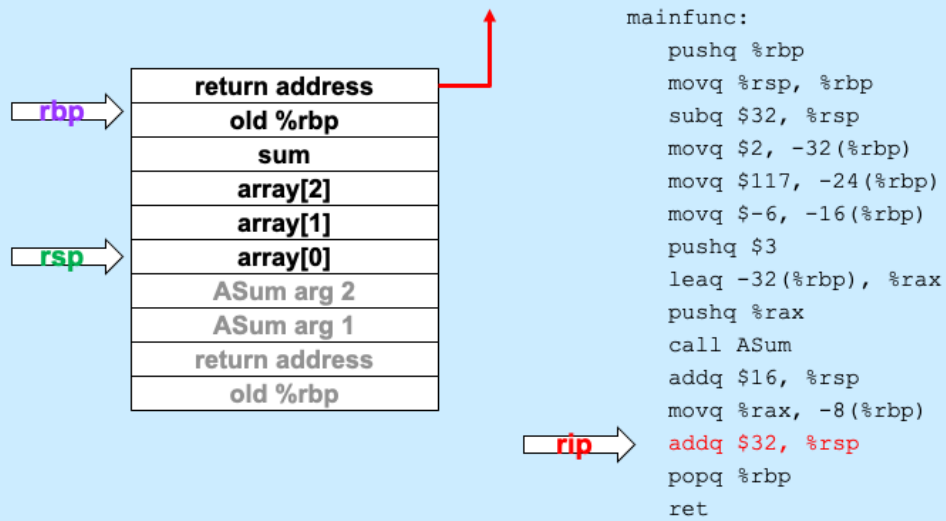
mainfunc no longer needs the arguments it had pushed onto the stack for *ASum*, so it pops them off the stack by adding their total size to %rsp.

Save Return Value



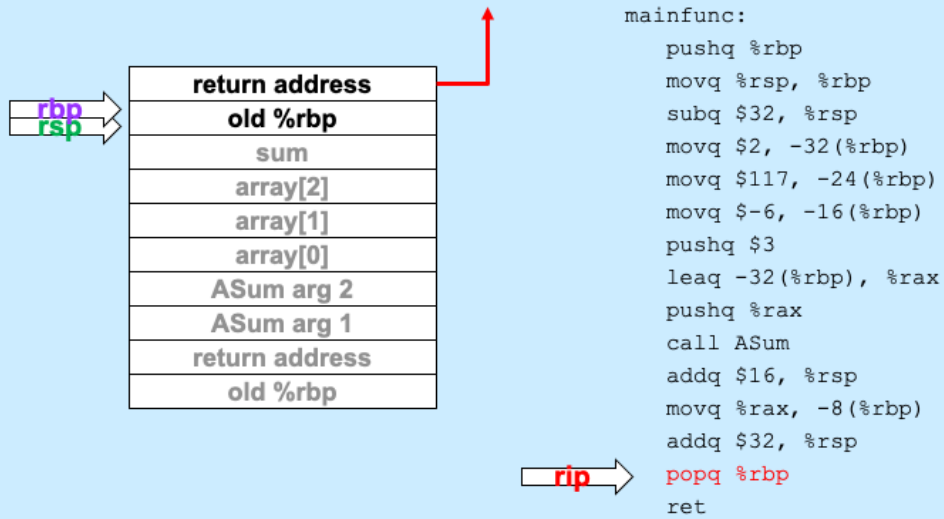
The value returned by *ASum* (in `%rax`) is copied into the local variable `sum` (which is in *mainfunc*'s stack frame).

Pop Local Variables



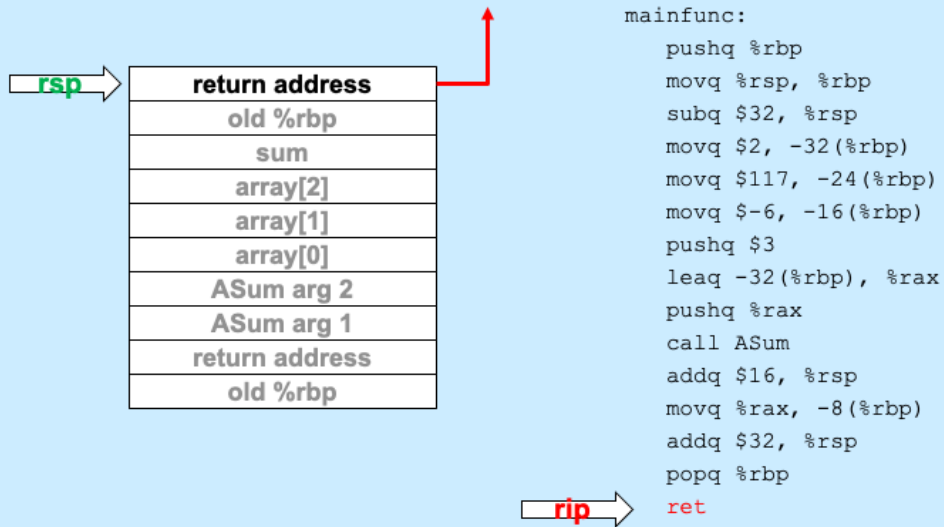
mainfunc is about to return, so it pops its local variables off the stack (by adding their total size to `%rsp`).

Prepare to Return



In preparation for returning, *mainfunc* restores its caller's `%rbp` by popping it off the stack.

Return



Finally, *mainfunc* returns by popping its caller's return address off the stack and into `%rip`.

Using Registers

- **ASum modifies registers:**

- %rsp
- %rbp
- %rcx
- %rax
- %rdx

- **Suppose its caller uses these registers**

```
...  
movq $33, %rcx  
movq $167, %rdx  
pushq $6  
pushq array  
call ASum  
    # assumes unmodified %rcx and %rdx  
addq $16, %rsp  
addq %rax,%rcx    # %rcx was modified!  
addq %rdx, %rcx   # %rdx was modified!
```

ASum:

```
pushq %rbp  
movq %rsp, %rbp  
movq $0, %rcx  
movq $0, %rax  
movq 16(%rbp), %rdx
```

loop:

```
cmpq 24(%rbp), %rcx  
jge done  
addq (%rdx,%rcx,8), %rax  
incq %rcx  
ja loop
```

done:

```
popq %rbp  
ret
```

ASum modified a number of registers. But suppose its caller was using these registers and depended on their values' being unchanged?

Register Values Across Function Calls

- **ASum modifies registers:**

- %rsp
- %rbp
- %rcx
- %rax
- %rdx

- **May the caller of ASum depend on its registers being the same on return?**

- **ASum saves and restores %rbp and makes no net changes to %rsp**
 - » their values are unmodified on return to its caller
- **%rax, %rcx, and %rdx are not saved and restored**
 - » their values might be different on return

ASum:

```
pushq %rbp
movq %rsp, %rbp
movq $0, %rcx
movq $0, %rax
movq 16(%rbp), %rdx
```

loop:

```
cmpq 24(%rbp), %rcx
jge done
addq (%rdx,%rcx,8), %rax
incq %rcx
ja loop
```

done:

```
popq %rbp
ret
```

Register-Saving Conventions

- **Caller-save registers**

- if the caller wants their values to be the same on return from function calls, it must save and restore them

```
pushq %rcx
call func
popq %rcx
```

- **Callee-save registers**

- if the callee wants to use these registers, it must first save them, then restore their values before returning

```
func:
    pushq %rbx
    movq $6, %rbx
    ...
    popq %rbx
```

Certain registers are designated as *caller-save*: if the caller depends on their values being the same on return as they were before the function was called, it must save and restore their values. Thus the called function (the "callee"), is free to modify these registers.

Other registers are designated as *callee-save*: if the callee function modifies their values, it must restore them to their original values before returning. Thus the caller may depend upon their values being unmodified on return from the function call.

x86-64 General-Purpose Registers: Usage Conventions

%rax	Return value	%r8	Caller saved
%rbx	Callee saved	%r9	Caller saved
%rcx	Caller saved	%r10	Caller saved
%rdx	Caller saved	%r11	Caller Saved
%rsi	Caller saved	%r12	Callee saved
%rdi	Caller saved	%r13	Callee saved
%rsp	Stack pointer	%r14	Callee saved
%rbp	Base pointer	%r15	Callee saved

Based on a slide supplied by CMU.

Here is a list of which registers are callee-save, which are caller-save, and which have special purposes. Note that this is merely a convention and not an inherent aspect of the x86-64 architecture.

Recursive Function

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

- **Registers**

- **%rax, %rdx** used without first saving
- **%rbx** used, but saved at beginning & restored at end

```
pcount_r:
    pushq %rbp
    movq %rsp, %rbp
    pushq %rbx
    movq 16(%rbp), %rbx
    movq $0, %rax
    testq %rbx, %rbx
    je .L3
    movq %rbx, %rax
    shrq $1, %rax
    pushq %rax
    call pcount_r
    addq $8, %rsp
    movq %rbx, %rdx
    andq $1, %rdx
    leaq (%rdx,%rax), %rax
.L3:
    popq %rbx
    popq %rbp
    ret
```

Adapted from a slide supplied by CMU.

We now walk through an example of the use of stacks for implementing recursive functions. Our example function, `pcount_r`, computes the number of one bits in its unsigned long argument.

Recursive Call #1


```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

• Actions

- save old value of %rbx on stack
- store x in %rbx

%rbx x

```
pcount_r:
    pushq %rbp
    movq %rsp, %rbp
    pushq %rbx
    movq 16(%rbp), %rbx
    . . .
```



Adapted from a slide supplied by CMU.

Recursive Call #2

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

```
    . . .
    movq $0, %rax
    testq %rbx, %rbx
    je .L3
    . . .
.L3:
    . . .
    ret
```

- **Actions**

- if `x == 0`, return
 - » with `%rax` set to 0

`%rbx` x

Adapted from a slide supplied by CMU.

Recall that "testq b,a" sets the condition codes based on the value of "a&b". Thus "testq %rbx,%rbx" sets ZF to zero if and only if %rbx contains zero.


Recursive Call #3

```
/* Recursive popcount */  
long pcount_r(unsigned long x) {  
    if (x == 0)  
        return 0;  
    else return  
        (x & 1) + pcount_r(x >> 1);  
}
```

```
...  
movq %rbx, %rax  
shrq $1, %rax  
pushq %rax  
call pcount_r  
...
```

- **Actions**
 - push $x \gg 1$ on stack as arg
 - make recursive call
- **Effect**
 - $\%rax$ set to function result
 - $\%rbx$ still has value of x

$\%rbx$ x



Adapted from a slide supplied by CMU.

Recursive Call #4

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

```
. . .
movq    %rbx, %rdx
addq    $8, %rsp
andq    $1, %rdx
leaq    (%rdx,%rax), %rax
. . .
```

- **Assume**
 - %rax holds value from recursive call
 - %rbx holds x
- **Actions**
 - pop argument from stack
 - compute (x & 1) + computed value
- **Effect**
 - %rax set to function result

%rbx

x

Adapted from a slide supplied by CMU.


Recursive Call #5

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

```
...
L3:
    popq    %rbx
    popq    %rbp
    ret
```

• Actions

- restore values of %rbx and %rbp
- return (and pop return address)



Adapted from a slide supplied by CMU.

Observations About Recursion

- **Handled without special consideration**
 - **stack frames mean that each function call has private storage**
 - » saved registers & local variables
 - » saved return pointer
 - **register-saving conventions prevent one function call from corrupting another's data**
 - **stack discipline follows call / return pattern**
 - » if P calls Q, then Q returns before P
 - » last-in, first-out
- **Also works for mutual recursion**
 - P calls Q; Q calls P

Supplied by CMU.

Passing Arguments in Registers

- **Observations**

- **accessing registers is much faster than accessing primary memory**
 - » if arguments were in registers rather than on the stack, speed would increase
- **most functions have just a few arguments**

- **Actions**

- **change calling conventions so that the first six arguments are passed in registers**
 - » in caller-save registers
- **any additional arguments are pushed on the stack**

Why Bother with a Base Pointer?

- **It (%rbp) points to the beginning of the stack frame**
 - making it easy for people to figure out where things are in the frame
 - but people don't execute the code ...
- **The stack pointer always points somewhere within the stack frame**
 - it moves about, but the compiler knows where it is pointing
 - » a local variable might be at 8(%rsp) for one instruction, but at 16(%rsp) for a subsequent one
 - » tough for people, but easy for the compiler
- **Thus the base pointer is superfluous**
 - it can be used as a general-purpose register

If one gives gcc the `-O0` flag (which turns off all optimization) when compiling, the base pointer (%rbp) will be used as in IA32: it is set to point to the stack frame and the arguments are copied from the registers into the stack frame. This clearly slows down the execution of the function, but makes the code easier for humans to read (and was done for the traps assignment).

x86-64 General-Purpose Registers: Updated Usage Conventions

%rax	Return value	%r8	Argument #5
%rbx	Callee saved	%r9	Argument #6
%rcx	Argument #4	%r10	Caller saved
%rdx	Argument #3	%r11	Caller Saved
%rsi	Argument #2	%r12	Callee saved
%rdi	Argument #1	%r13	Callee saved
%rsp	Stack pointer	%r14	Callee saved
%rbp	Callee saved	%r15	Callee saved

Supplied by CMU.

Recursive Function (Improved)

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

- **Registers**

- the single argument (x) is passed in %rdi
- %rbx is a callee-saved register and thus is saved and restored
- %rax is caller-saved
- %rbp isn't used

```
pcount_r:
    movq    $0, %rax
    testq   %rdi, %rdi
    je      .L8
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    $1, %rdi
    call    pcount_r
    andq    $1, %rbx
    addq    %rbx, %rax
    popq    %rbx
.L8:
    ret
```

The assembler code in the slide is from gcc with the -O1 flag (turning on a small amount of optimization). The argument is passed in %rdi and there is no use of %rbp as a base pointer.

Summary

- **What's pushed on the stack**
 - **return address**
 - **saved registers**
 - » **caller-saved by the caller**
 - » **callee-saved by the callee**
 - **local variables**
 - **function parameters**
 - » **those too large to be in registers (structs)**
 - » **those beyond the six that we have registers for**
 - **large return values (structs)**
 - » **caller allocates space on stack**
 - » **callee copies return value to that space**

Quiz 2

Suppose function A is compiled using the convention that %rbp is used as the base pointer, pointing to the beginning of the stack frame. Function B is compiled using the convention that there's no need for a base pointer. Will there be any problems if A calls B or if B calls A?

- a) Neither case will work**
- b) A calling B works, but B calling A doesn't**
- c) B calling A works, but A calling B doesn't**
- d) Both work**

Tail Recursion

```
int factorial(int x) {  
    if (x == 1)  
        return x;  
    else  
        return  
            x*factorial(x-1);  
}
```

```
int factorial(int x) {  
    return f2(x, 1);  
}  
  
int f2(int a1, int a2) {  
    if (a1 == 1)  
        return a2;  
    else  
        return  
            f2(a1-1, a1*a2);  
}
```

The slide shows two implementations of the factorial function. Both use recursion. In the version on the left, the result of each recursive call is used within the invocation that issued the call. In the second, the result of each recursive call is simply returned. This is known as *tail recursion*.

No Tail Recursion (1)


x: 6
return addr
x: 5
return addr
x: 4
return addr
x: 3
return addr
x: 2
return addr
x: 1
return addr

Here we look at the stack usage for the version without tail recursion. Note that we have as many stack frames as the value of the argument; the results of the calls are combined after the stack reaches its maximum size.

No Tail Recursion (2)

x: 6	ret: 720
return addr	
x: 5	ret: 120
return addr	
x: 4	ret: 24
return addr	
x: 3	ret: 6
return addr	
x: 2	ret: 2
return addr	
x: 1	ret: 1
return addr	

Tail Recursion



With tail recursion, since the result of the recursive call is not used by the issuing stack frame, it's possible to reuse the issuing stack frame to handle the recursive invocation. Thus rather than push a new stack frame on the stack, the current one is written over. Thus the entire sequence of recursive calls can be handled within a single stack frame.

Code: gcc -O1

```
f2:
    movl    %esi, %eax
    cmpl    $1, %edi
    je      .L5
    subq    $8, %rsp
    movl    %edi, %esi
    imull   %eax, %esi
    subl    $1, %edi
    call    f2      # recursive call!
    addq    $8, %rsp
.L5:
    ret
```

This is the result of compiling the tail-recursive version of factorial using gcc with the -O1 flag. This flag turns on a moderate level of code optimization, but not enough to cause the stack frame to be reused.

Not Using the Stack ...

```
f2:
    movl    %esi, %eax
    cmpl    $1, %edi
    je      .L5
    movl    %edi, %esi
    imull   %eax, %esi
    subl    $1, %edi
    call    f2      # recursive call!
.L5:
    ret
```

The code that moved the stack pointer down and then up again is clearly not needed, so we've removed it.

Not Recursive!

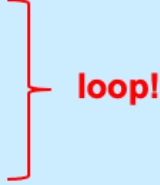
```
f2:
    movl    %esi, %eax
    cmpl    $1, %edi
    je      .L5
    movl    %edi, %esi
    imull   %eax, %esi
    subl    $1, %edi
    ja      f2      # goto!
.L5:
    ret
```

Since nothing is pushed onto the stack and nothing is popped from the stack – the code does everything in the registers %eax, %edi, and %esi -- there's really no need to call f2 recursively. Instead, it can simply jump to f2.

Code: gcc -O2

```
f2:
    cmpl    $1, %edi
    movl    %esi, %eax
    je      .L8

.L12:
    imull   %edi, %eax
    subl    $1, %edi
    cmpl    $1, %edi
    jne     .L12
.L8:
    ret
```



Here we've compiled the program using the `-O2` flag, which turns on additional optimization (at the cost of increased compile time), with the result that the recursive calls are automatically optimized away — they are replaced with a loop.

Why not always compile with `-O2`? For “production code” that is bug-free (assuming this is possible), this is a good idea. But this and other aggressive optimizations make it difficult to relate the runtime code with the source code. Thus, a runtime error might occur at some point in the program's execution, but it is impossible to determine exactly which line of the source code was in play when the error occurred.

Note that gcc can do some very impressive optimization when given the `-O2` flag. For example, in both the “normal” recursive implementation of factorial, as well as the `pcount_r` function discussed earlier, it produces assembler code that uses loops rather than recursive calls.