

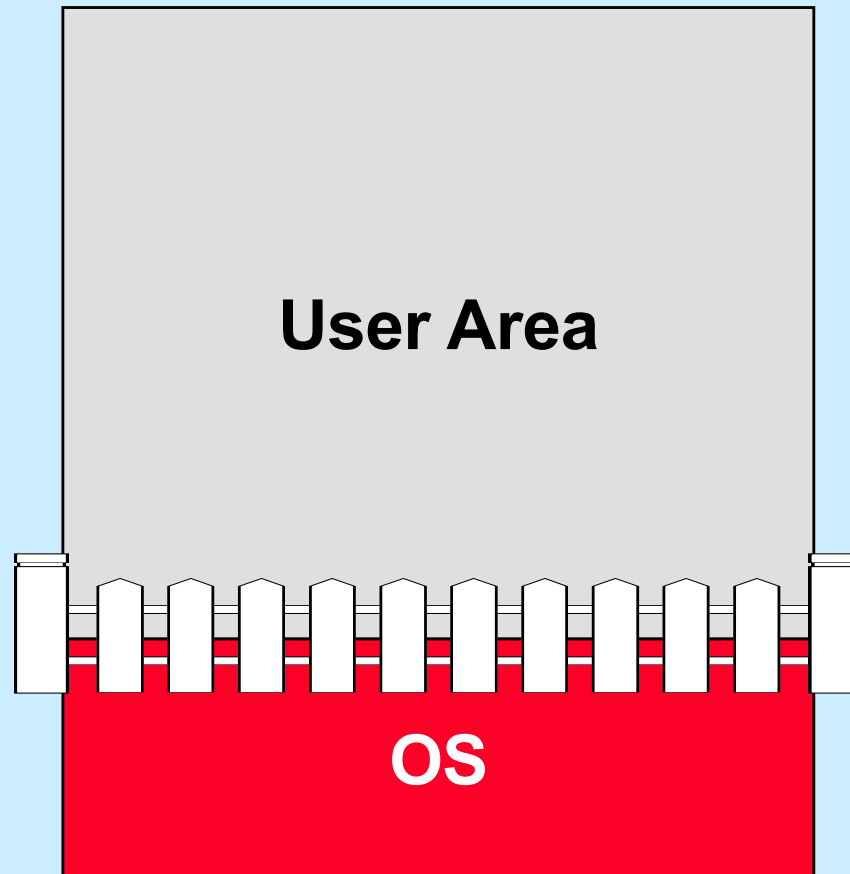
# CS 33

## Virtual Memory

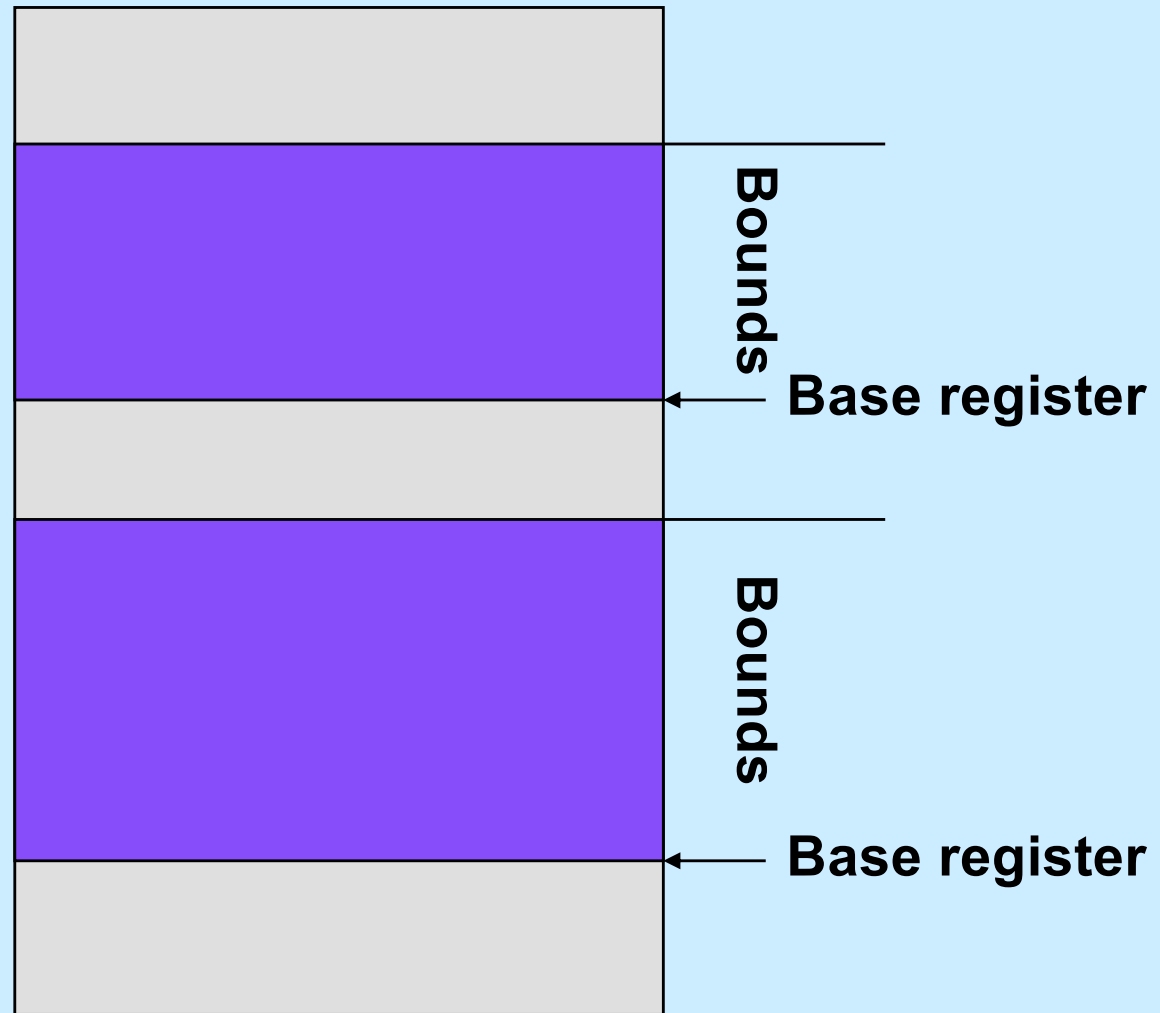
# The Address-Space Concept

- **Protect processes from one another**
- **Protect the OS from user processes**
- **Provide efficient management of available storage**

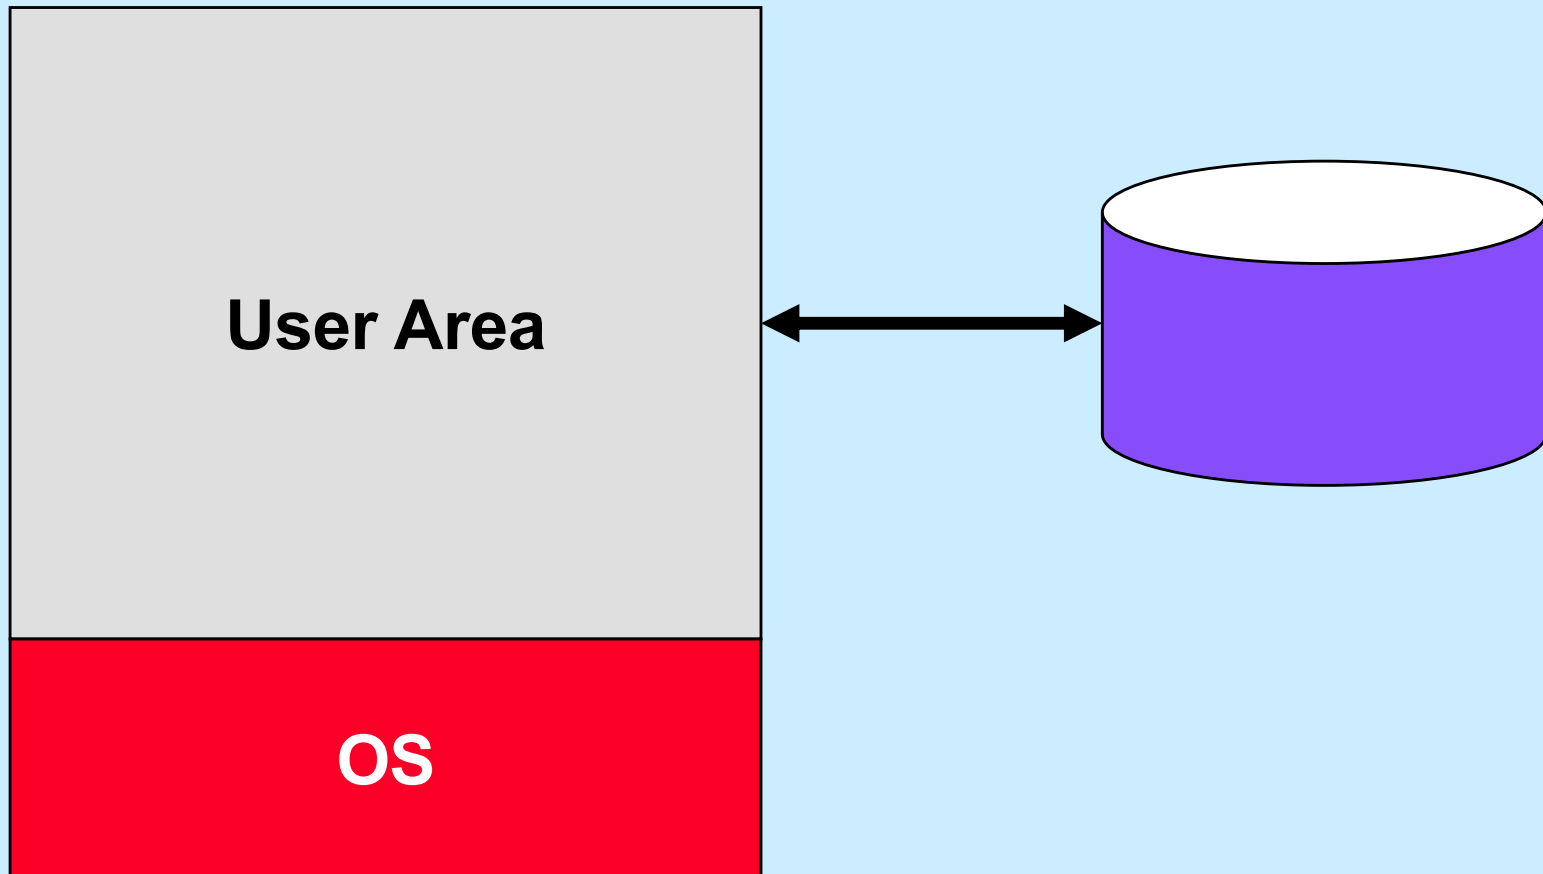
# Memory Fence



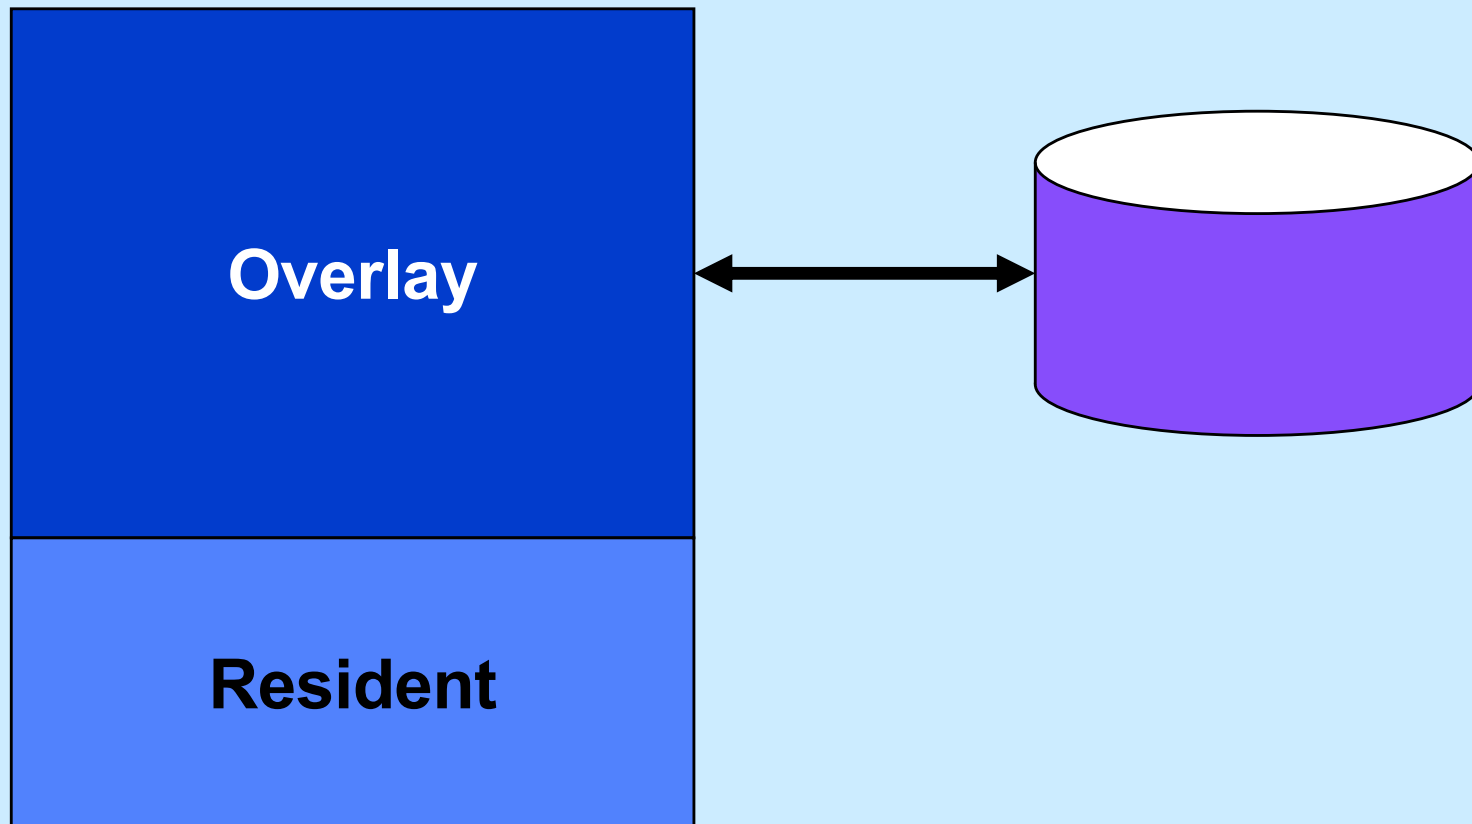
# Base and Bounds Registers



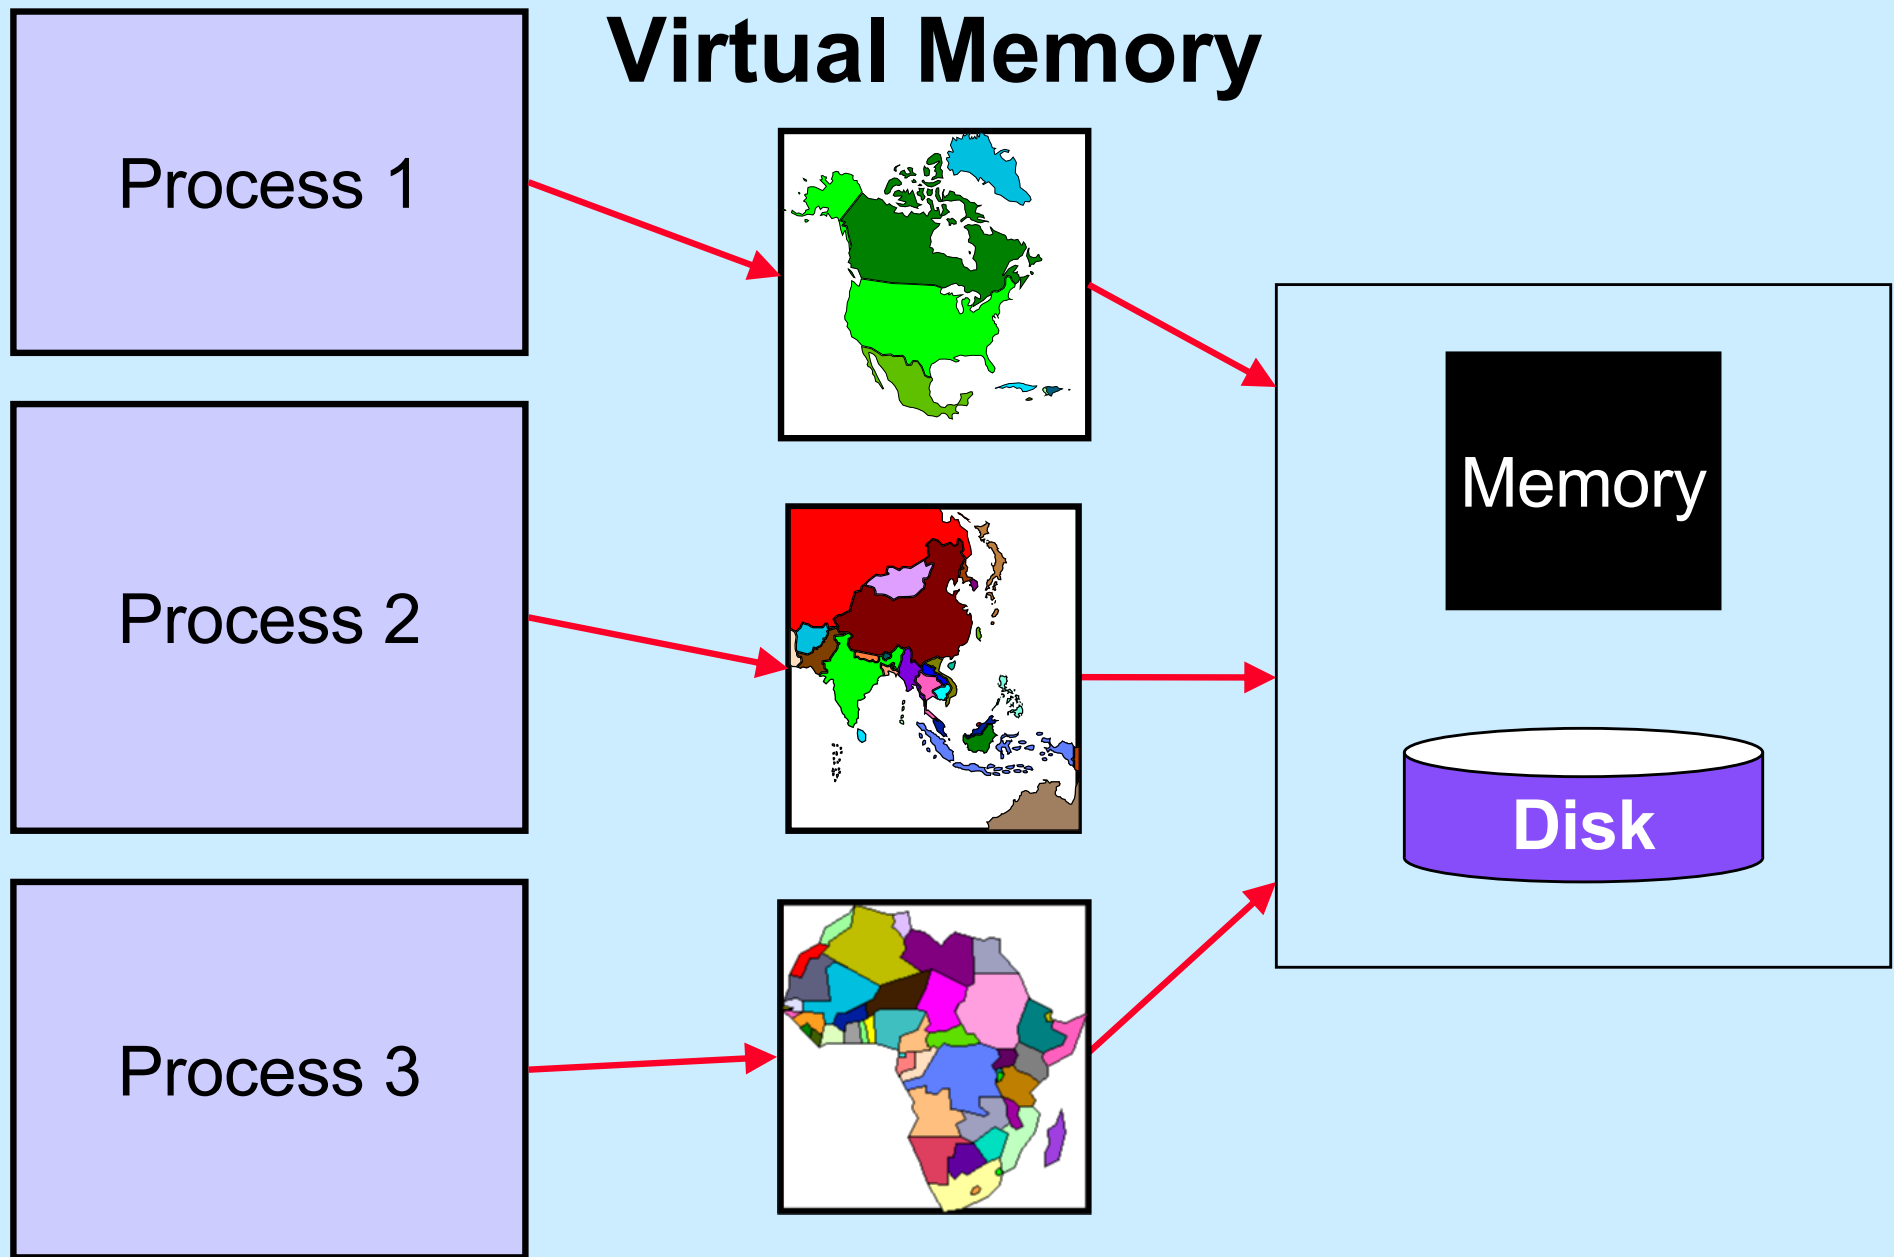
# Swapping



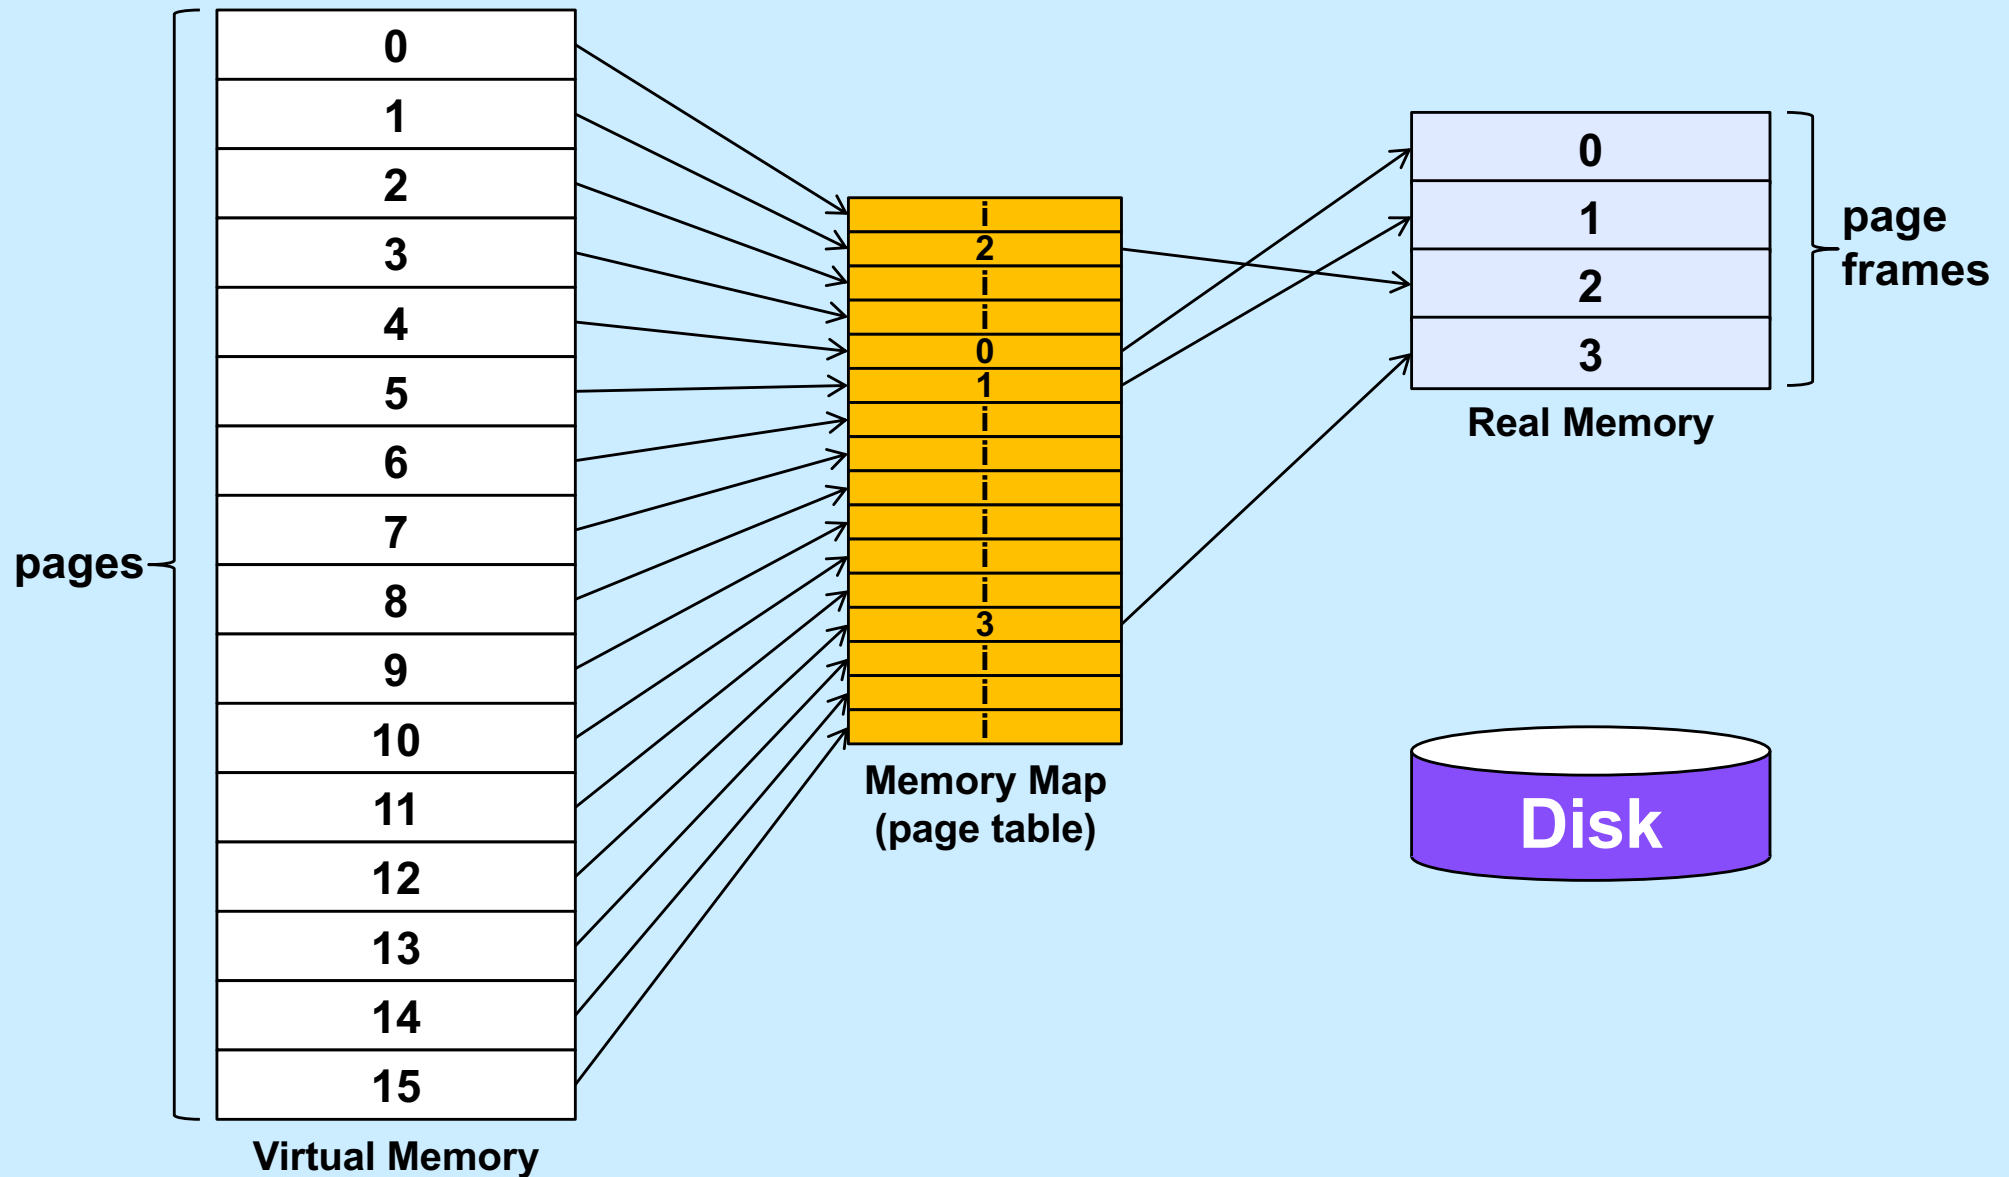
# Overlays



# Virtual Memory

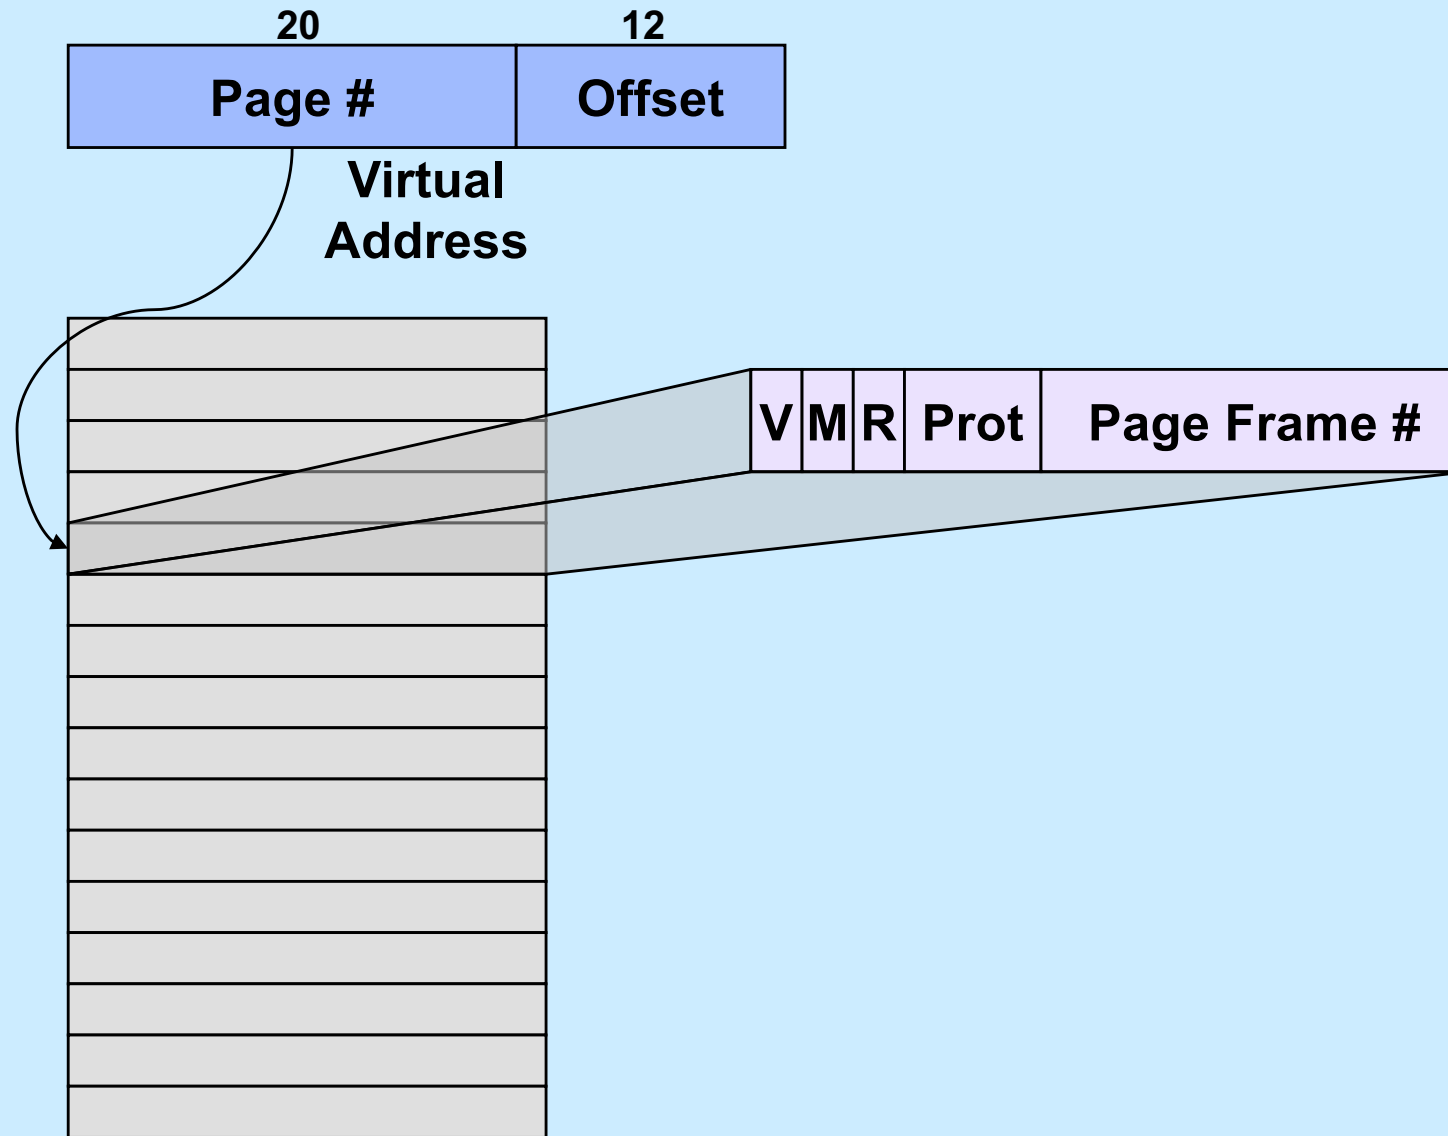


# Memory Maps





# Page Tables



# Quiz 1

**How many  $2^{12}$ -byte pages fit in a 32-bit address space?**

- a) a little over a 1000**
- b) a little over a million**
- c) a little over a billion**
- d) none of the above**

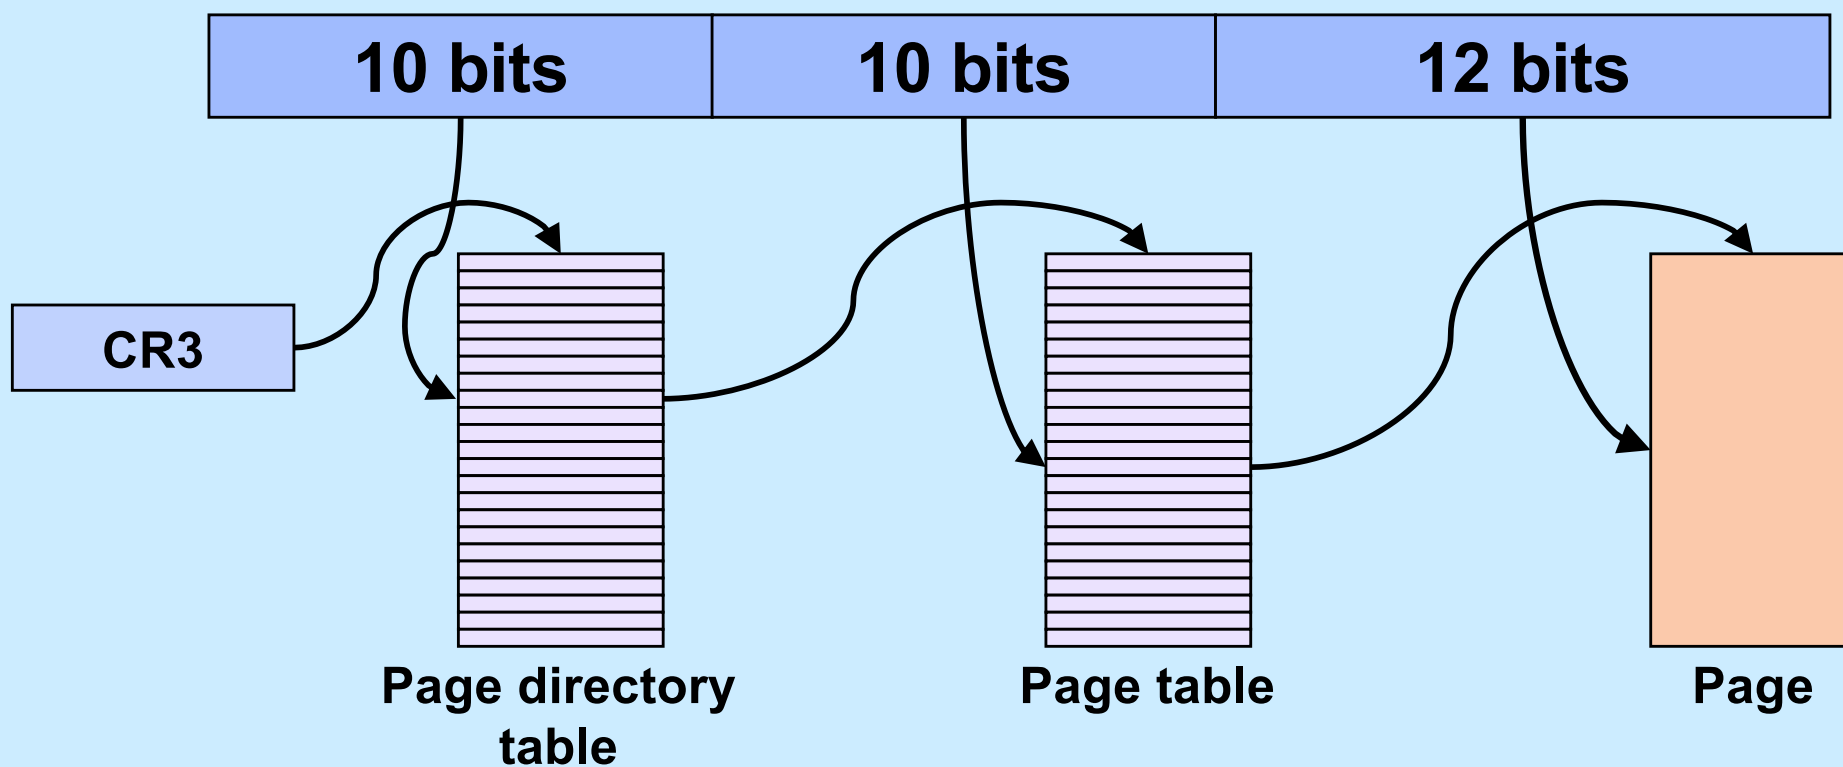
# VM is Your Friend ...

- **Not everything has to be in memory at once**
    - pages brought in (and pushed out) when needed
    - unallocated parts of the address space consume no memory
      - » e.g., hole between stack and dynamic areas
  - **What's mine is not yours** (and vice versa)
    - address spaces are disjoint
  - **Sharing is ok though ...**
    - address spaces don't have to be disjoint
      - » a single page frame may be mapped into multiple processes
  - **I don't trust you (or me)**
    - access to individual pages can be restricted
      - » read, write, execute, or any combination
-

# Page-Table Size

- **Consider a full  $2^{32}$ -byte address space**
  - assume 4096-byte ( $2^{12}$ -byte) pages
  - 4 bytes per page-table entry
  - the page table would consist of  $2^{32}/2^{12}$  ( $= 2^{20}$ ) entries
  - its size would be  $2^{22}$  bytes (or 4 megabytes)
    - » at \$100/gigabyte
      - around \$0.40
- **For a  $2^{64}$ -byte address space**
  - assume 4096-byte ( $2^{12}$ -byte) pages
  - 8 bytes per page-table entry
  - the page table would consist of  $2^{64}/2^{12}$  ( $= 2^{52}$ ) entries
  - its size would be  $2^{55}$  bytes (or 32 petabytes)
    - » at \$1/gigabyte
      - over \$33 million

# IA32 Paging

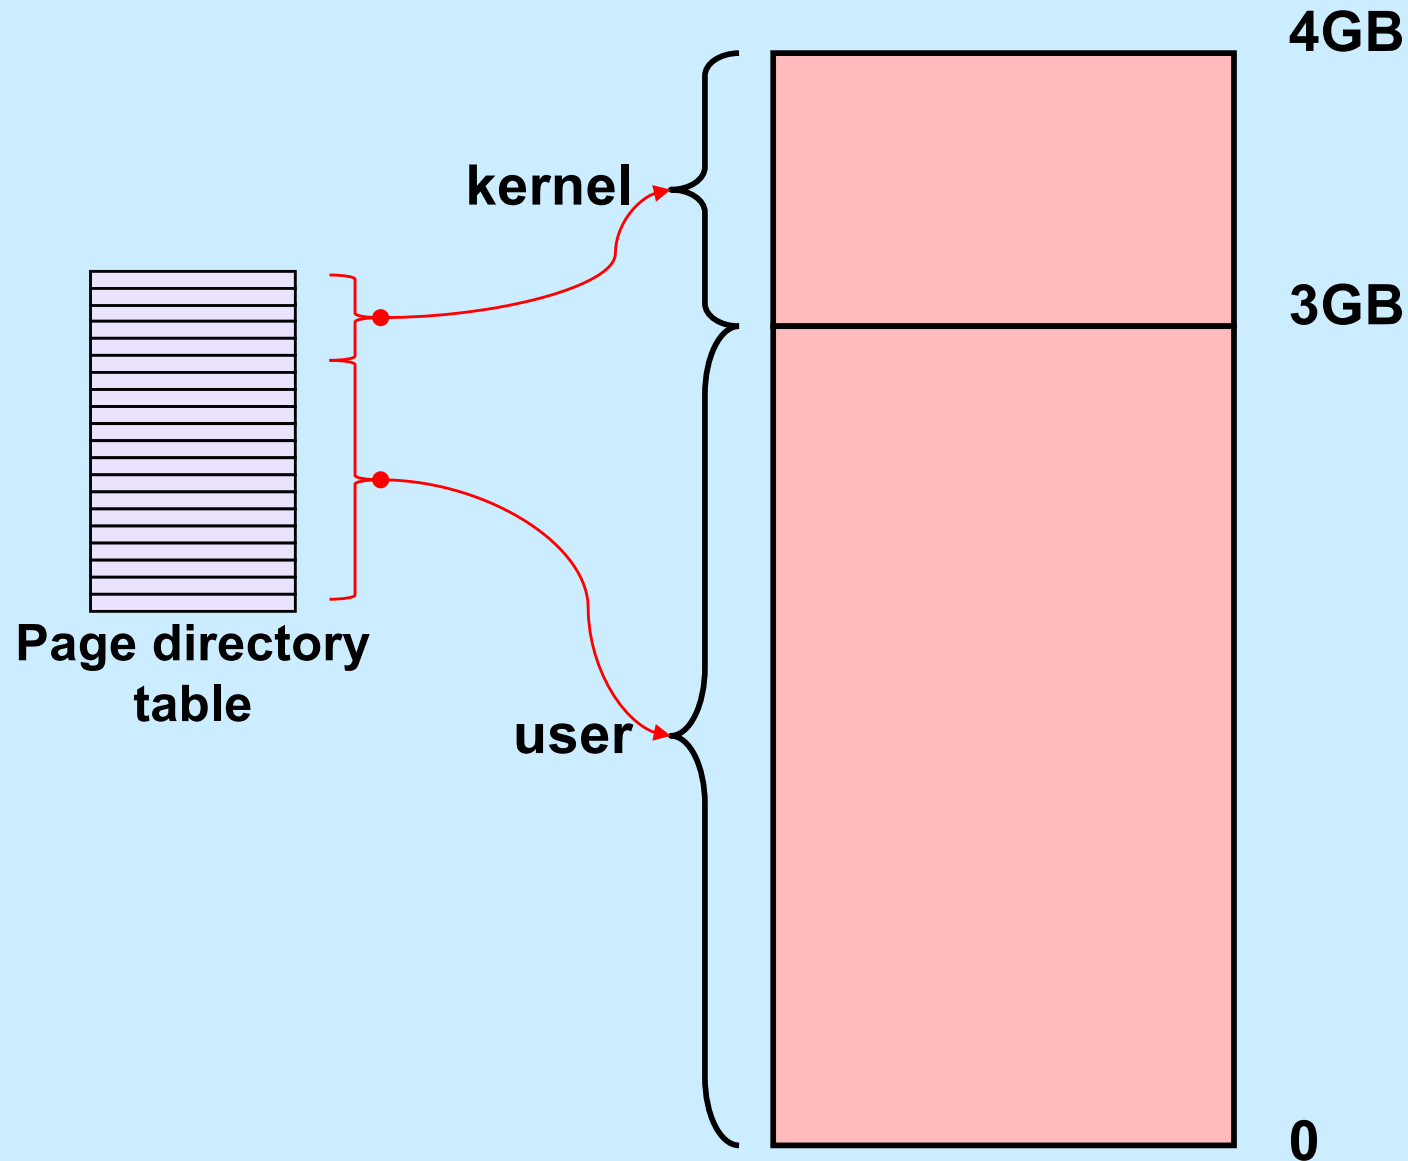


# Quiz 2

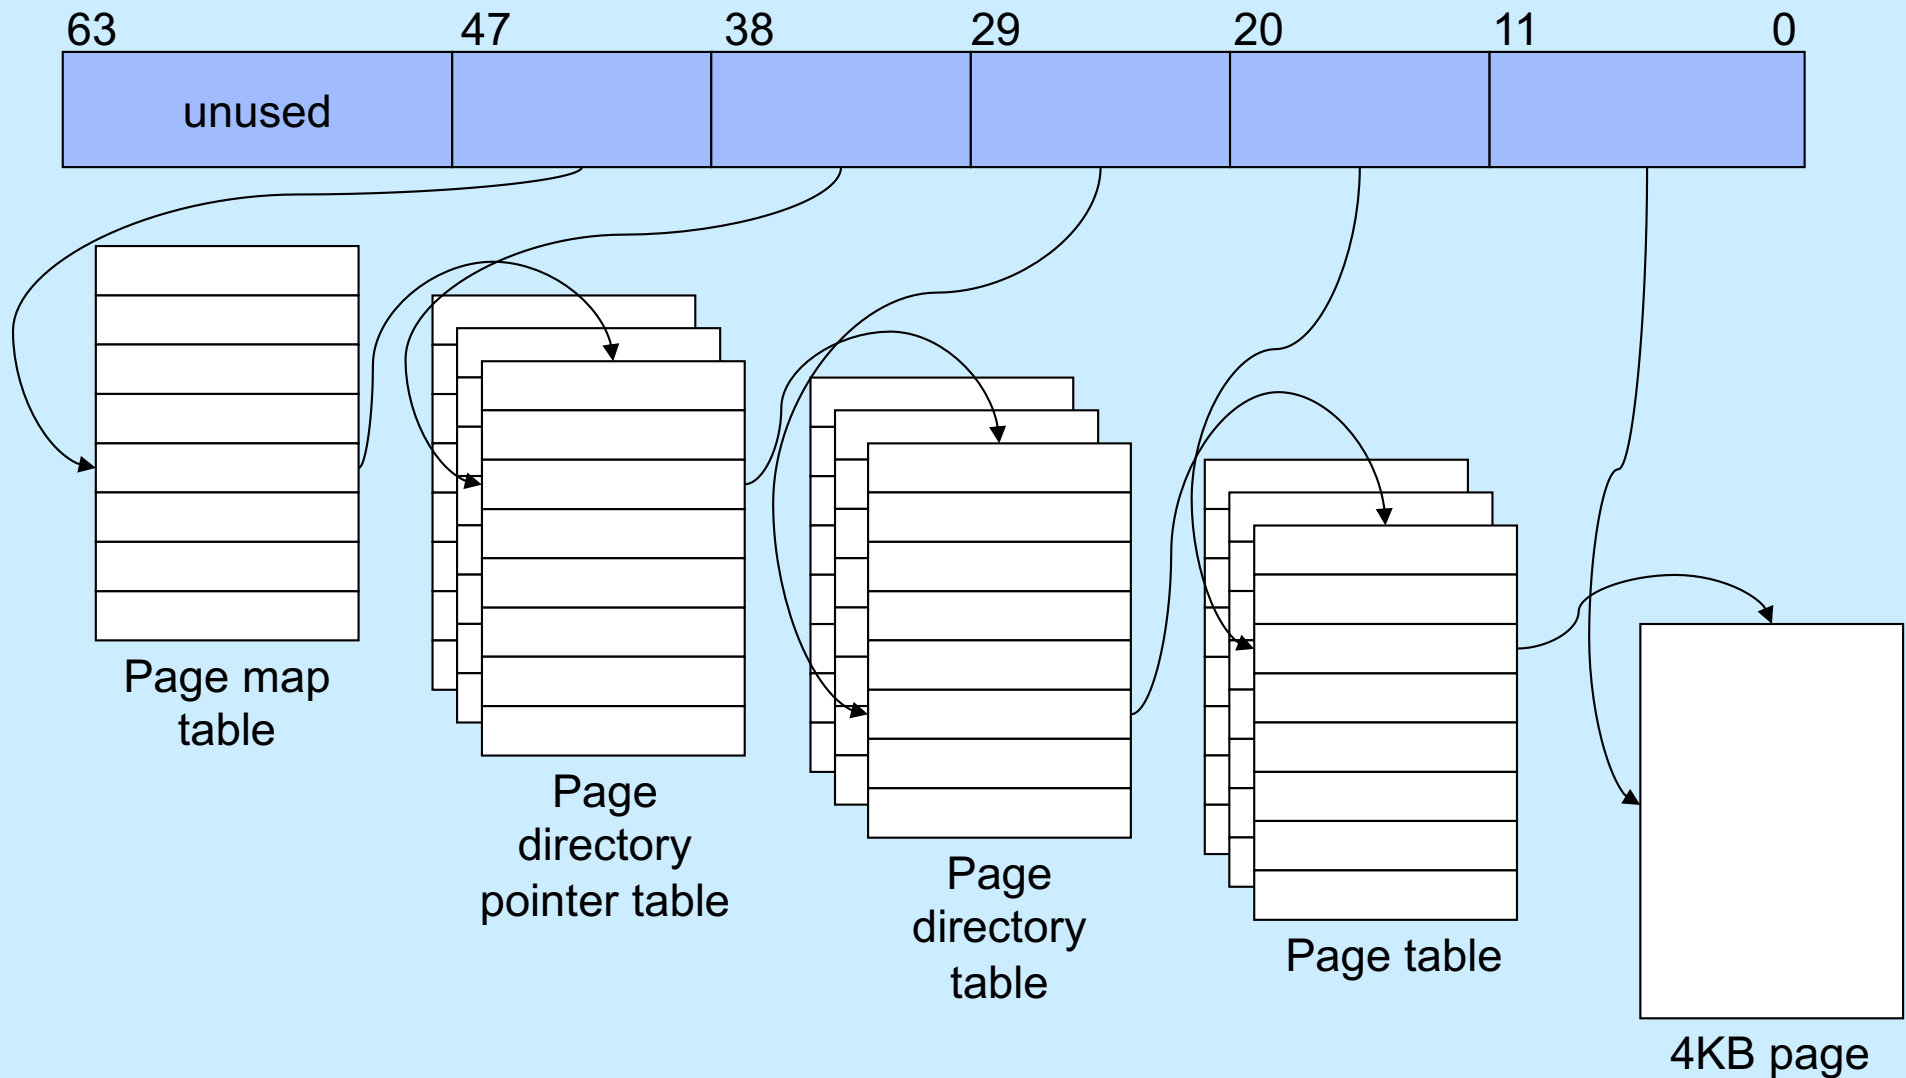
**Can a page start at a virtual address that's not divisible by the page size?**

- a) yes**
- b) no**

# Linux Intel IA32 VM Layout

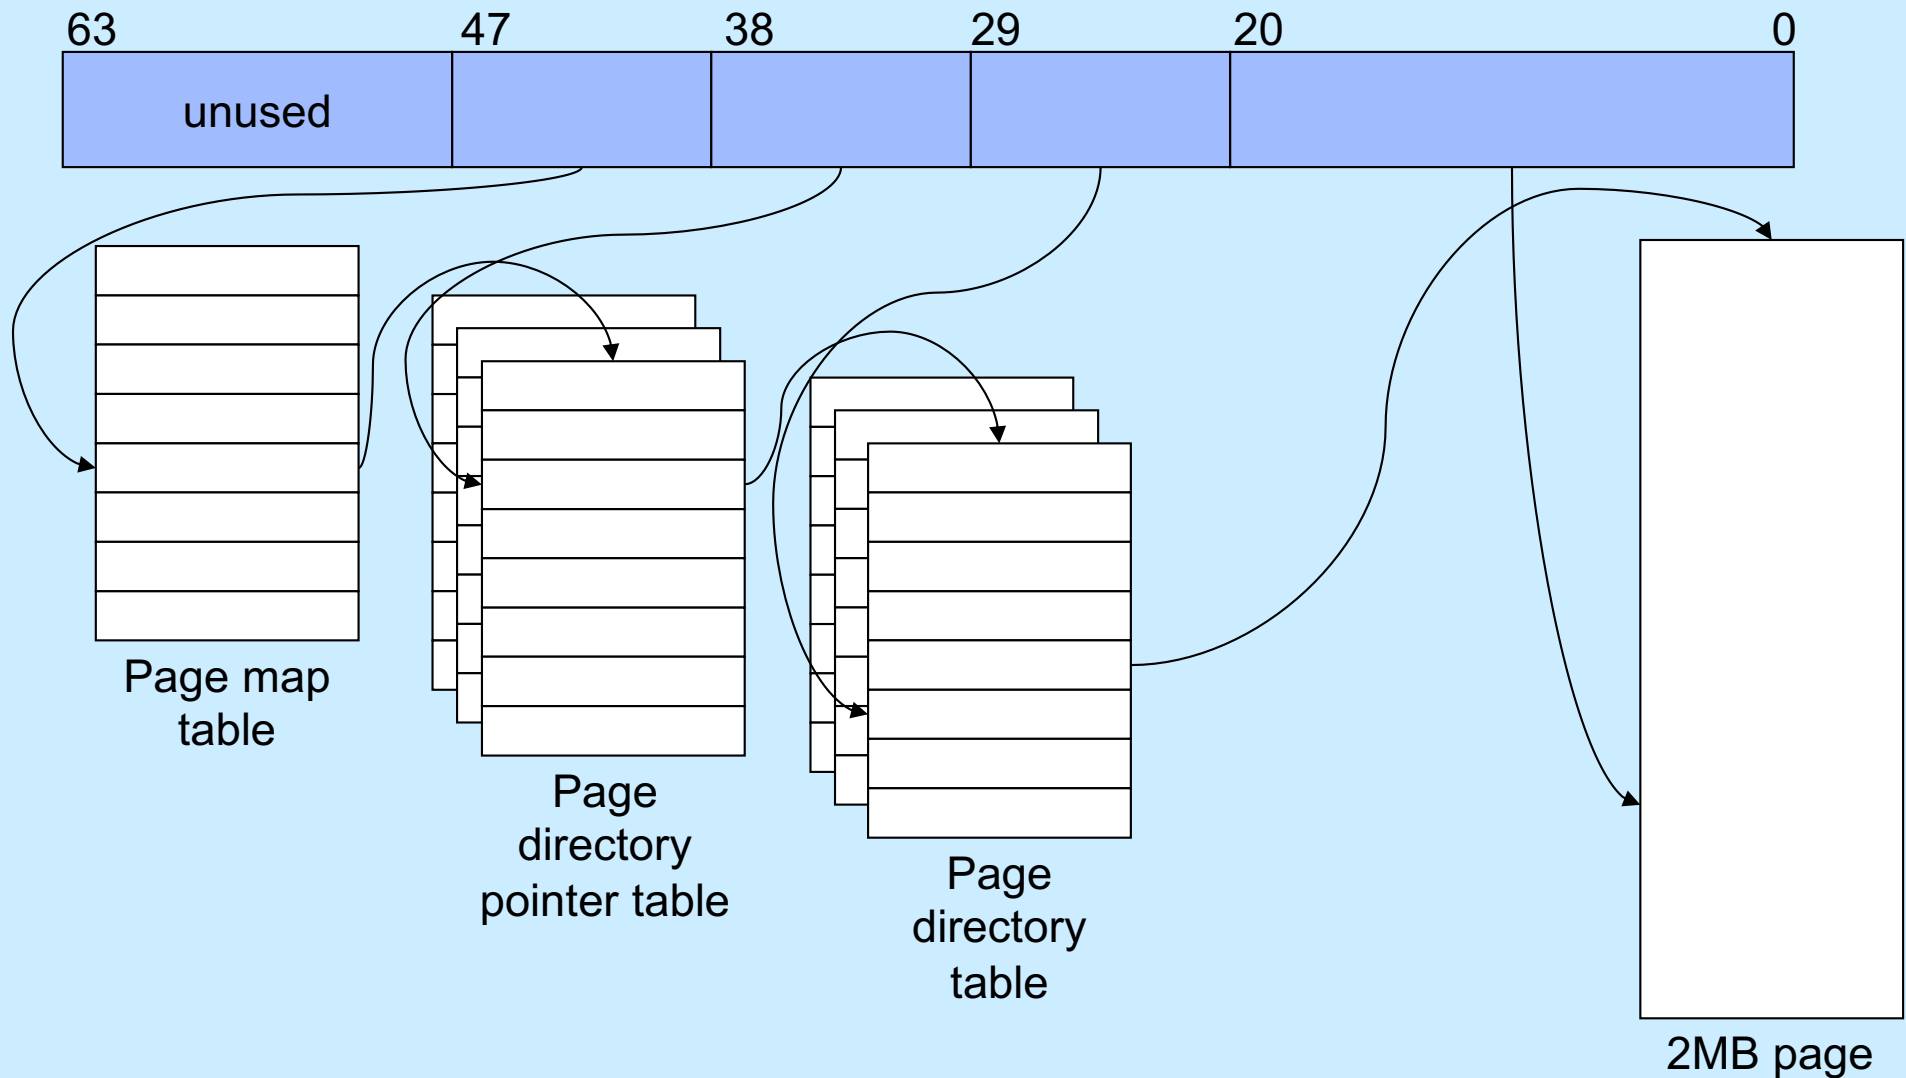


# x86-64 Virtual Address Format 1

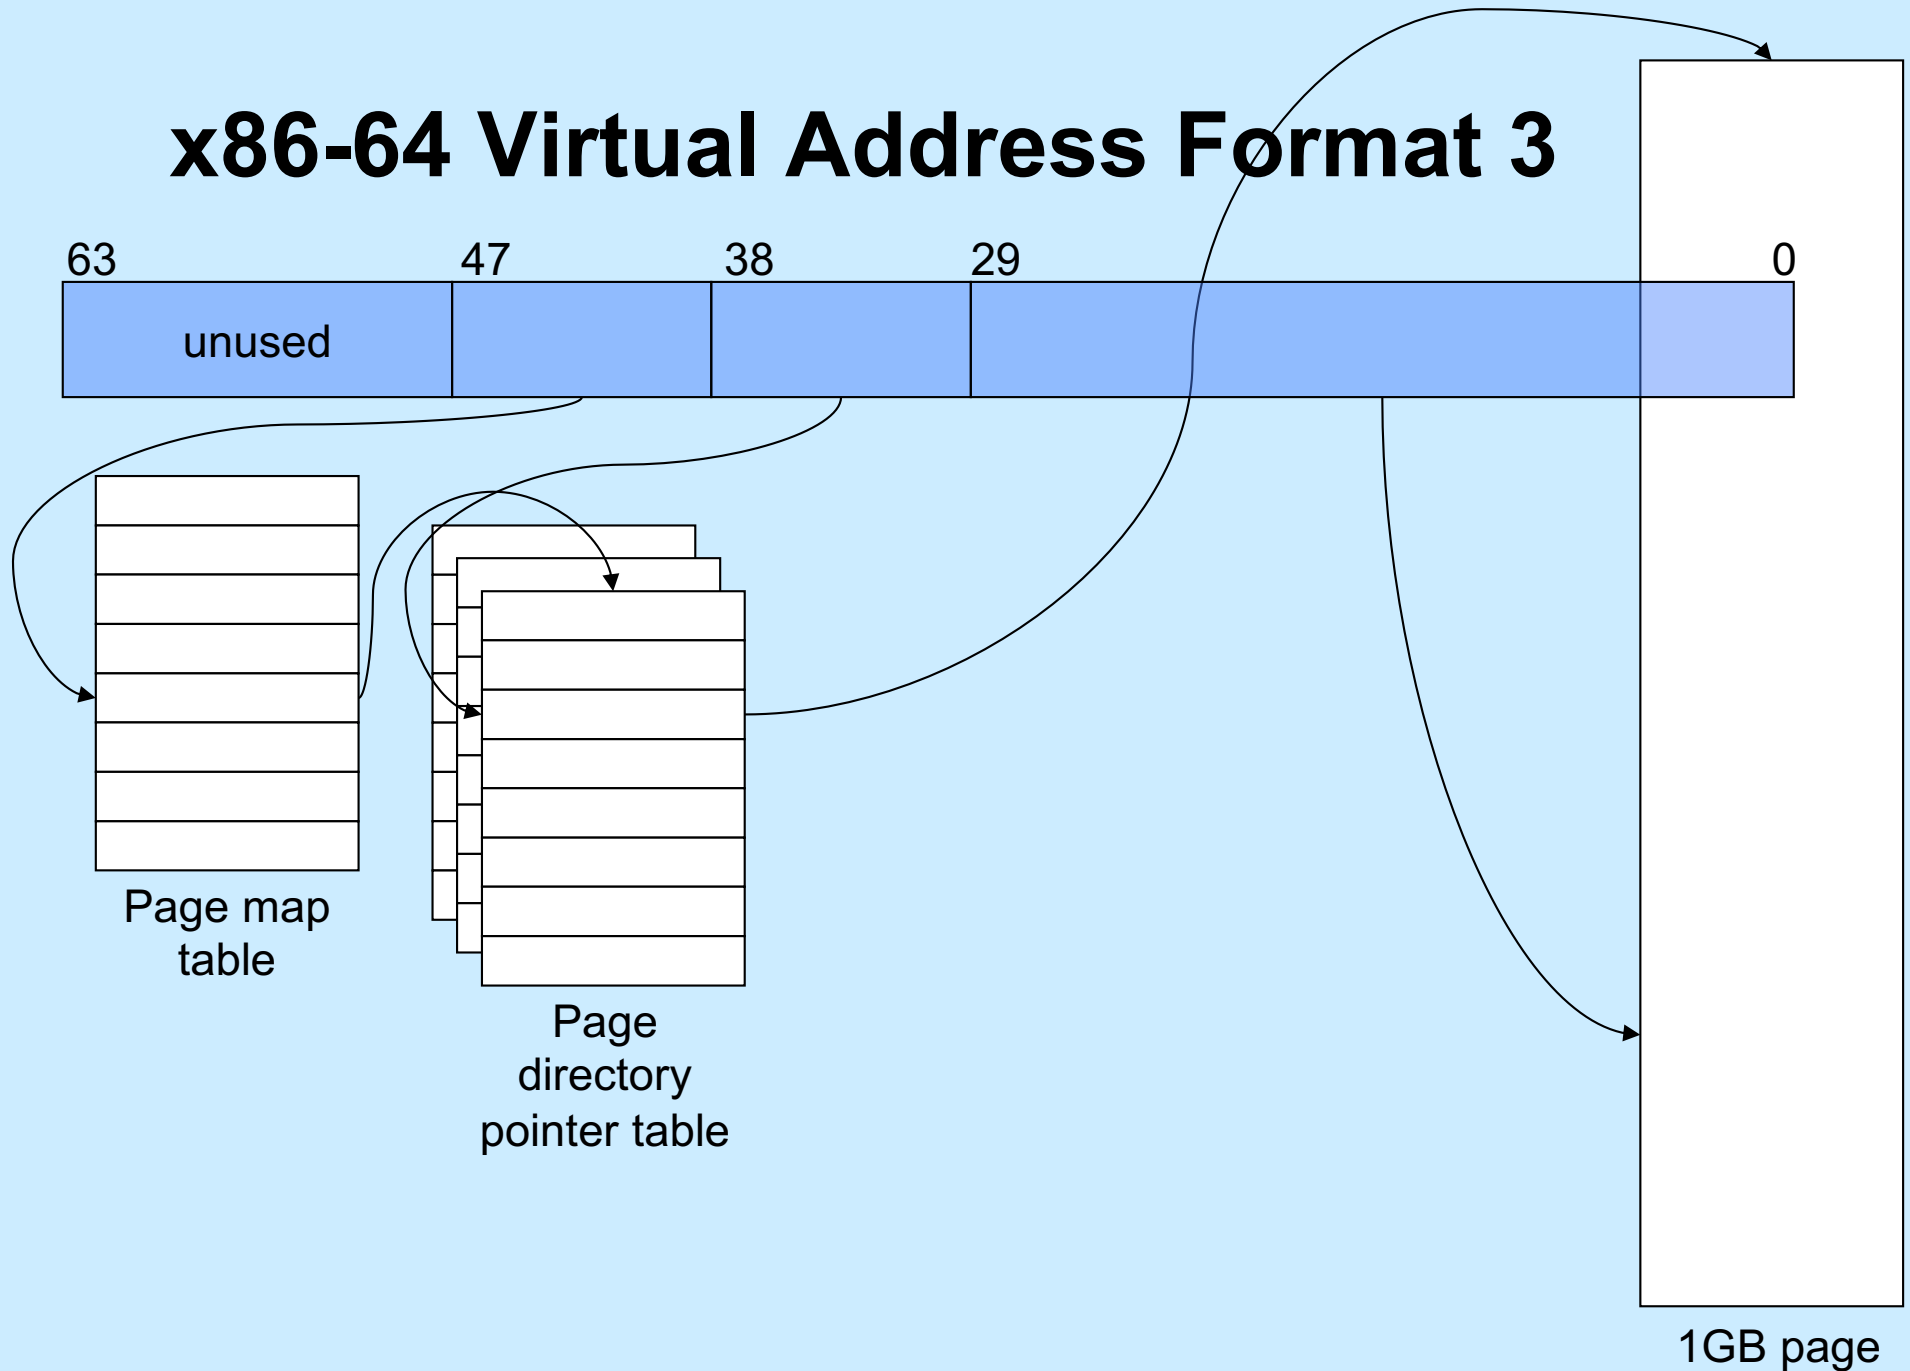




# x86-64 Virtual Address Format 2



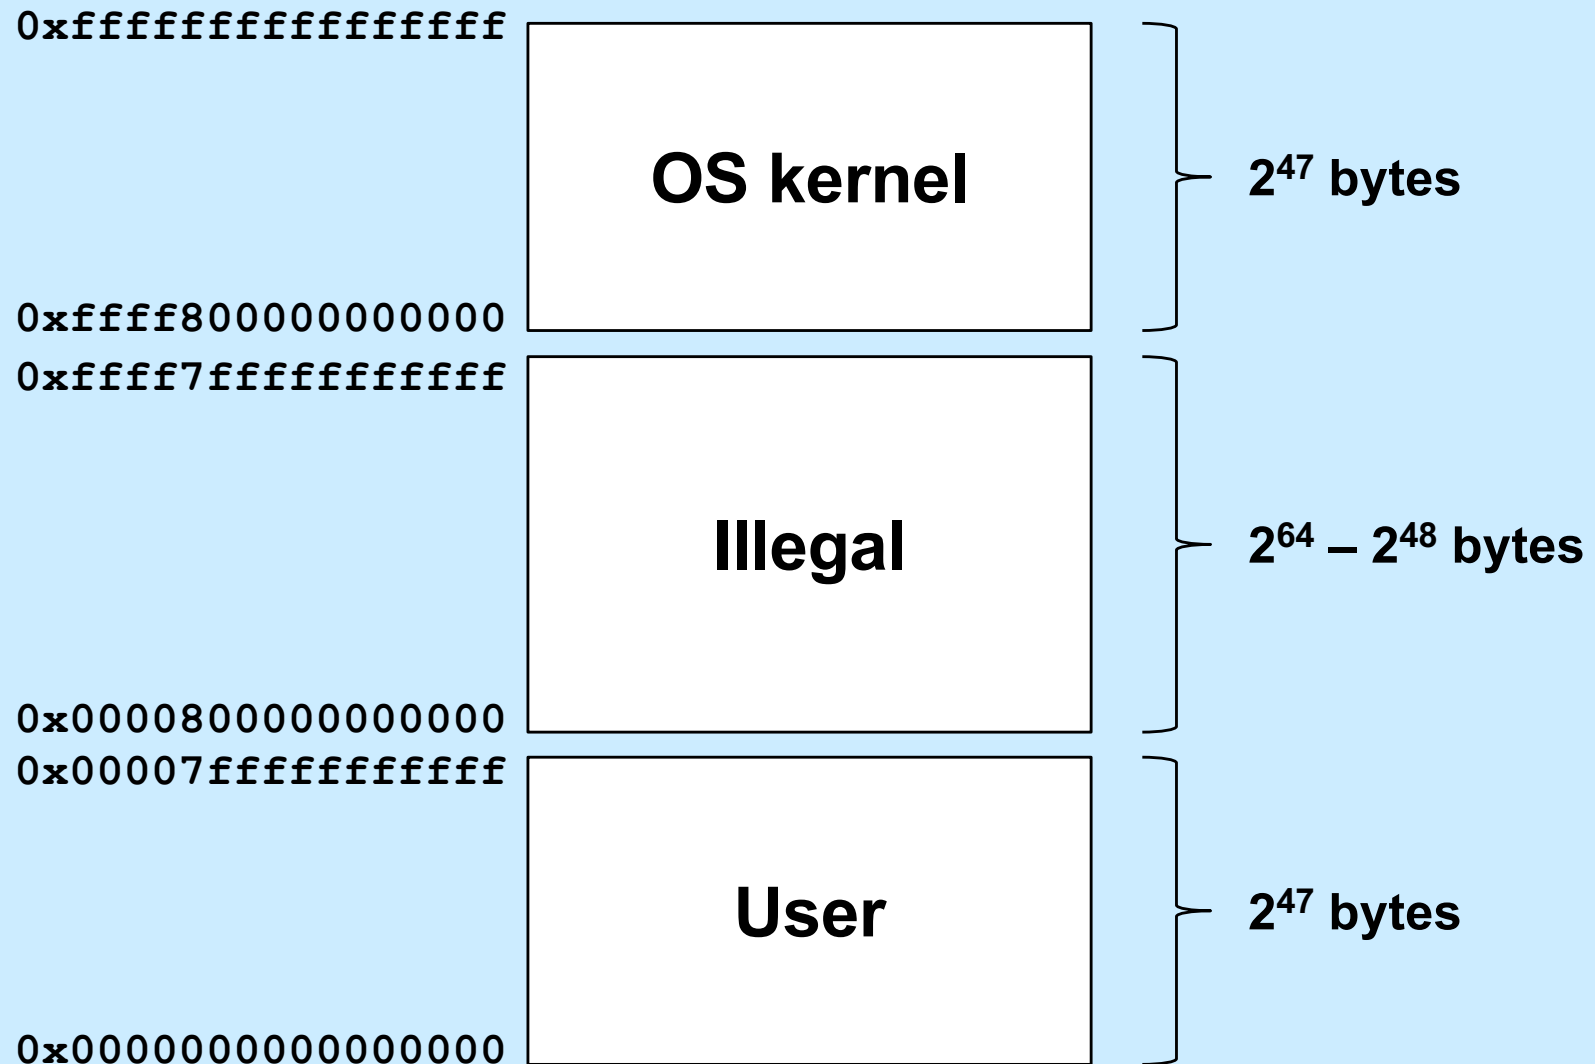
# x86-64 Virtual Address Format 3



# Why Multiple Page Sizes?

- **Fragmentation**
  - for region composed of 4KB pages, average internal fragmentation is 2KB
  - for region composed of 1GB pages, average internal fragmentation is 512MB
- **Page-table overhead**
  - larger page sizes have fewer page tables
    - » less overhead in representing mappings


# x86-64 Address Space



# Performance

- **Page table resides in real memory (DRAM)**
- **A 32-bit virtual-to-real translation requires two accesses to page tables, plus the access to the ultimate real address**
  - three real accesses for each virtual access
  - 3X slowdown!
- **A 64-bit virtual-to-real translation requires four accesses to page tables, plus the access to the ultimate real address**
  - 5X slowdown!

# Translation Lookaside Buffers



# Quiz 3

**Recall that there is a 5x slowdown on memory references via virtual memory on the x86-64. If all references are translated via the TLB, the slowdown will be**

- a) 1x**
- b) 2x**
- c) 3x**
- d) 4x**

# OS Role in Virtual Memory


- **Memory is like a cache**
  - quick access if what's wanted is mapped via page table
  - slow if not — OS assistance required
- **OS**
  - make sure what's needed is mapped in
  - make sure what's no longer needed is not mapped in




# Mechanism

- **Program references memory**
  - if reference is mapped, access is quick
    - » even quicker if translation in TLB and referent in on-chip cache
  - if not, page-translation fault occurs and OS is invoked
    - » determines desired page
    - » maps it in, if legal reference


# The “Pageout Daemon”



# Managing Page Frames

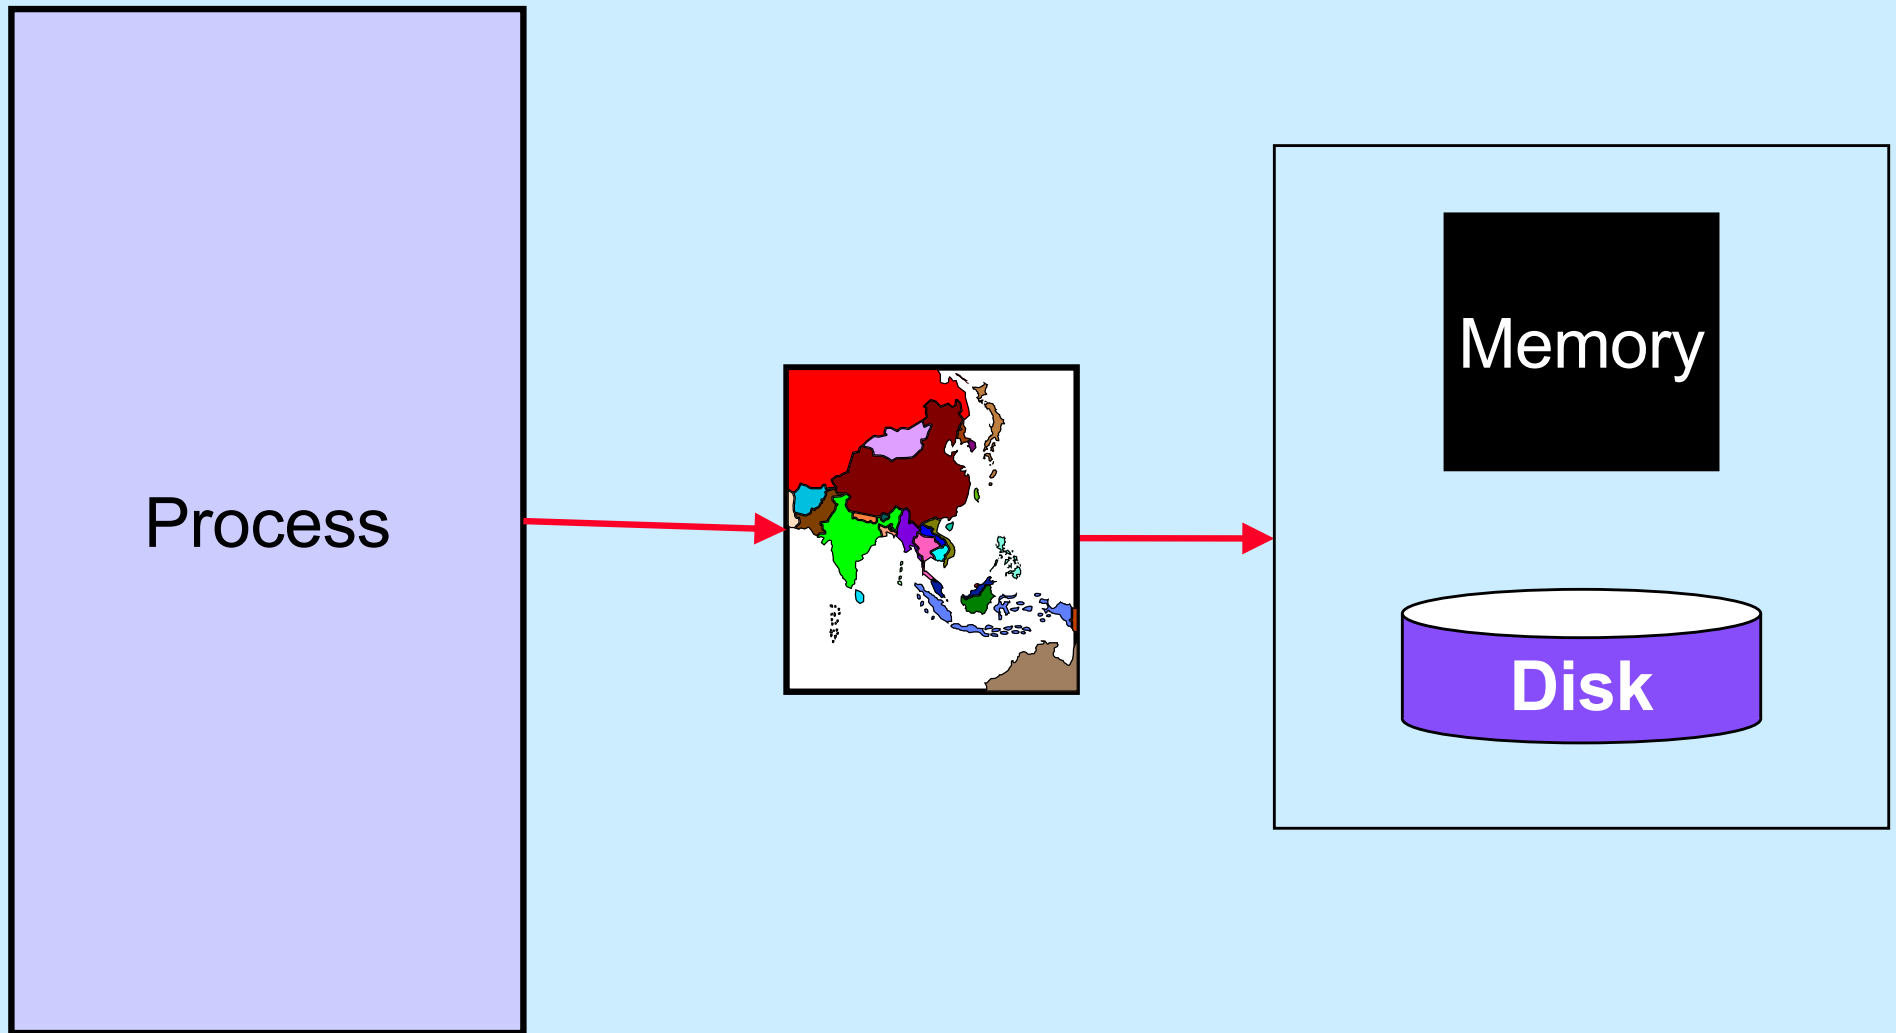


# Clock Algorithm

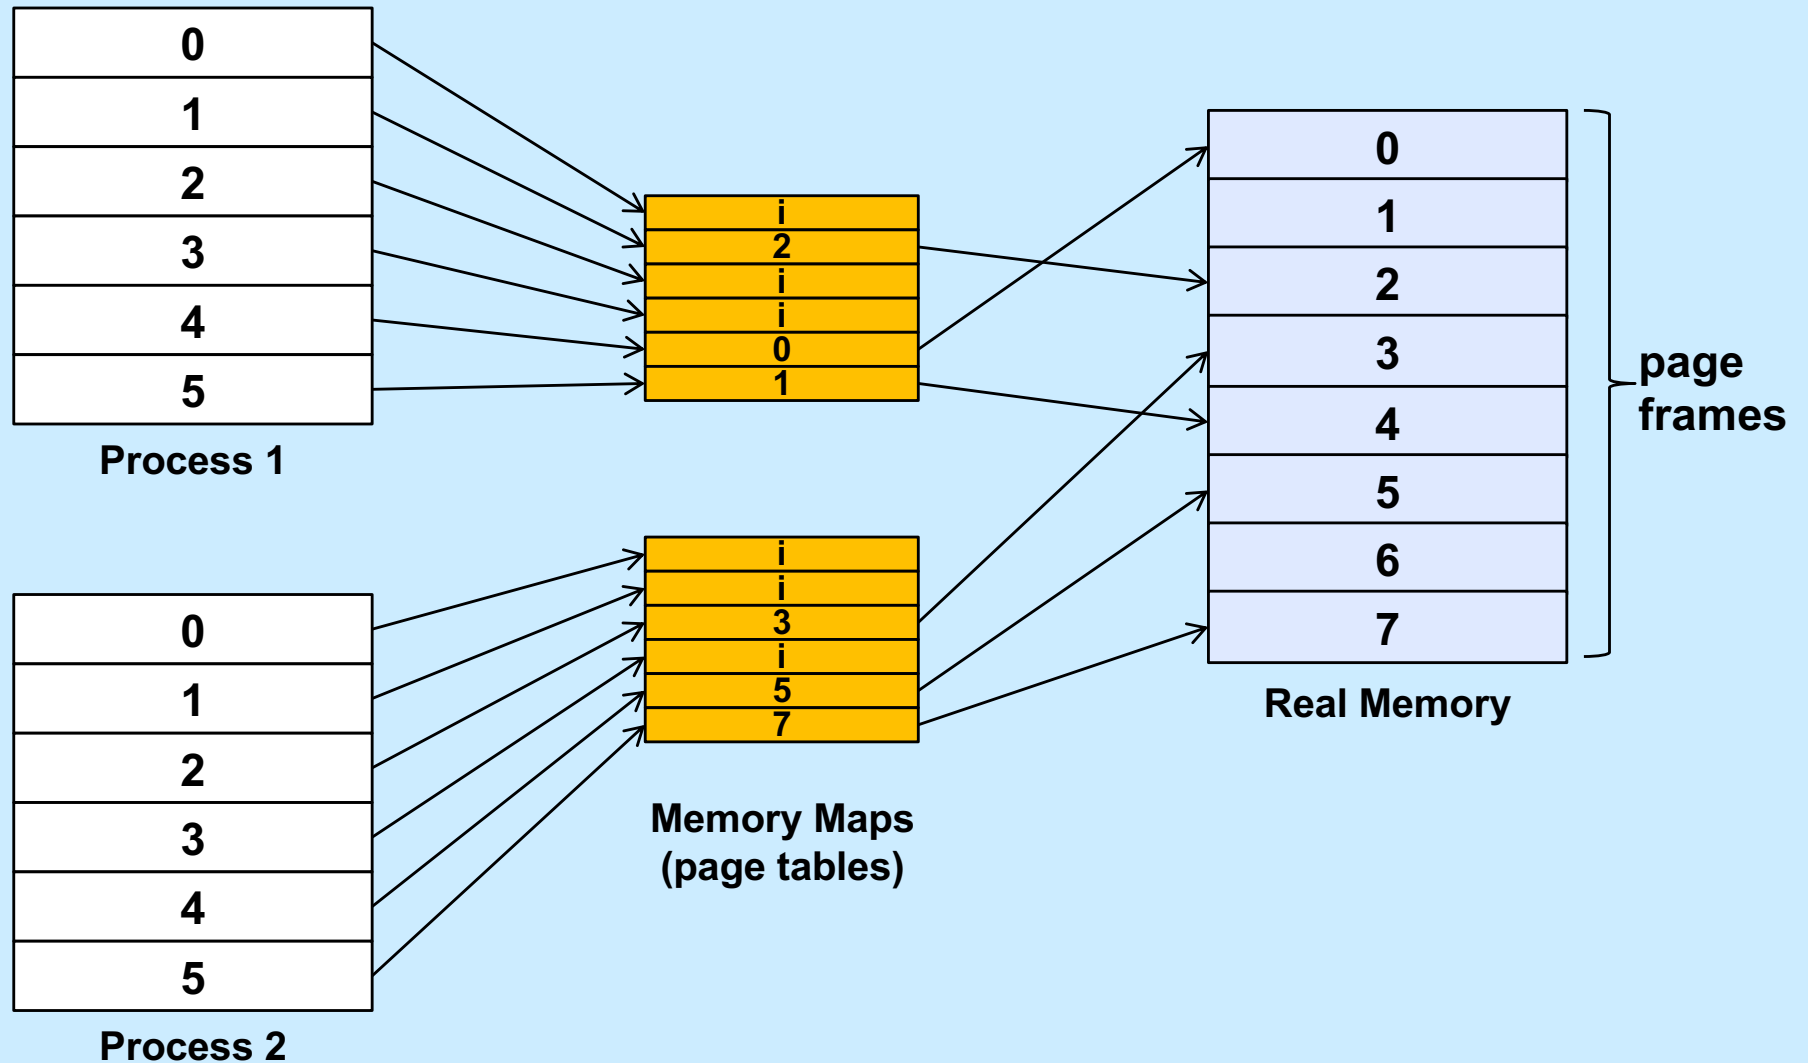


# Why is virtual memory used?

# More VM than RM

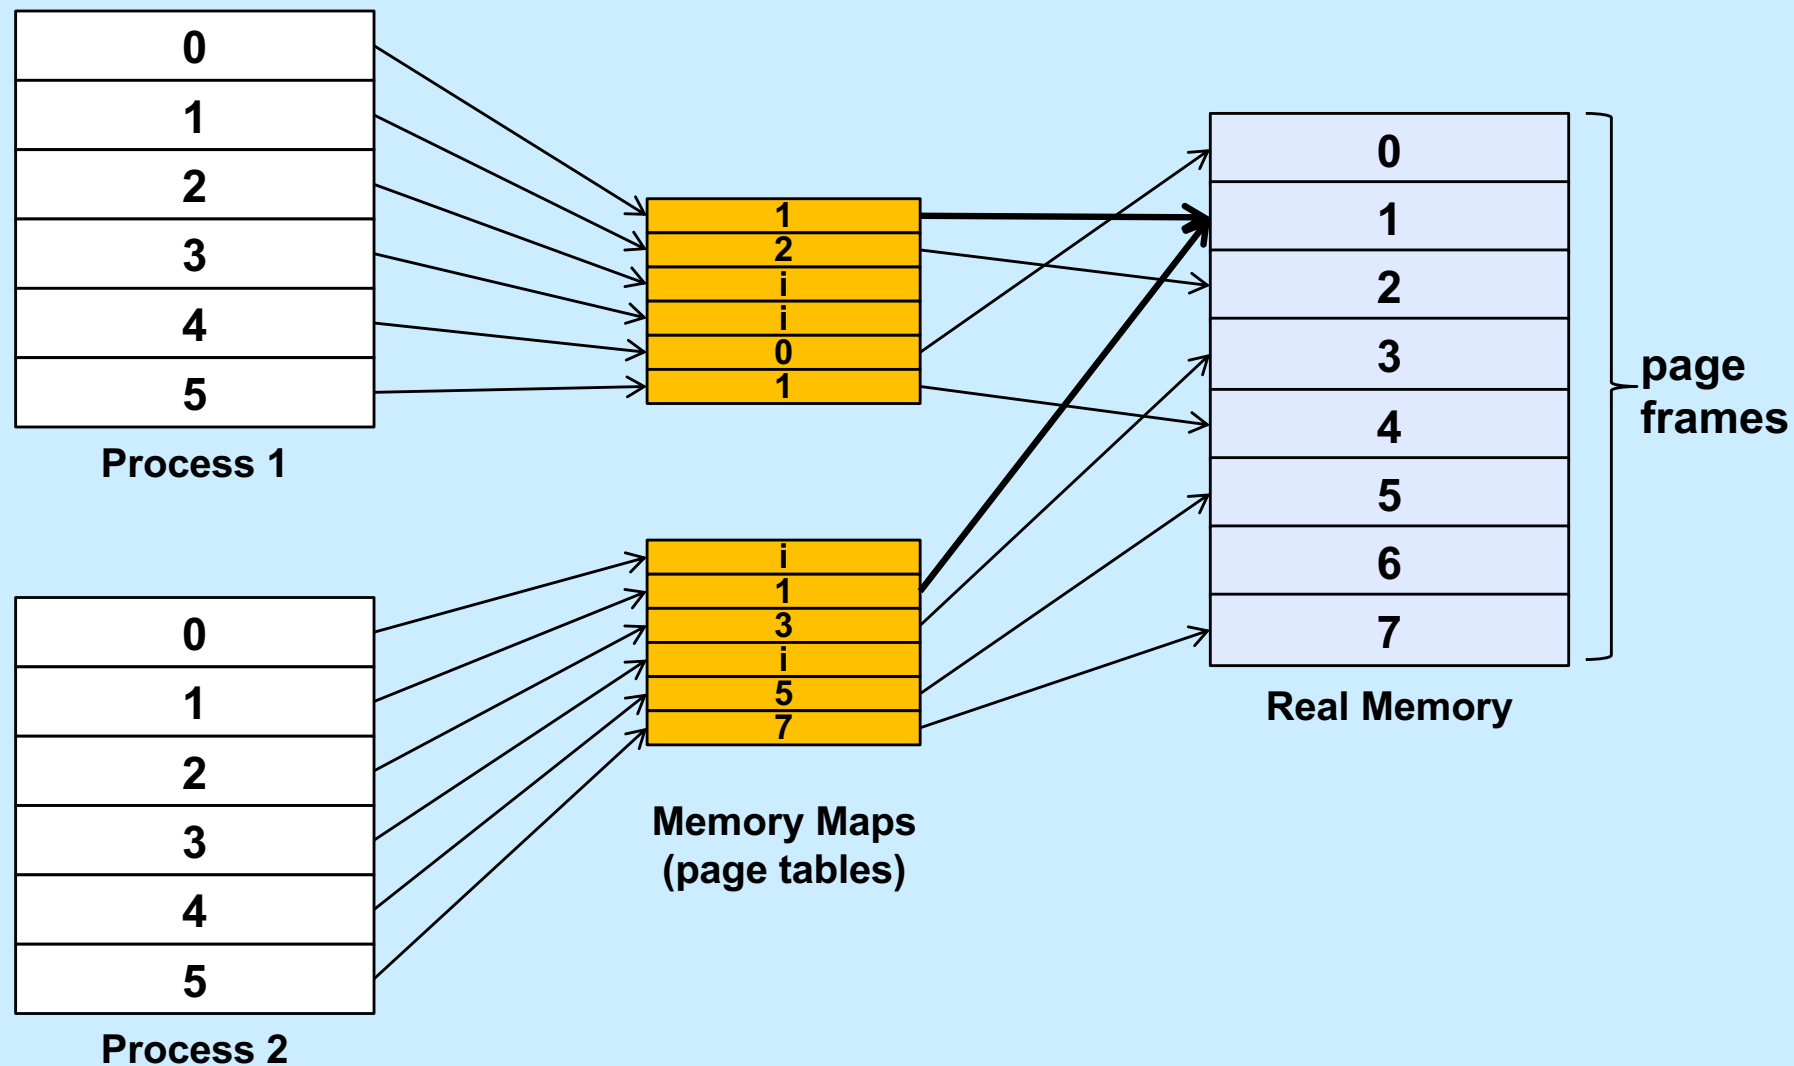


# Isolation



**Virtual Memory**

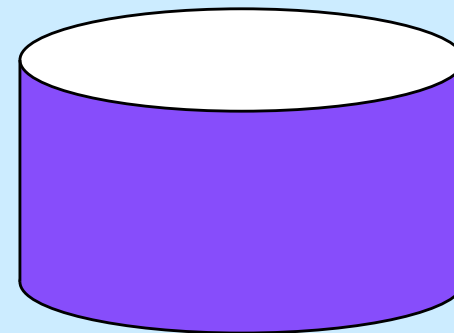

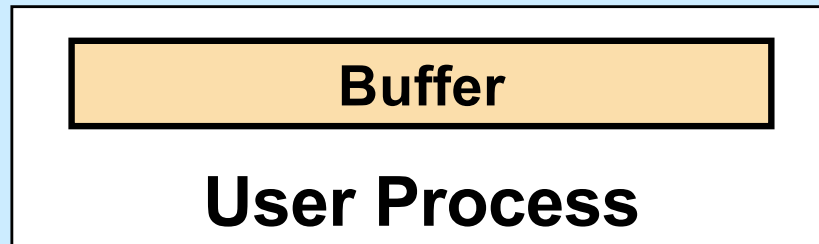
# Sharing



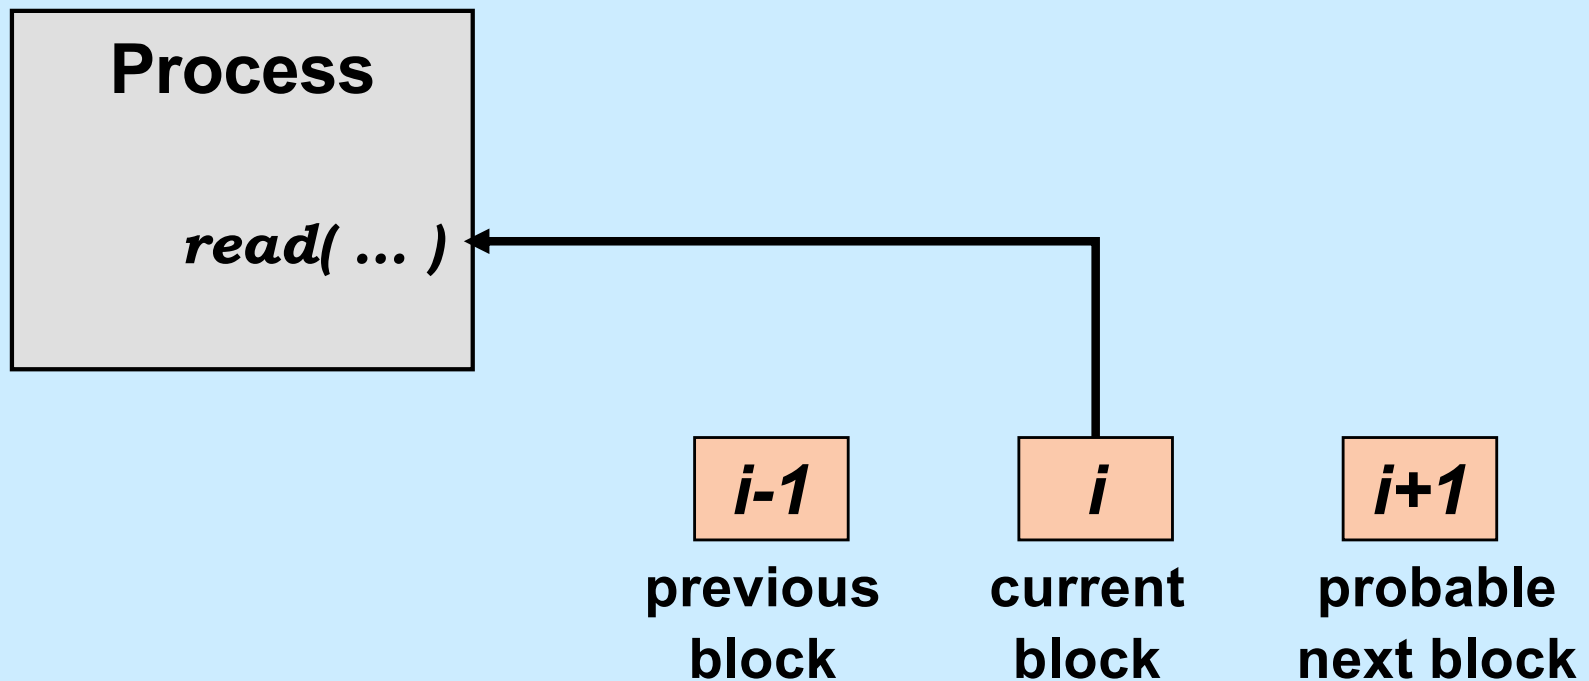
**Virtual Memory**



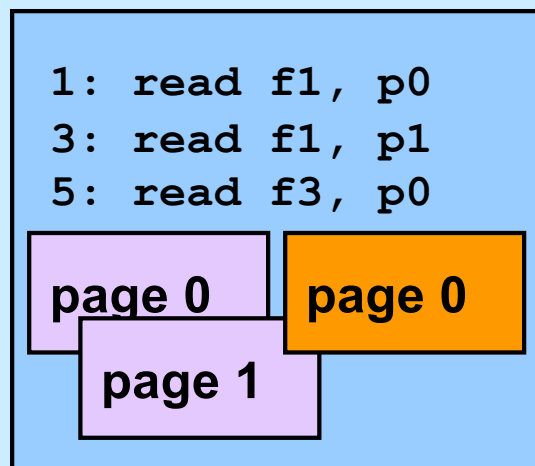
# File I/O



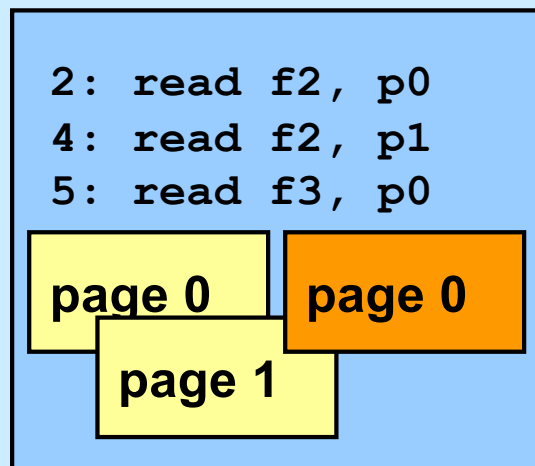
# Multi-Buffered I/O



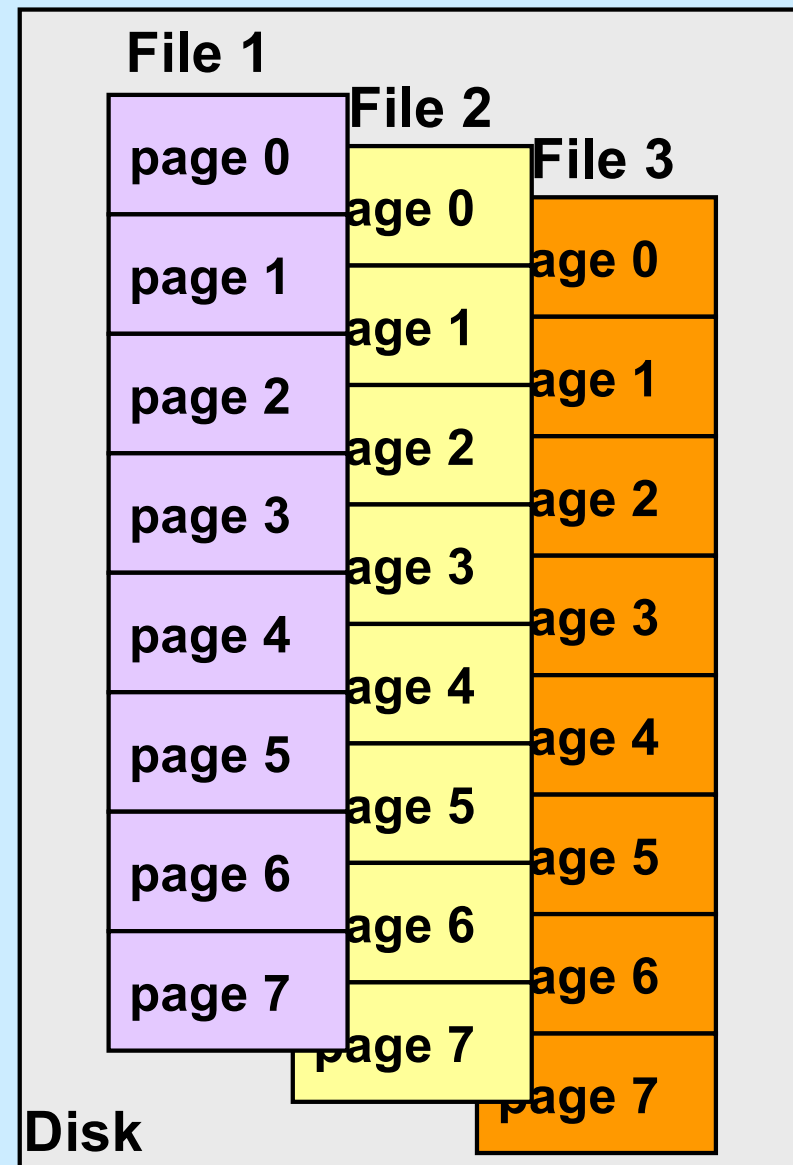

# Traditional I/O



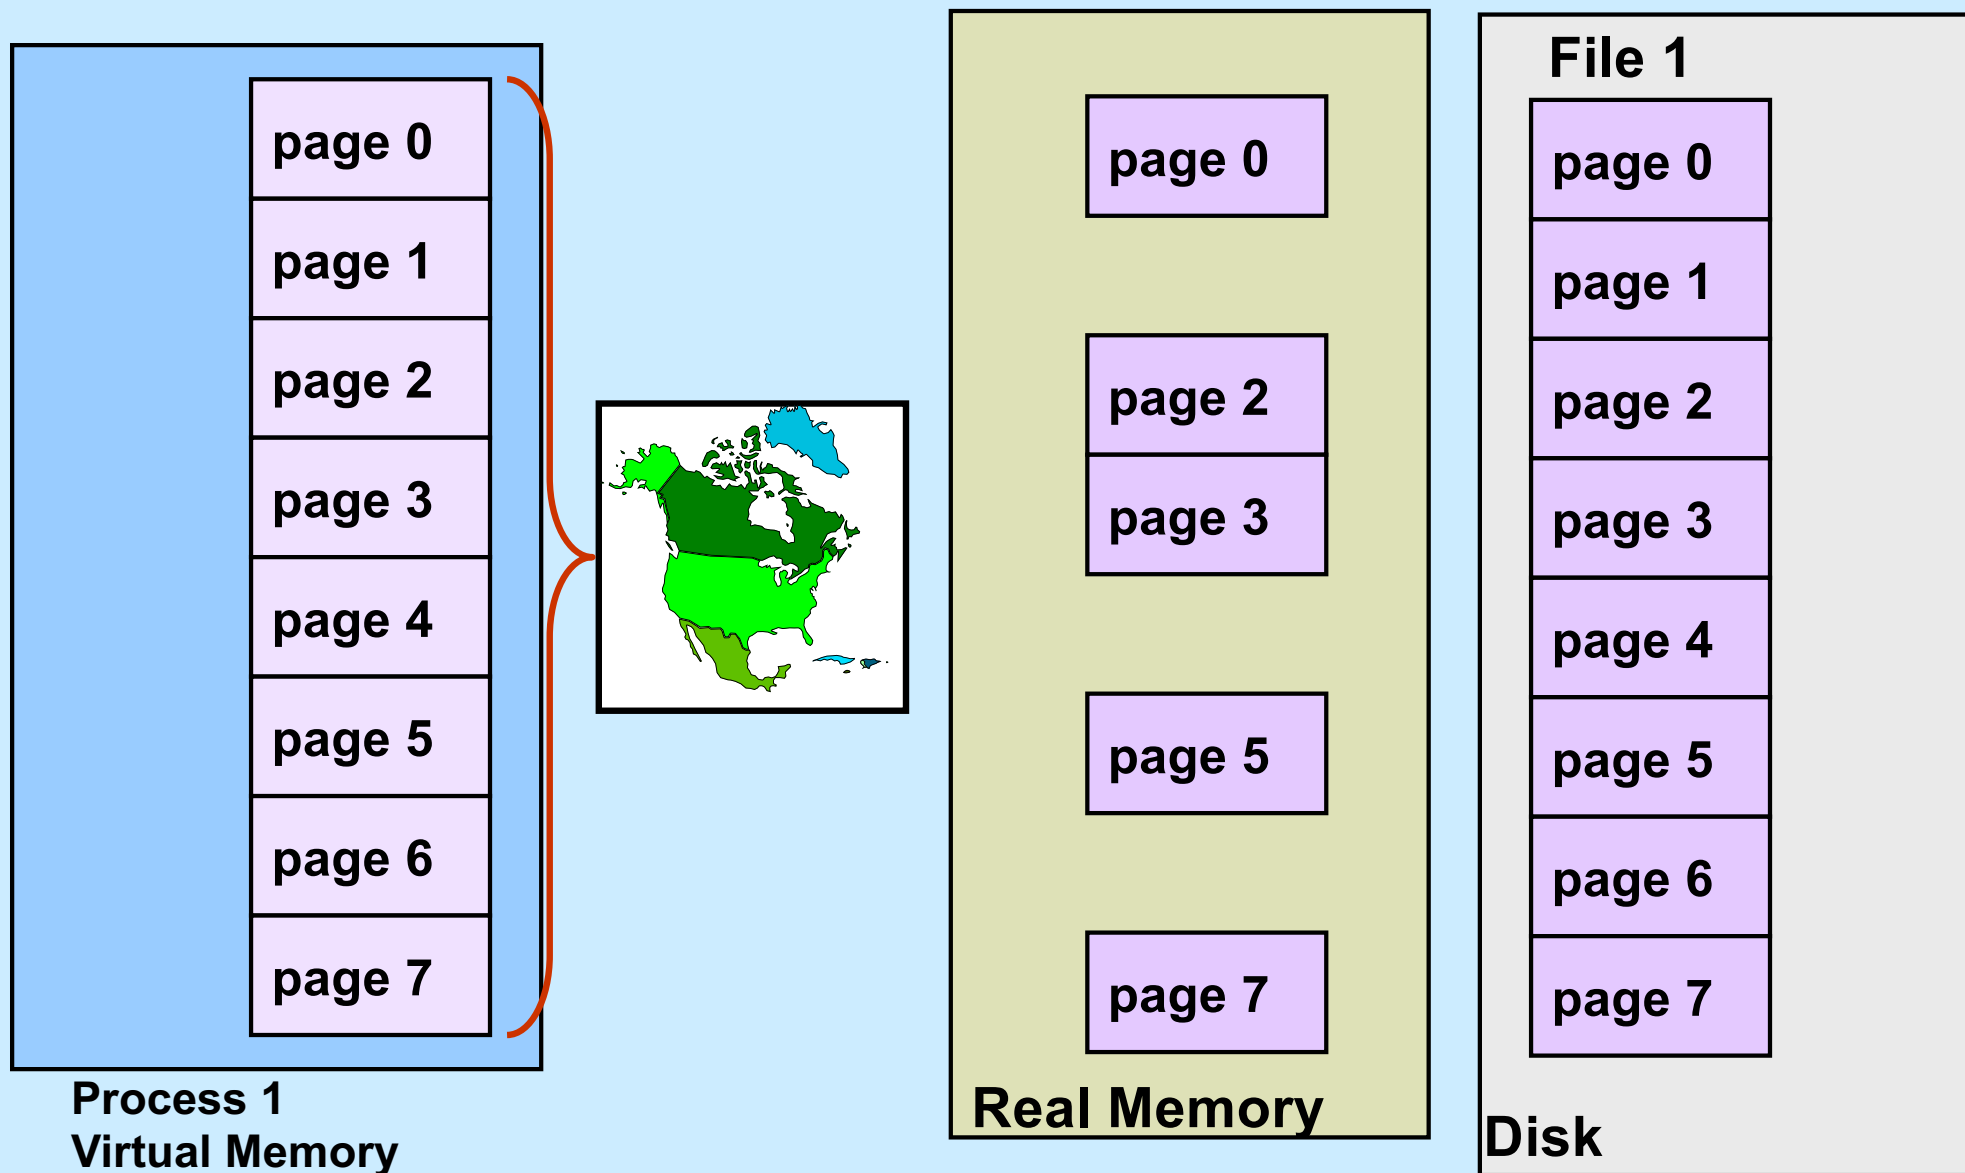
User Process 1



User Process 2



# Mapped File I/O



Process 1  
Virtual Memory

Real Memory

Disk

# Multi-Process Mapped File I/O

