

# CS 33

## Multithreaded Programming II

# Example (1)

```
#include <stdio.h>
#include <pthread.h>
#include <string.h>
```

```
#define M    3
#define N    4
#define P    5
```

```
long A[M][N];
long B[N][P];
long C[M][P];
```

```
void *matmult(void *);
```

```
main( ) {
    long i;
    pthread_t thr[M];
    int error;

    // initialize the matrices
    ...
}
```

## Example (2)

```
for (i=0; i<M; i++) { // create worker threads
    if (error = pthread_create(
        &thr[i],
        0,
        matmult,
        (void *)i)) {
        fprintf(stderr, "pthread_create: %s", strerror(error));
        exit(1);
    }
}

for (i=0; i<M; i++) // wait for workers to finish their jobs
    pthread_join(thr[i], 0)

/* print the results ... */
}
```

## Example (3)

```
void *matmult(void *arg) {  
    long row = (long) arg;  
    long col;  
    long i;  
    long t;  
  
    for (col=0; col < P; col++) {  
        t = 0;  
        for (i=0; i<N; i++)  
            t += A[row][i] * B[i][col];  
        C[row][col] = t;  
    }  
    return (0);  
}
```

# Compiling It

```
% gcc -o mat mat.c -pthread
```

# Termination

```
pthread_exit((void *) value);
```

```
return((void *) value);
```

```
pthread_join(thread, (void **) &value);
```

# Detached Threads

```
start_servers( ) {  
    pthread_t thread;  
    int i;  
  
    for (i=0; i<nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
void *server(void * arg) {  
    ...  
}
```

# Worker Threads

```
int main() {  
    pthread_t thread[10];  
    for (int i=0; i<10; i++)  
        pthread_create(&thread[i], 0,  
            worker, (void *)i);  
    return 0;  
}
```

```
void *worker(...) {...}
```



# Termination

```
pthread_exit((void *) value);
```

```
return((void *) value);
```

```
pthread_join(thread, (void **) &value);
```

```
exit(code); // terminates process!
```

# Complications

```
void relay(int left, int right) {  
    pthread_t LRthread, RLthread;  
  
    pthread_create(&LRthread,  
        0,  
        copy,  
        left, right);           // Can't do this ...  
    pthread_create(&RLthread,  
        0,  
        copy,  
        right, left);           // Can't do this ...  
}
```

# Multiple Arguments

```
typedef struct args {  
    int src;  
    int dest;  
} args_t;
```

```
void relay(int left, int right) {  
    args_t LRargs, RLargs;  
    pthread_t LRthread, RLthread;  
    ...  
    pthread_create(&LRthread, 0, copy, &LRargs);  
    pthread_create(&RLthread, 0, copy, &RLargs);  
}
```

# Multiple Arguments

```
typedef struct args {  
    int src;  
    int dest;  
} args_t;
```


```
void relay(int left, int right) {  
    args_t LRargs, RLargs;  
    pthread_t LRthread, RLthread;  
    ...  
    pthread_create(&LRthread, 0, copy, &LRargs);  
    pthread_create(&RLthread, 0, copy, &RLargs);  
}
```

## Quiz 1

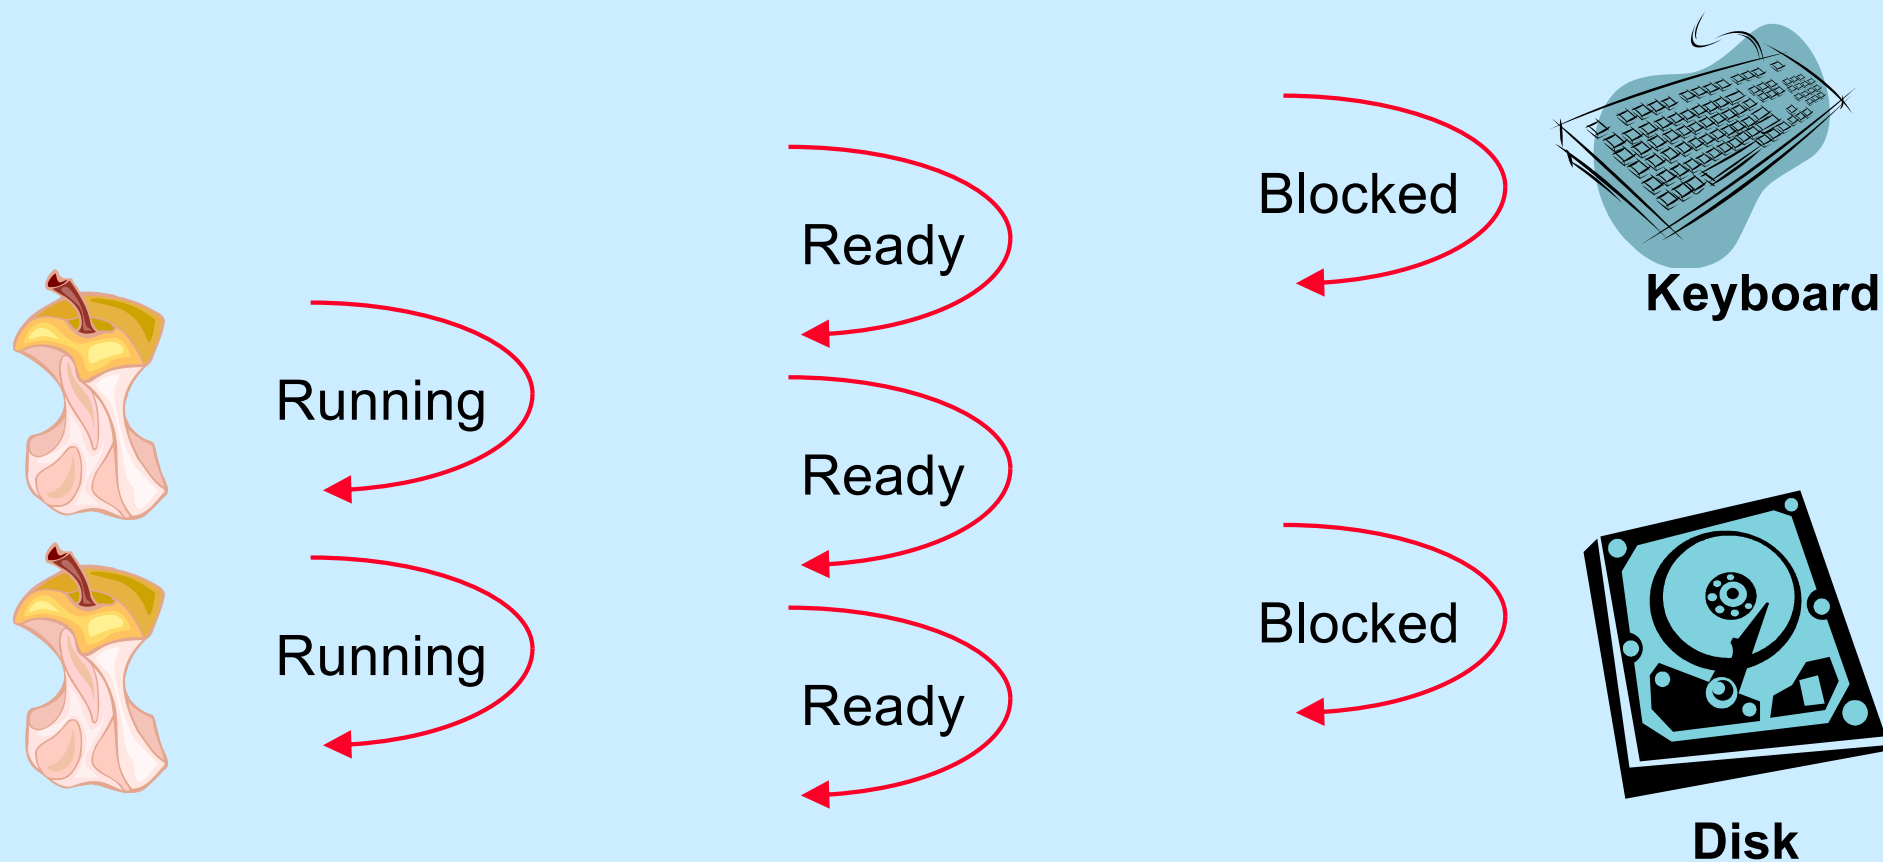
**Does this work?**

- a) yes**
- b) no**

# Execution



# Multiplexing Processors



# Quiz 2

```
pthread_create(&tid, 0, tproc, (void *)1);  
pthread_create(&tid, 0, tproc, (void *)2);
```

```
printf("T0\n");
```

```
...
```

```
void *tproc(void *arg) {  
    printf("T%d\n", (long) arg);  
    return 0;  
}
```

**In which order are things printed?**

- a) T0, T1, T2
- b) T1, T2, T0
- c) T2, T1, T0
- d) indeterminate

# Cost of Threads

```
int main(int argc, char *argv[]) {
    ...
    val = niters/nthreads;

    for (i=0; i<nthreads; i++)
        pthread_create(&thread, 0, work, (void *)val);
    pthread_exit(0);
    return 0;
}

void *work(void *arg) {
    long n = (long)arg; int i, j; volatile long x;

    for (i=0; i<n; i++) {
        x = 0;
        for (j=0; j<1000; j++)
            x = x*j;
    }
    return 0;
}
```



# Cost of Threads

```
int main(int argc, char *argv[]) {
    ...
    val = niters/nthreads;

    for (i=0; i<nthreads; i++)
        pthread_create(&thread, 0, work, (void *)val);
    pthread_exit(0);
    return 0;
}

void *work(void *arg) {
    long n = (long)arg; int i, j; volatile long x;

    for (i=0; i<n; i++) {
        x = 0;
        for (j=0; j<1000; j++)
            x = x*j;
    }
    return (void *)x;
}
```

## Not a Quiz

This code runs in time  $n$  on a 4-core processor when  $nthreads$  is 8. It runs in time  $p$  on the same processor when  $nthreads$  is 400.

- a)  $n \ll p$  (slower)
- b)  $n \approx p$  (same speed)
- c)  $n \gg p$  (faster)

# Problem

```
pthread_create(&thread, 0, start, 0);
```

```
...
```

```
void *start(void *arg) {  
    long BigArray[128*1024*1024];  
    ...  
    return 0;  
}
```

# Thread Attributes

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
  
...  
  
/* establish some attributes */  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...
```

# Stack Size

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 130*1024*1024);  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...
```

# Mutual Exclusion



# Threads and Mutual Exclusion

## Thread 1:

```
x = x+1;  
/*  
    movl x,%eax  
    incr %eax  
    movl %eax,x  
*/
```

## Thread 2:

```
x = x+1;  
/*  
    movl x,%eax  
    incr %eax  
    movl %eax,x  
*/
```

# Quiz 3

Suppose gcc produces the following code. Will it still be the case that x's value might not be incremented by 2?

- a) yes
- b) no

**Thread 1:**

```
x = x+1;  
/*  
    incr x  
*/
```

**Thread 2:**

```
x = x+1;  
/*  
    incr x  
*/
```


# POSIX Threads Mutual Exclusion

```
pthread_mutex_t m =  
    PTHREAD_MUTEX_INITIALIZER;  
    // shared by both threads  
int x; // ditto  
  
pthread_mutex_lock(&m);  
  
x = x+1;  
  
pthread_mutex_unlock(&m);
```



# A Queue

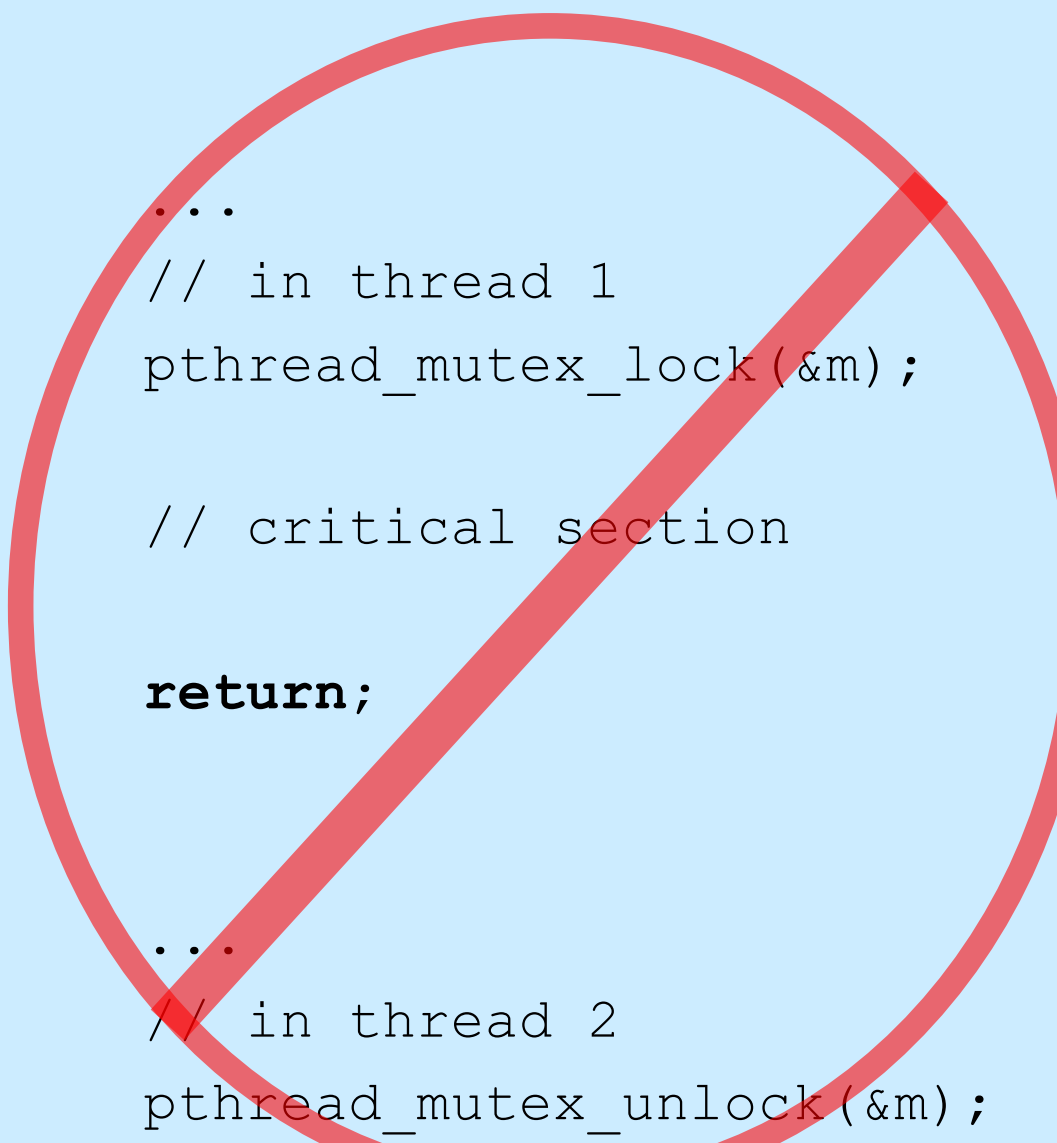
```
void enqueue(node_t *item) {  
    pthread_mutex_lock(&mutex);  
    item->next = NULL;  
    if (tail == NULL) {  
        head = item;  
    } else {  
        tail->next = item;  
    }  
    tail = item;  
    pthread_mutex_unlock(&mutex);  
}
```



```
node_t *dequeue() {  
    node_t *ret;  
    pthread_mutex_lock(&mutex);  
    if (head == NULL) {  
        ret = NULL;  
    } else {  
        ret = head;  
        head = head->next;  
        if (head == NULL)  
            tail = NULL;  
    }  
    pthread_mutex_unlock(&mutex);  
    return ret;  
}
```

# Correct Usage

```
pthread_mutex_lock(&m);  
  
// critical section  
  
pthread_mutex_unlock(&m);
```




```
...  
// in thread 1  
pthread_mutex_lock(&m);  
  
// critical section  
  
return;  
  
...  
// in thread 2  
pthread_mutex_unlock(&m);
```

# Taking Multiple Locks

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

```
proc2( ) {  
    pthread_mutex_lock(&m2);  
    /* use object 2 */  
    pthread_mutex_lock(&m1);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m1);  
    pthread_mutex_unlock(&m2);  
}
```

# Preventing Deadlock



# Taking Multiple Locks, Safely

```
proc1( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

```
proc2( ) {  
    pthread_mutex_lock(&m1);  
    /* use object 1 */  
    pthread_mutex_lock(&m2);  
    /* use objects 1 and 2 */  
    pthread_mutex_unlock(&m2);  
    pthread_mutex_unlock(&m1);  
}
```

# Practical Issues with Mutexes

- **Used a lot in multithreaded programs**
  - **speed is really important**
    - » **shouldn't slow things down much in the success case**
  - **checking for errors slows things down (a lot)**
    - » **thus errors aren't checked by default**

# Set Up

```
int pthread_mutex_init(pthread_mutex_t *mutexp,  
    pthread_mutexattr_t *attrp)
```

```
int pthread_mutex_destroy(pthread_mutex_t *mutexp)
```

```
int pthread_mutexattr_init(pthread_mutexattr_t *attrp)
```

```
int pthread_mutexattr_destroy(pthread_mutexattr_t *attrp)
```

# Stupid (i.e., Common) Mistakes ...

```
pthread_mutex_lock(&m1);  
pthread_mutex_lock(&m1);  
    // really meant to lock m2 ...
```

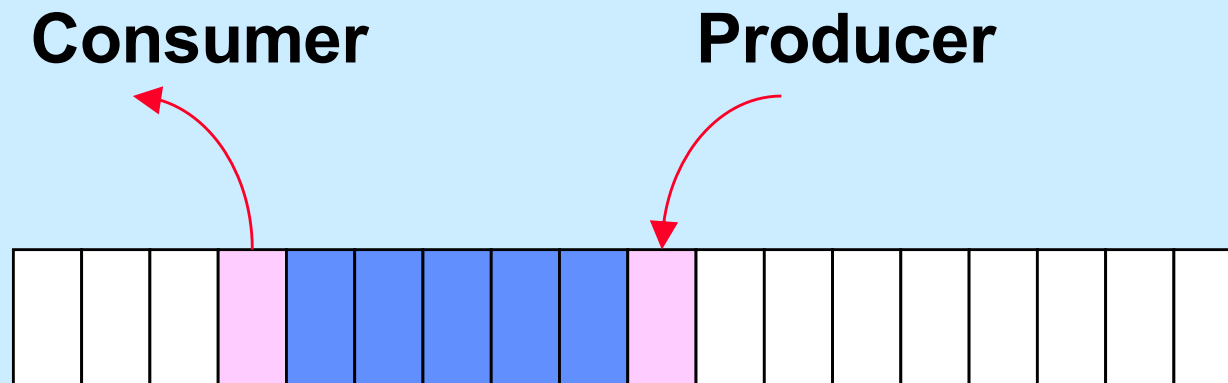
```
pthread_mutex_lock(&m1);  
    ...  
pthread_mutex_unlock(&m2);  
    // really meant to unlock m1 ...
```



# Runtime Error Checking

```
pthread_mutexattr_t err_chk_attr;  
pthread_mutexattr_init(&err_chk_attr);  
pthread_mutexattr_settype(&err_chk_attr,  
    PTHREAD_MUTEX_ERRORCHECK);  
  
pthread_mutex_t mut1;  
pthread_mutex_init(&mut1, &err_chk_attr);  
  
pthread_mutex_lock(&mut1);  
  
if (pthread_mutex_lock(&mut1) == EDEADLK)  
    fprintf(stderr, "error caught at runtime\n");  
  
if (pthread_mutex_unlock(&mut2) == EPERM)  
    fprintf(stderr, "another error: you didn't lock it!\n");
```

# Producer-Consumer Problem



# Guarded Commands

```
when (guard) [  
  /*
```

```
  /*
```

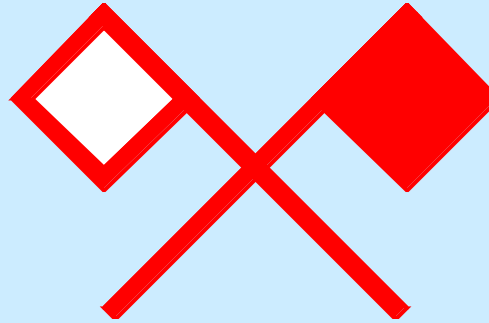
```
    once the guard is true, execute this  
    code atomically
```

```
  */
```

```
  . . .
```

```
]
```

# Semaphores



- **P(S) operation:**

```
when (S > 0) [  
    S = S - 1;  
]
```

- **V(S) operation:**

```
[S = S + 1;]
```

# Quiz 4

```
semaphore S = 1;  
int count = 0;
```

```
void proc( ) {  
    P(S);  
    count++;  
  
    ...  
    count--;  
    V(S);  
}
```

The function proc is called concurrently by n threads. What's the maximum value that count will take on?

- a) 1
- b) 2
- c) n
- d) indeterminate

- **P(S) operation:**  
    **when** (S > 0) [  
        S = S - 1;  
    ]
- **V(S) operation:**  
    [S = S + 1;]

# Producer/Consumer with Semaphores

```
Semaphore empty = BSIZE;  
Semaphore occupied = 0;  
int nextin = 0;  
int nextout = 0;
```

```
void Produce(char item) {  
    P(empty);  
    buf[nextin] = item;  
    if (++nextin >= BSIZE)  
        nextin = 0;  
    V(occupied);  
}
```

```
char Consume( ) {  
    char item;  
    P(occupied);  
    item = buf[nextout];  
    if (++nextout >= BSIZE)  
        nextout = 0;  
    V(empty);  
    return item;  
}
```

# POSIX Semaphores

```
#include <semaphore.h>
```

```
int sem_init(sem_t *semaphore, int pshared, int init);
```

```
int sem_destroy(sem_t *semaphore);
```

```
int sem_wait(sem_t *semaphore);
```

```
    /* P operation */
```

```
int sem_trywait(sem_t *semaphore);
```

```
    /* conditional P operation */
```

```
int sem_post(sem_t *semaphore);
```

```
    /* V operation */
```

# Producer-Consumer with POSIX Semaphores

```
sem_init(&empty, 0, BSIZE);  
sem_init(&occupied, 0, 0);  
int nextin = 0;  
int nextout = 0;
```

```
void produce(char item) {  
  
    sem_wait(&empty);  
    buf[nextin] = item;  
    if (++nextin >= BSIZE)  
        nextin = 0;  
    sem_post(&occupied);  
}
```

```
char consume( ) {  
    char item;  
    sem_wait(&occupied);  
    item = buf[nextout];  
    if (++nextout >= BSIZE)  
        nextout = 0;  
    sem_post(&empty);  
    return item;  
}
```