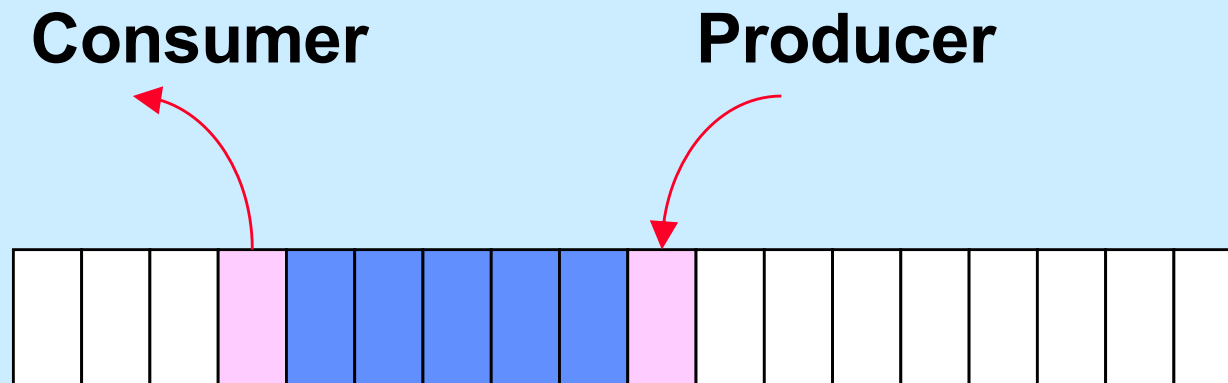


CS 33

Multithreaded Programming II

Producer-Consumer Problem



Guarded Commands

```
when (guard) [  
  /*
```

```
  /*
```

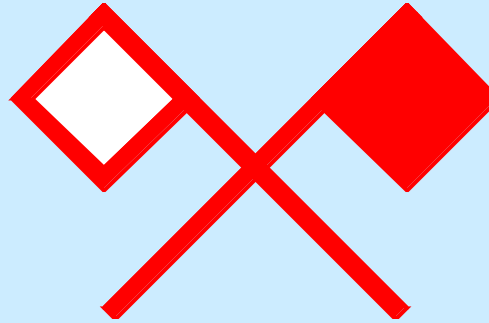
```
    once the guard is true, execute this  
    code atomically
```

```
  */
```

```
  . . .
```

```
]
```

Semaphores



- **P(S) operation:**

```
when (S > 0) [  
    S = S - 1;  
]
```

- **V(S) operation:**

```
[S = S + 1;]
```

Quiz 1

```
semaphore S = 1;  
int count = 0;
```

```
void proc( ) {  
    P(S);  
    count++;  
  
    ...  
    count--;  
    V(S);  
}
```

The function proc is called concurrently by n threads. What's the maximum value that count will take on?

- a) 1
- b) 2
- c) n
- d) indeterminate

- **P(S) operation:**
 when (S > 0) [
 S = S - 1;
]
- **V(S) operation:**
 [S = S + 1;]

Producer/Consumer with Semaphores

```
Semaphore empty = BSIZE;  
Semaphore occupied = 0;  
int nextin = 0;  
int nextout = 0;
```

```
void Produce(char item) {  
    P(empty);  
    buf[nextin] = item;  
    if (++nextin >= BSIZE)  
        nextin = 0;  
    V(occupied);  
}
```

```
char Consume( ) {  
    char item;  
    P(occupied);  
    item = buf[nextout];  
    if (++nextout >= BSIZE)  
        nextout = 0;  
    V(empty);  
    return item;  
}
```

POSIX Semaphores

```
#include <semaphore.h>
```

```
int sem_init(sem_t *semaphore, int pshared, int init);
```

```
int sem_destroy(sem_t *semaphore);
```

```
int sem_wait(sem_t *semaphore);
```

```
    /* P operation */
```

```
int sem_trywait(sem_t *semaphore);
```

```
    /* conditional P operation */
```

```
int sem_post(sem_t *semaphore);
```

```
    /* V operation */
```

Producer-Consumer with POSIX Semaphores

```
sem_init(&empty, 0, BSIZE);  
sem_init(&occupied, 0, 0);  
int nextin = 0;  
int nextout = 0;
```

```
void produce(char item) {  
  
    sem_wait(&empty);  
    buf[nextin] = item;  
    if (++nextin >= BSIZE)  
        nextin = 0;  
    sem_post(&occupied);  
}
```

```
char consume( ) {  
    char item;  
    sem_wait(&occupied);  
    item = buf[nextout];  
    if (++nextout >= BSIZE)  
        nextout = 0;  
    sem_post(&empty);  
    return item;  
}
```


Start/Stop



- **Start/Stop interface**

```
void wait_for_start(state_t *s);
```

```
void start(state_t *s);
```

```
void stop(state_t *s);
```

Start/Stop



- **Start/Stop interface**

```
void wait_for_start(state_t *s) {  
    if (s->state == stopped)  
        sleep();  
}  
  
void start(state_t *s) {  
    state = started;  
    wakeup_all();  
}  
  
void stop(state_t *s) {  
    state = stopped;  
}
```

Start/Stop



- **Start/Stop interface**

```
void wait_for_start(state_t *s) {  
    pthread_mutex_lock(&s->mutex);  
    if (s->state == stopped) {  
        pthread_mutex_unlock(&s->mutex);  
        sleep();  
    } else pthread_mutex_unlock(&s->mutex);  
}  
  
void start(state_t *s) {  
    pthread_mutex_lock(&s->mutex);  
    state = started;  
    wakeup_all();  
    pthread_mutex_unlock(&s->mutex);  
}
```

Start/Stop



- **Start/Stop interface**

```
void wait_for_start(state_t *s) {  
    pthread_mutex_lock(&s->mutex);  
    if (s->state == stopped) {  
        sleep();  
    }  
    pthread_mutex_unlock(&s->mutex);  
}  
  
void start(state_t *s) {  
    pthread_mutex_lock(&s->mutex);  
    state = started;  
    wakeup_all();  
    pthread_mutex_unlock(&s->mutex);  
}
```

Start/Stop



- **Start/Stop interface**

```
void wait_for_start(state_t *s) {  
    pthread_mutex_lock(&s->mutex);  
    while(s->state == stopped)  
        pthread_cond_wait(&s->queue, &s->mutex);  
    pthread_mutex_unlock(&s->mutex);  
}  
  
void start(state_t *s) {  
    pthread_mutex_lock(&s->mutex);  
    s->state = started;  
    pthread_cond_broadcast(&s->queue);  
    pthread_mutex_unlock(&s->mutex);  
}
```

Condition Variables

```
when (guard) [  
    statement 1;  
    ...  
    statement n;  
]
```

```
// code modifying the guard:  
...
```

```
pthread_mutex_lock(&mutex);  
while (!guard)  
    pthread_cond_wait(  
        &cond_var, &mutex);  
statement 1;  
...  
statement n;  
pthread_mutex_unlock(&mutex);
```

```
pthread_mutex_lock(&mutex);  
// code modifying the guard:  
...  
pthread_cond_broadcast(  
    &cond_var);  
pthread_mutex_unlock(&mutex);
```

Set Up

```
int pthread_cond_init(pthread_cond_t *cvp,  
    pthread_condattr_t *attrp)
```

```
int pthread_cond_destroy(pthread_cond_t *cvp)
```

```
int pthread_condattr_init(pthread_condattr_t *attrp)
```

```
int pthread_condattr_destroy(pthread_condattr_t *attrp)
```

PC with Condition Variables (1)

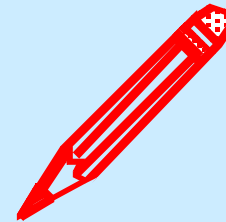
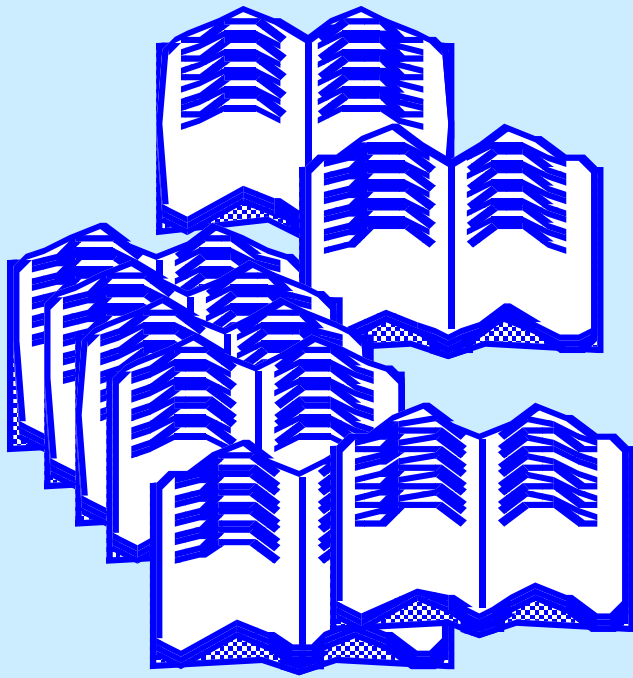
```
typedef struct buffer {  
    pthread_mutex_t m;  
    pthread_cond_t  more_space;  
    pthread_cond_t  more_items;  
    int             next_in;  
    int             next_out;  
    int             empty;  
    char            buf[BSIZE];  
} buffer_t;
```


PC with Condition Variables (2)

```
void produce(buffer_t *b,  
            char item) {  
  
    pthread_mutex_lock(&b->m);  
    while (!(b->empty > 0))  
        pthread_cond_wait(  
            &b->more_space, &b->m);  
    b->buf[b->nextin] = item;  
    if (++(b->nextin) == BSIZE)  
        b->nextin = 0;  
    b->empty--;  
    pthread_cond_signal(  
        &b->more_items);  
    pthread_mutex_unlock(&b->m);  
}
```

```
char consume(buffer_t *b) {  
    char item;  
    pthread_mutex_lock(&b->m);  
    while (!(b->empty < BSIZE))  
        pthread_cond_wait(  
            &b->more_items, &b->m);  
    item = b->buf[b->nextout];  
    if (++(b->nextout) == BSIZE)  
        b->nextout = 0;  
    b->empty++;  
    pthread_cond_signal(  
        &b->more_space);  
    pthread_mutex_unlock(&b->m);  
    return item;  
}
```

Readers-Writers Problem



Pseudocode

```
reader( ) {  
    when (writers == 0) [  
        readers++;  
    ]  
  
    /* read */  
  
    [readers--;]  
}
```

```
writer( ) {  
    when ((writers == 0) &&  
        (readers == 0)) [  
        writers++;  
    ]  
  
    /* write */  
  
    [writers--;]  
}
```

Pseudocode with Assertions

```
reader( ) {  
    when (writers == 0) [  
        readers++;  
    ]  
  
    assert((writers == 0) &&  
        (readers > 0));  
    /* read */  
  
    [readers--;]  
}
```

```
writer( ) {  
    when ((writers == 0) &&  
        (readers == 0)) [  
        writers++;  
    ]  
  
    assert((readers == 0) &&  
        (writers == 1));  
    /* write */  
  
    [writers--;]  
}
```

Solution with POSIX Threads

```
reader( ) {  
    pthread_mutex_lock(&m);  
    while (!(writers == 0))  
        pthread_cond_wait(  
            &readersQ, &m);  
    readers++;  
    pthread_mutex_unlock(&m);  
    /* read */  
    pthread_mutex_lock(&m);  
    if (--readers == 0)  
        pthread_cond_signal(  
            &writersQ);  
    pthread_mutex_unlock(&m);  
}
```

```
writer( ) {  
    pthread_mutex_lock(&m);  
    while (!(readers == 0) &&  
        (writers == 0))  
        pthread_cond_wait(  
            &writersQ, &m);  
    writers++;  
    pthread_mutex_unlock(&m);  
    /* write */  
    pthread_mutex_lock(&m);  
    writers--;  
    pthread_cond_signal(  
        &writersQ);  
    pthread_cond_broadcast(  
        &readersQ);  
    pthread_mutex_unlock(&m);  
}
```

New Pseudocode

```
reader( ) {  
    when (writers == 0) [  
        readers++;  
    ]  
  
    /* read */  
  
    [readers--;]  
}
```

```
writer( ) {  
    [writers++;]  
    when ((readers == 0) &&  
        (active_writers == 0)) [  
        active_writers++;  
    ]  
  
    /* write */  
  
    [writers--;  
    active_writers--;]  
}
```

Improved Reader

```
reader( ) {  
    pthread_mutex_lock(&m);  
  
    while (!(writers == 0)) {  
        pthread_cond_wait(  
            &readersQ, &m);  
    }  
    readers++;  
  
    pthread_mutex_unlock(&m);  
  
    /* read */
```

```
        pthread_mutex_lock(&m);  
  
        if (--readers == 0)  
            pthread_cond_signal(  
                &writersQ);  
  
        pthread_mutex_unlock(&m);  
    }
```

Improved Writer

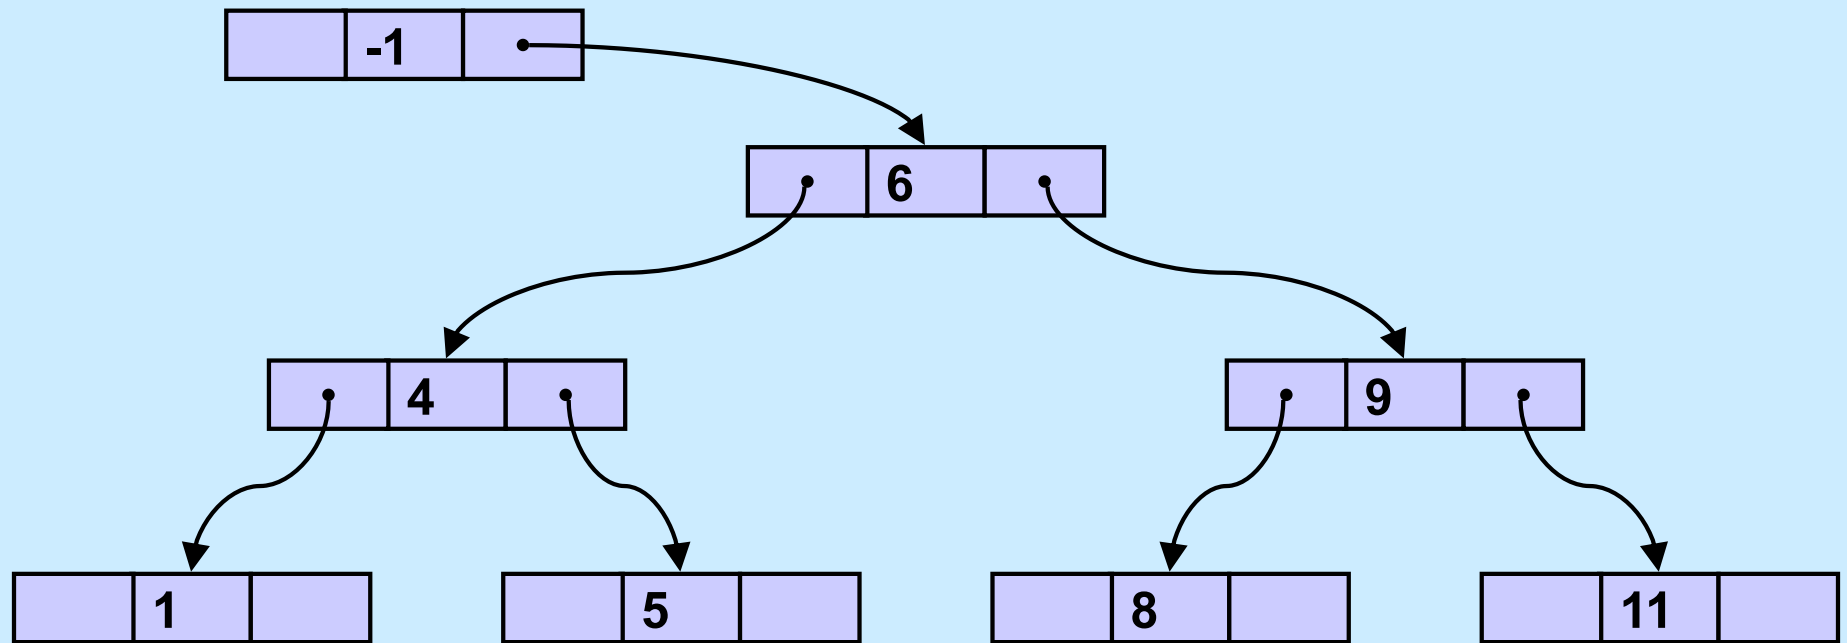
```
writer( ) {  
    pthread_mutex_lock(&m);  
  
    writers++;  
    while (!(readers == 0) &&  
           (active_writers == 0)) {  
        pthread_cond_wait(  
            &writersQ, &m);  
    }  
    active_writers++;  
  
    pthread_mutex_unlock(&m);  
  
    /* write */
```

```
    pthread_mutex_lock(&m);  
  
    writers--;  
    active_writers--;  
    if (writers)  
        pthread_cond_signal(  
            &writersQ);  
    else  
        pthread_cond_broadcast(  
            &readersQ);  
  
    pthread_mutex_unlock(&m);  
}
```

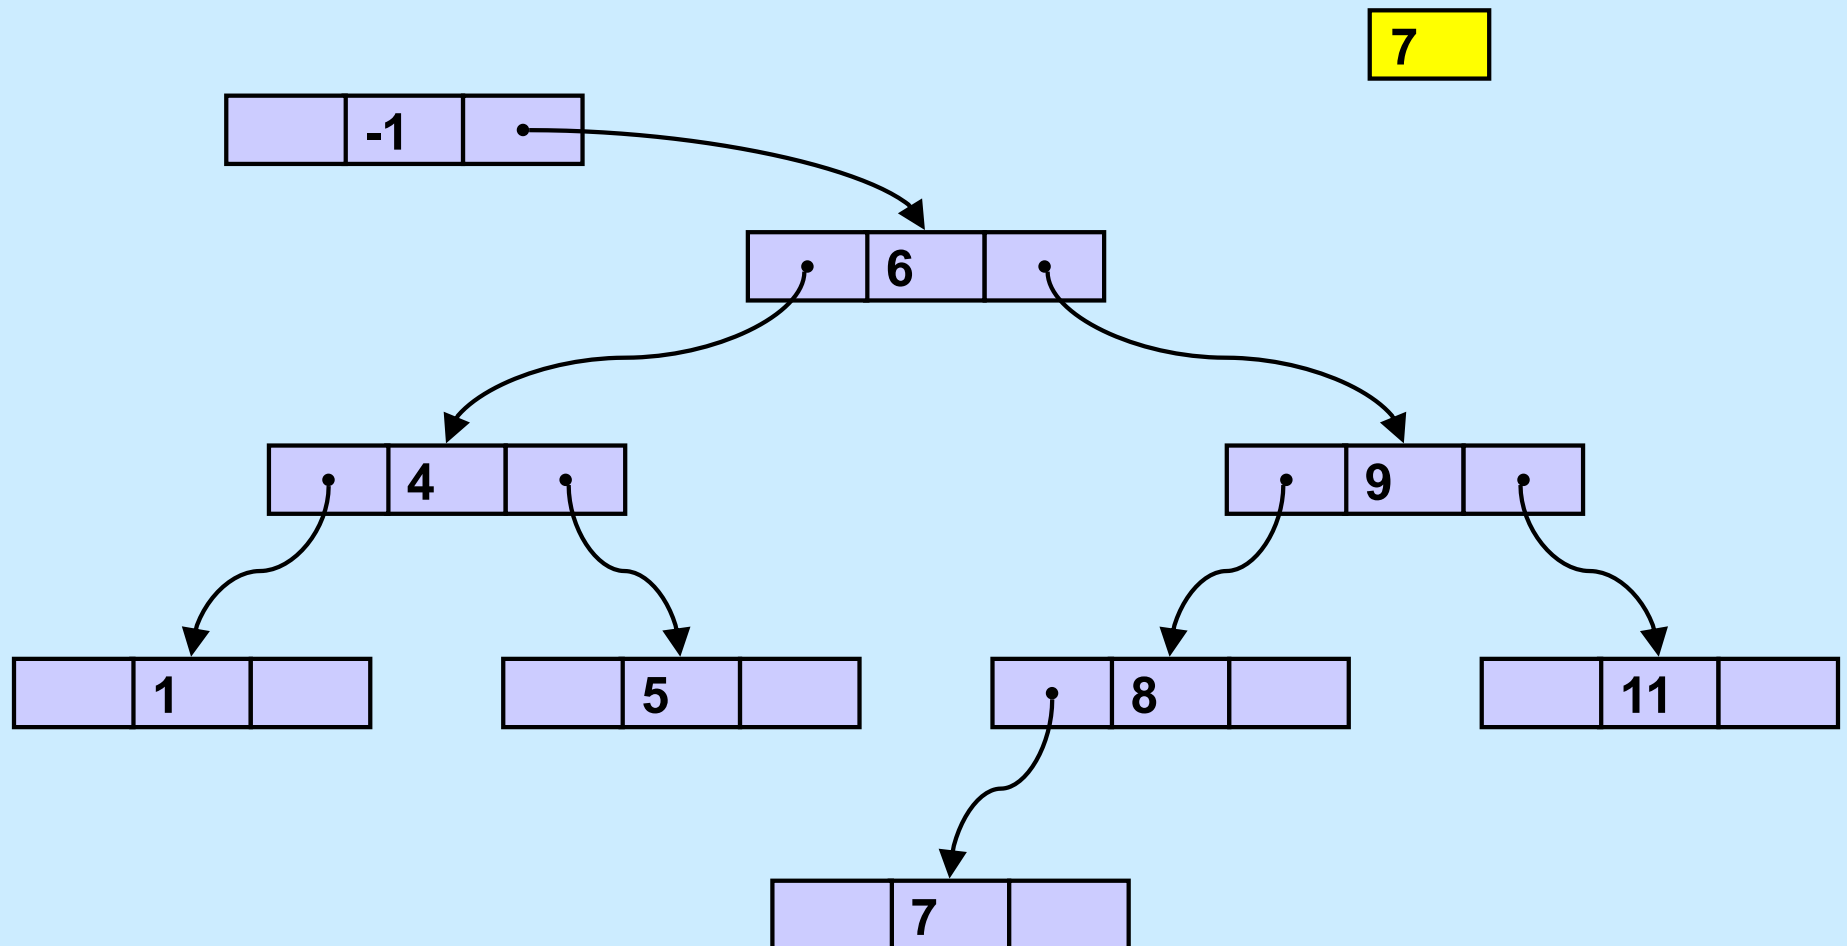

New, From POSIX!

```
int pthread_rwlock_init(pthread_rwlock_t *lock,  
    pthread_rwlockattr_t *att);  
int pthread_rwlock_destroy(pthread_rwlock_t *lock);  
int pthread_rwlock_rdlock(pthread_rwlock_t *lock);  
int pthread_rwlock_wrlock(pthread_rwlock_t *lock);  
int pthread_rwlock_tryrdlock(pthread_rwlock_t *lock);  
int pthread_rwlock_trywrlock(pthread_rwlock_t *lock);  
int pthread_timedrwlock_rdlock(pthread_rwlock_t *lock,  
    struct timespec *ts);  
int pthread_timedrwlock_wrlock(pthread_rwlock_t *lock,  
    struct timespec *ts);  
int pthread_rwlock_unlock(pthread_rwlock_t *lock);
```

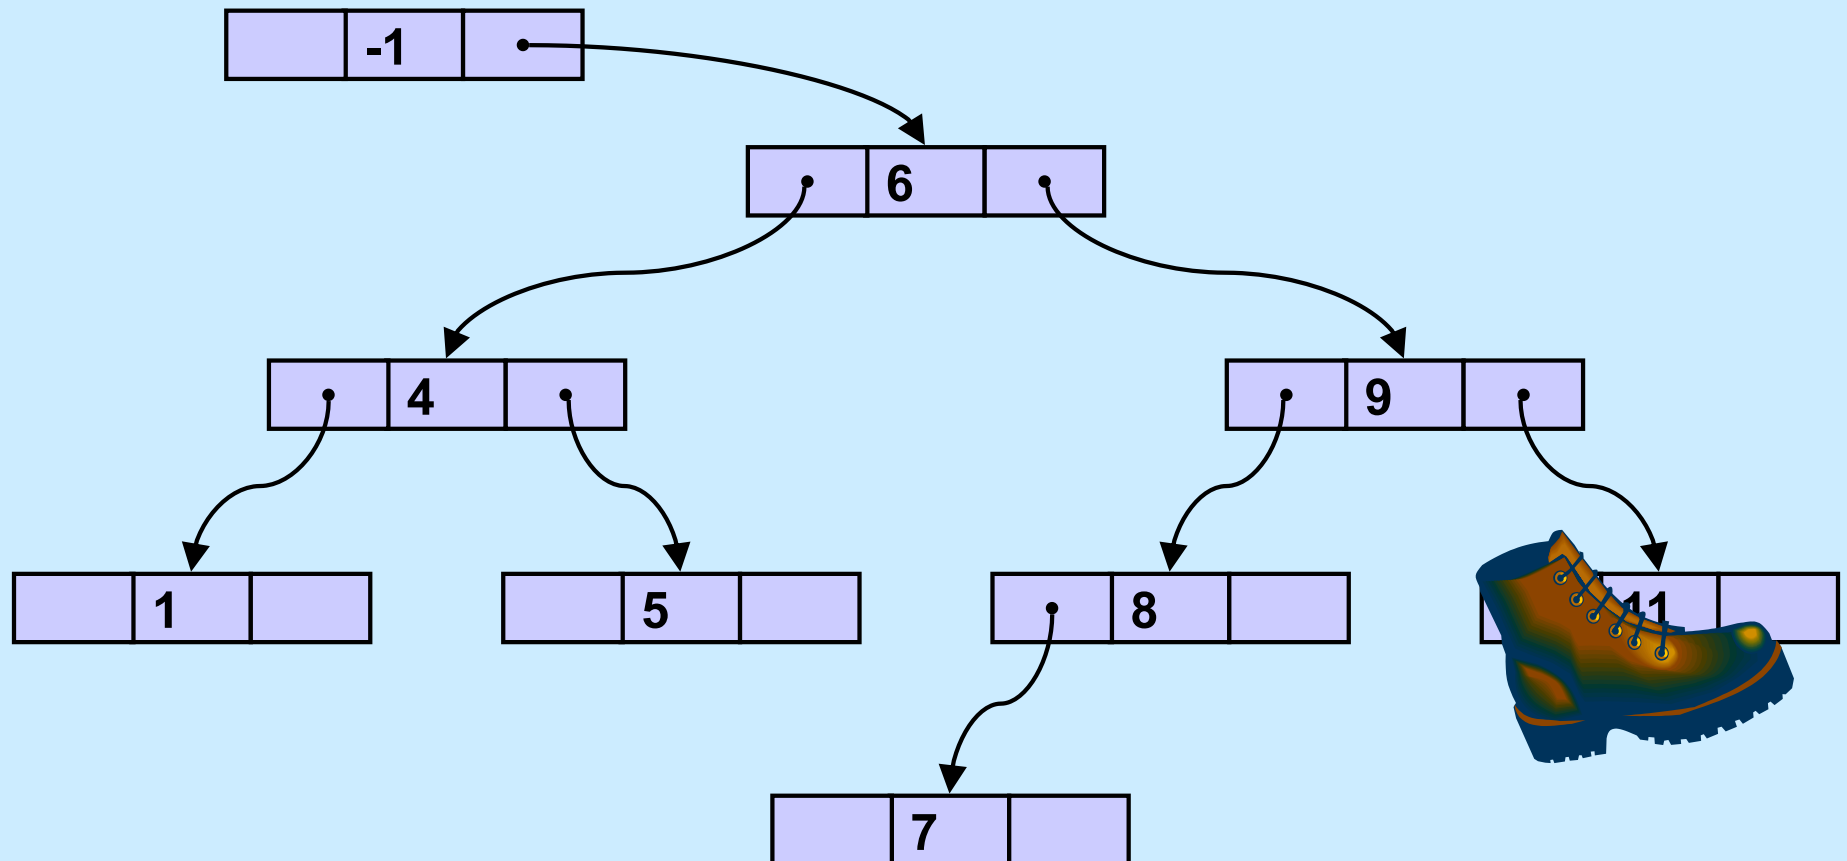
Binary Search Tree



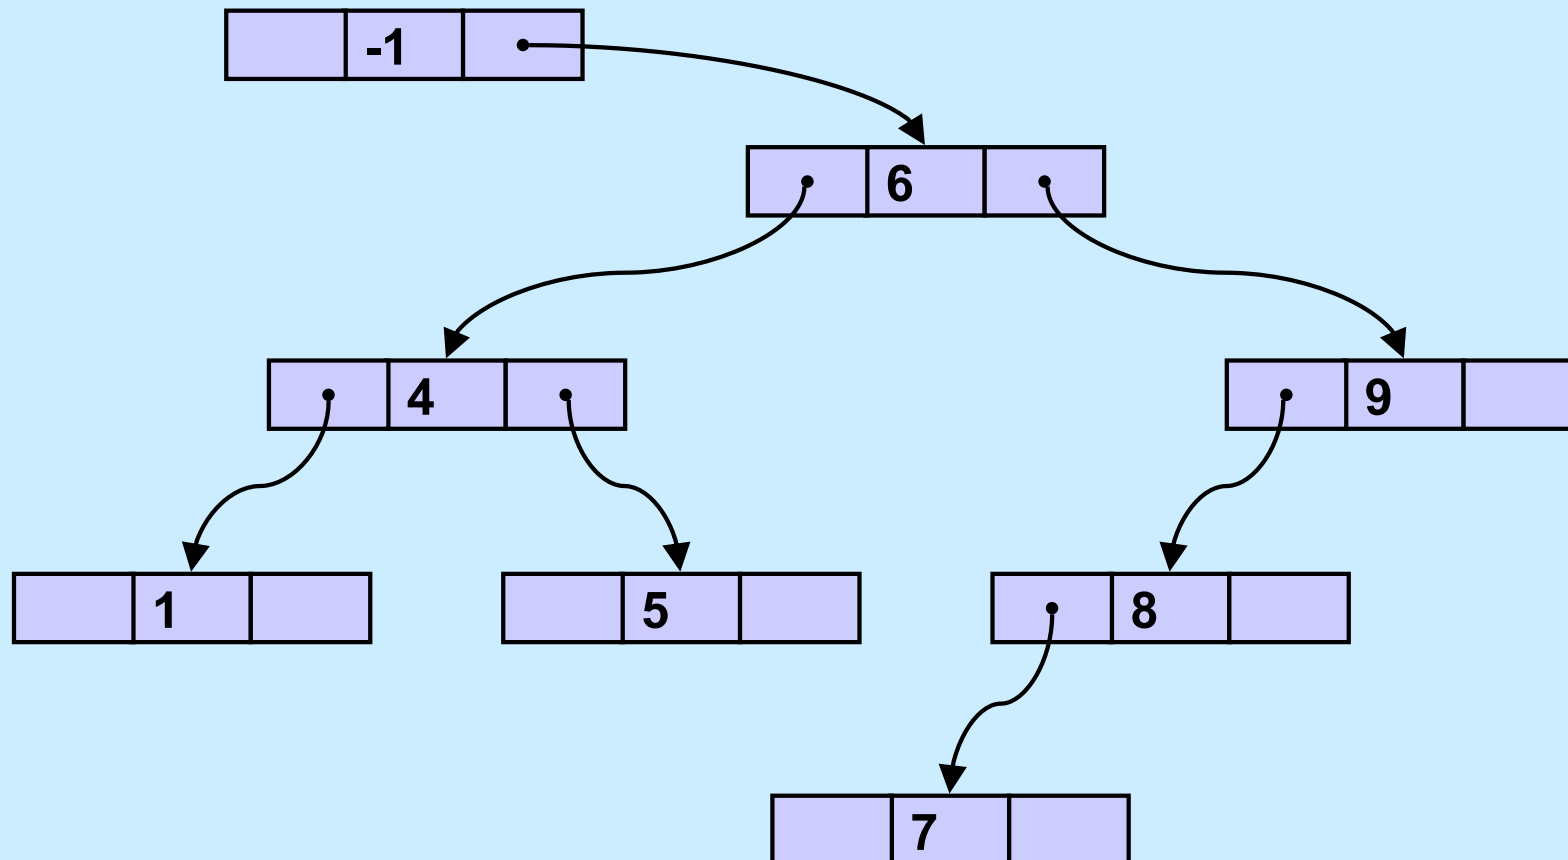
Binary Search Tree: Insertion



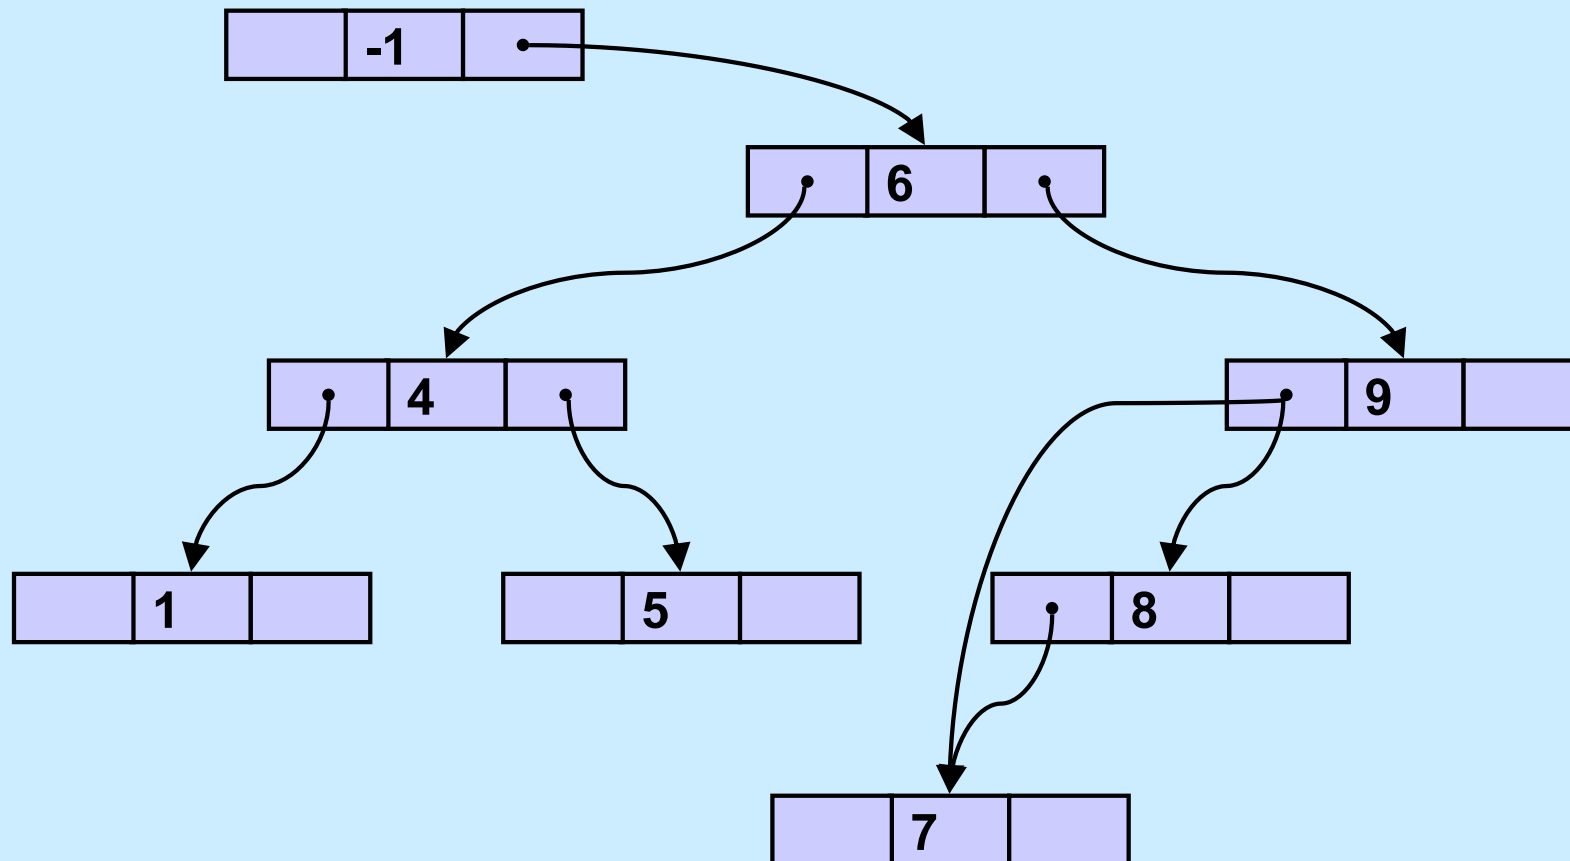
Binary Search Tree: Deletion of Leaf



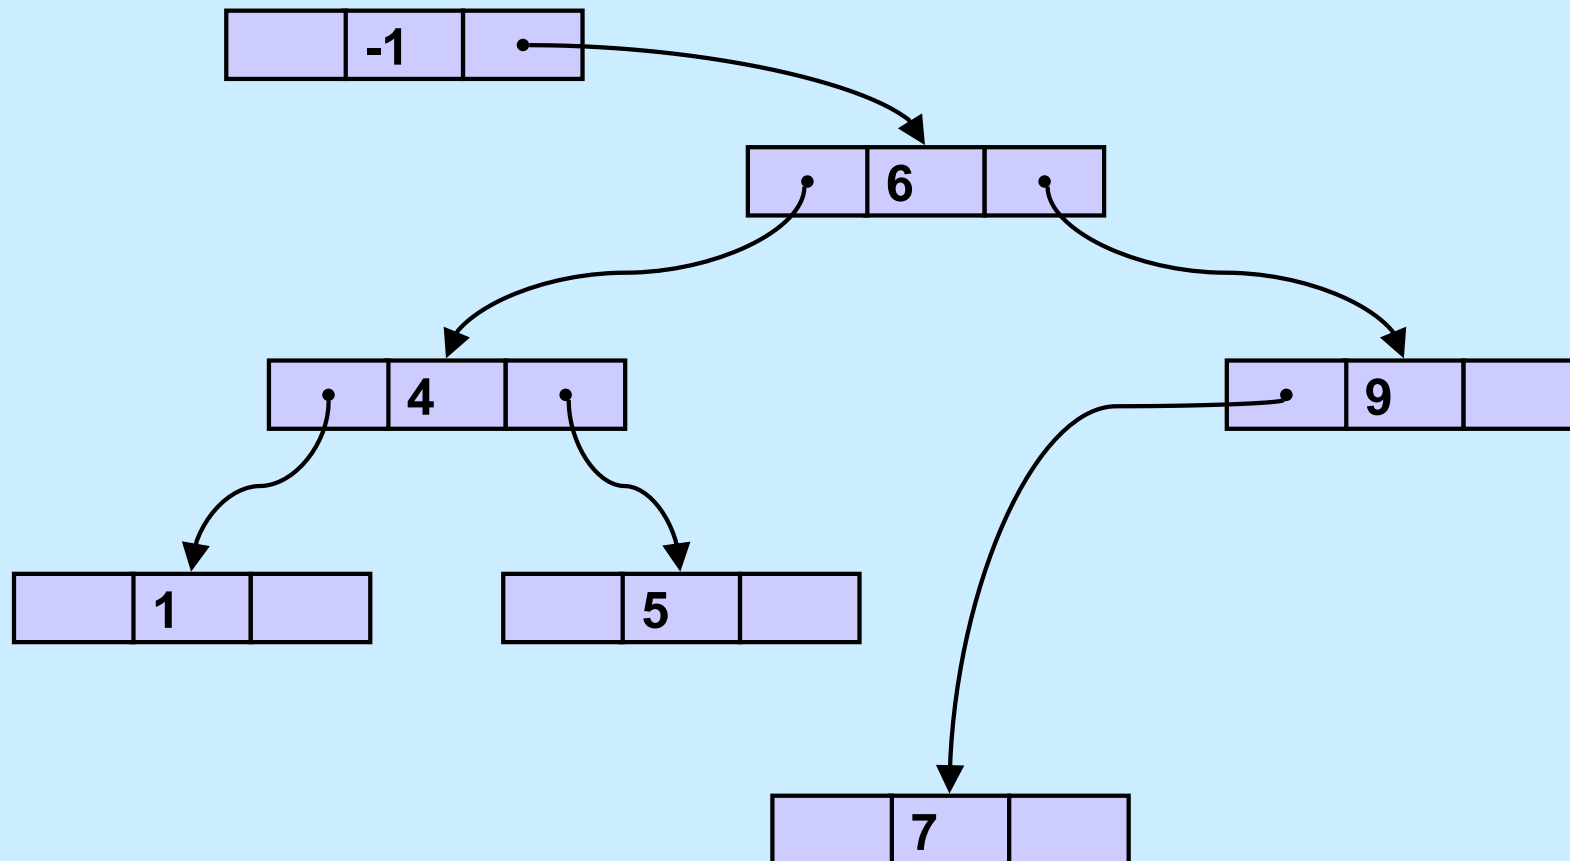
Binary Search Tree: Deletion of Leaf



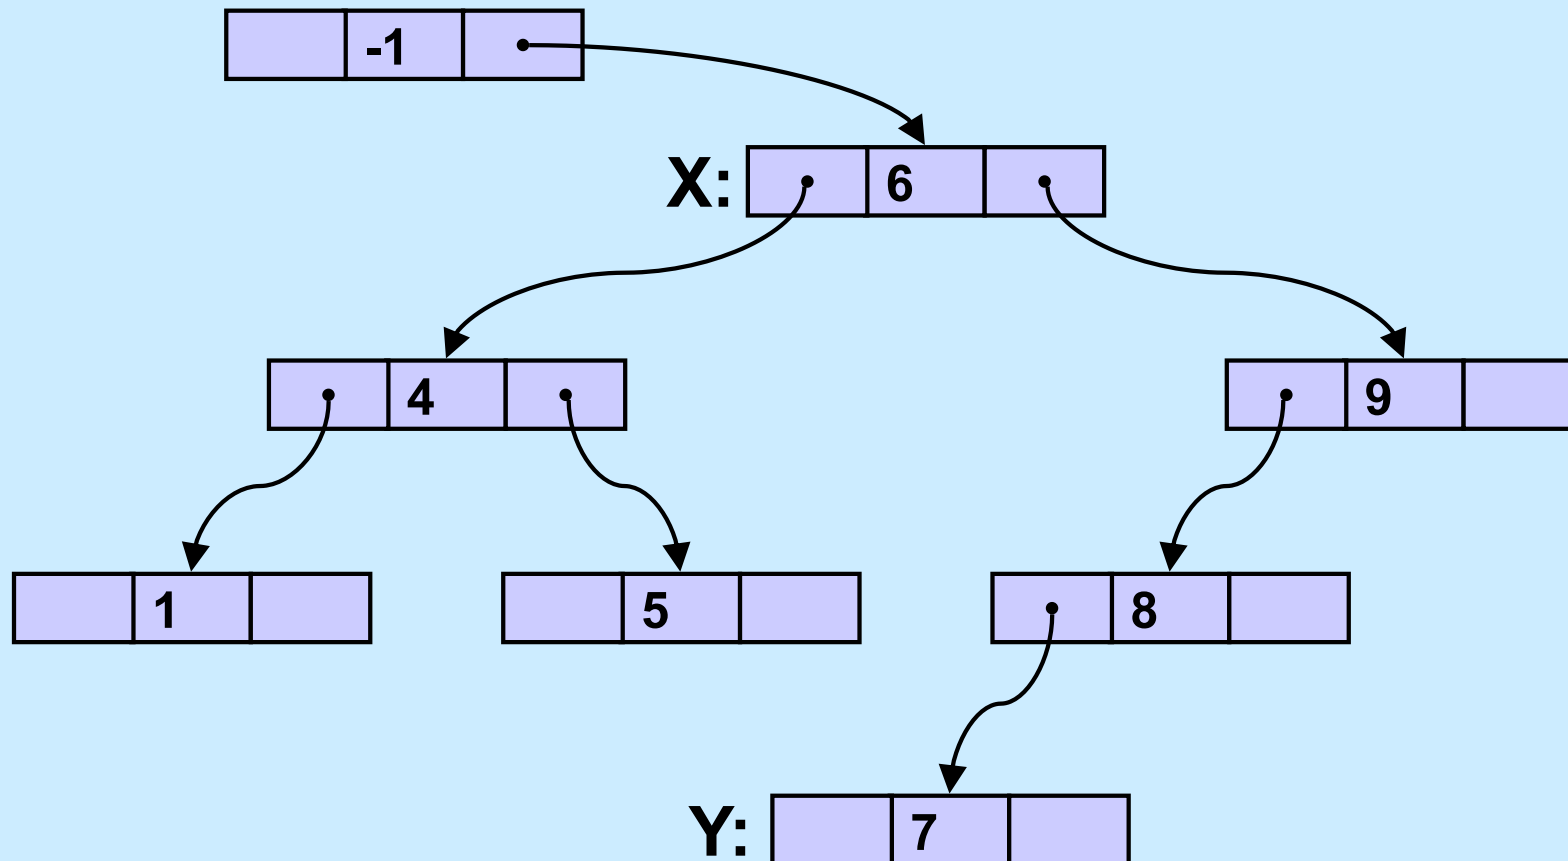
Binary Search Tree: Deletion of Node with One Child



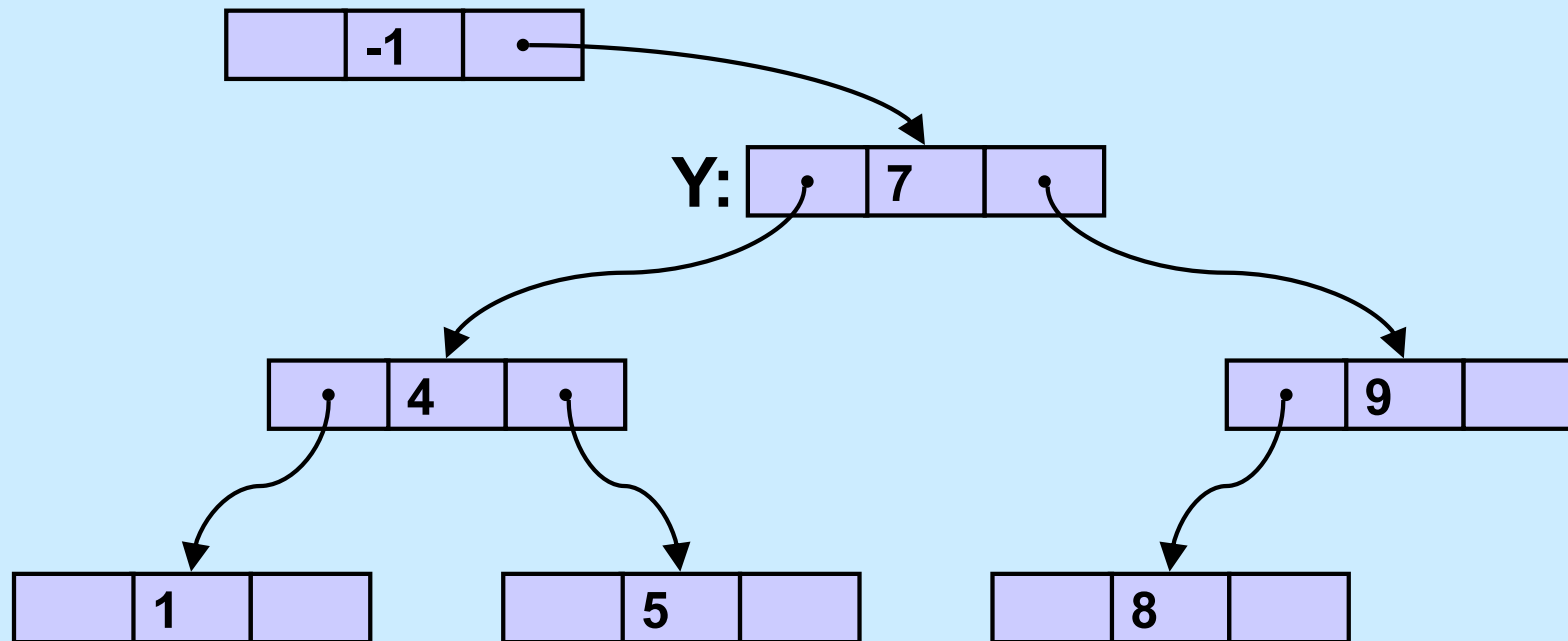
Binary Search Tree: Deletion of Node with One Child



Binary Search Tree: Deletion of Node with Two Children



Binary Search Tree: Deletion of Node with Two Children



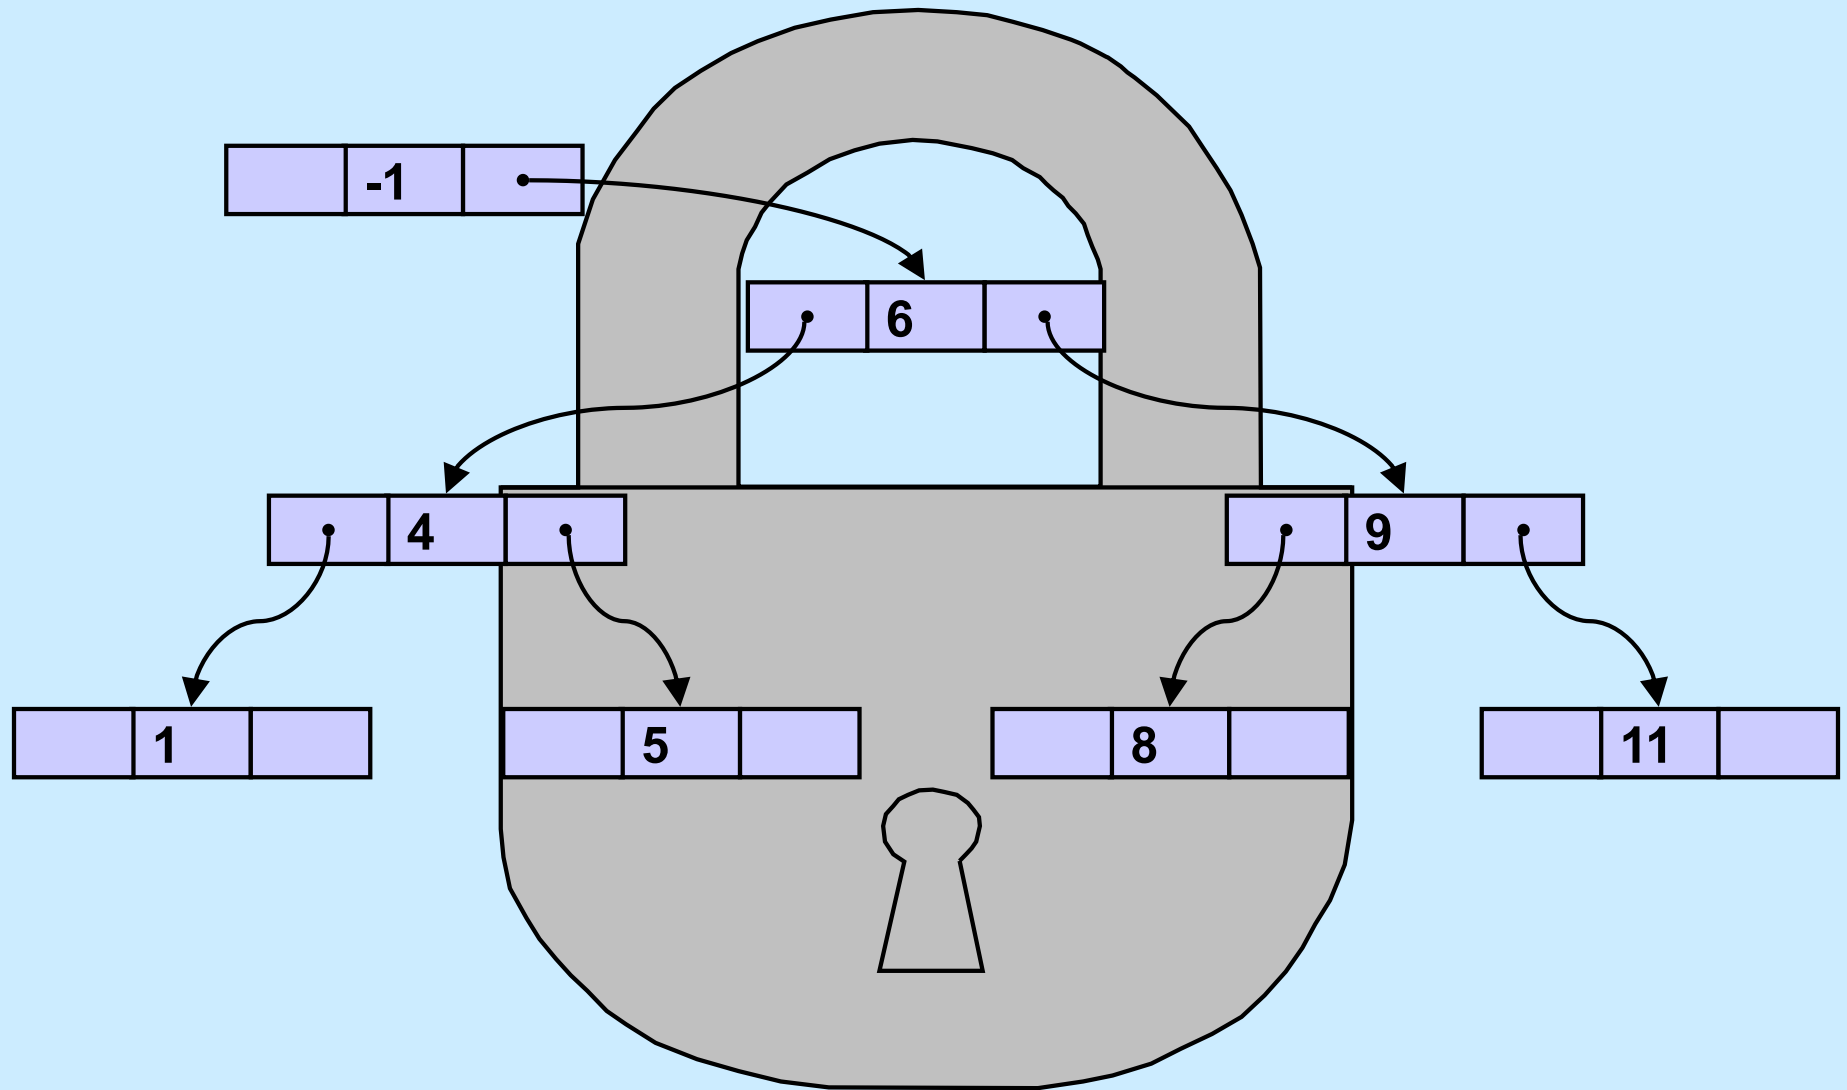
C Code: Search

```
Node *search(int key,
             Node *parent, Node **parentp) {
    Node *next;
    Node *result;
    if (key < parent->key) {
        if ((next = parent->lchild)
            == 0) {
            result = 0;
        } else {
            if (key == next->key) {
                result = next;
            } else {
                result = search(key,
                               next, parentpp);
            }
            return result;
        }
    } else {
        if ((next = parent->rchild)
            == 0) {
            result = 0;
        } else {
            if (key == next->key) {
                result = next;
            } else {
                result = search(key,
                               next, parentpp);
            }
            return result;
        }
    }
    if (parentpp != 0)
        *parentpp = parent;
    return result;
}
```

C Code: Add

```
int add(int key) {
    Node *parent, *target, *newnode;
    if ((target = search(key, &head, &parent)) != 0) {
        return 0;
    }
    newnode = malloc(sizeof(Node));
    newnode->key = key;
    newnode->lchild = newnode->rchild = 0;
    if (name < parent->name)
        parent->lchild = newnode;
    else
        parent->rchild = newnode;
    return 1;
}
```

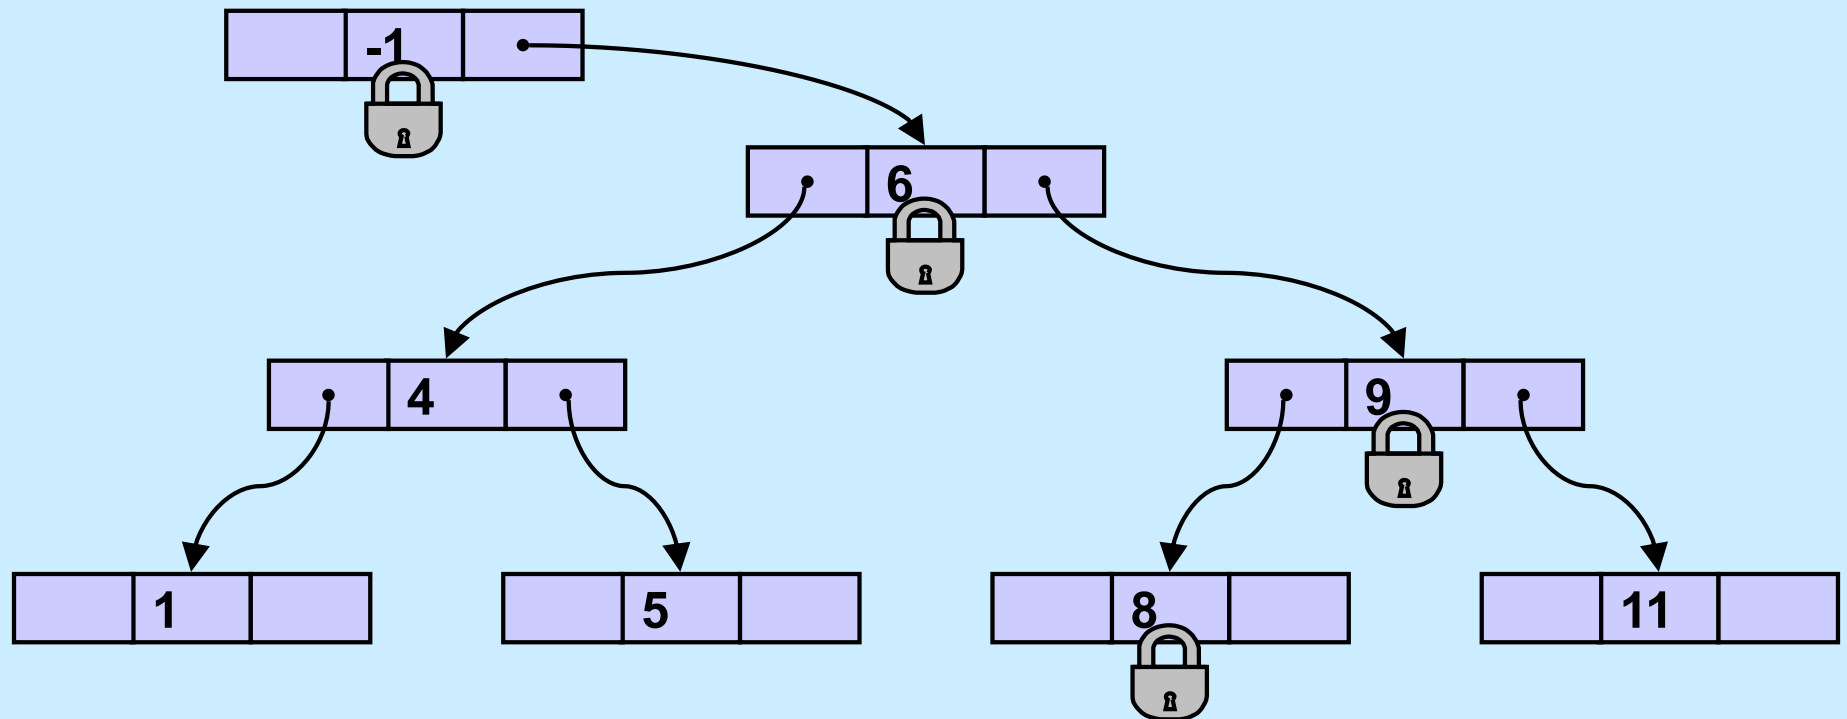
Binary Search Tree with Coarse-Grained Synchronization



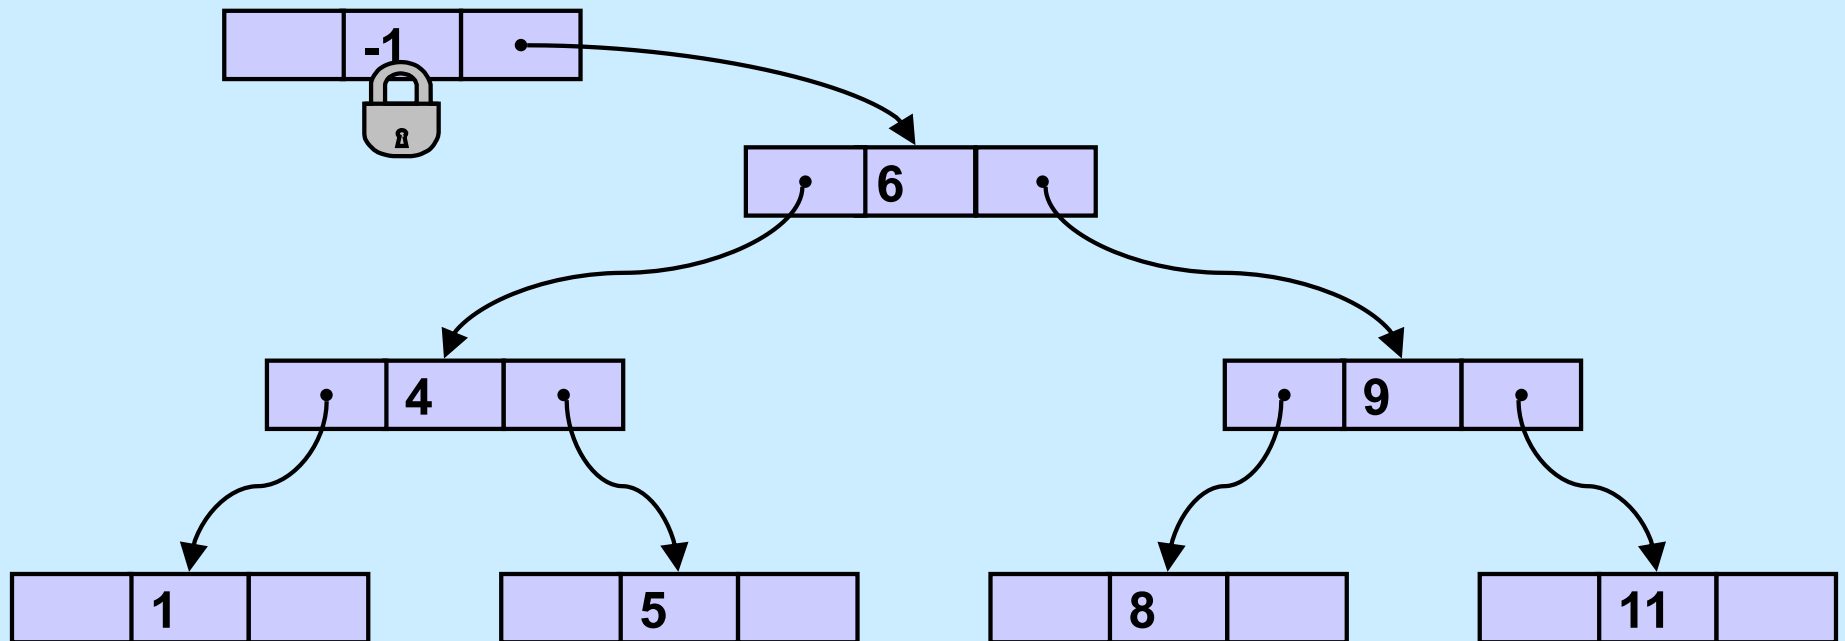
C Code: Add with Coarse-Grained Synchronization

```
int add(int key) {
    Node *parent, *target, *newnode;
    pthread_rwlock_wrlock(&tree_lock);
    if ((target = search(key, &head, &parent)) != 0) {
        pthread_rwlock_unlock(&tree_lock);
        return 0;
    }
    newnode = malloc(sizeof(Node));
    newnode->key = key;
    newnode->lchild = newnode->rchild = 0;
    if (name < parent->name)
        parent->lchild = newnode;
    else
        parent->rchild = newnode;
    pthread_rwlock_unlock(&tree_lock);
    return 1;
}
```

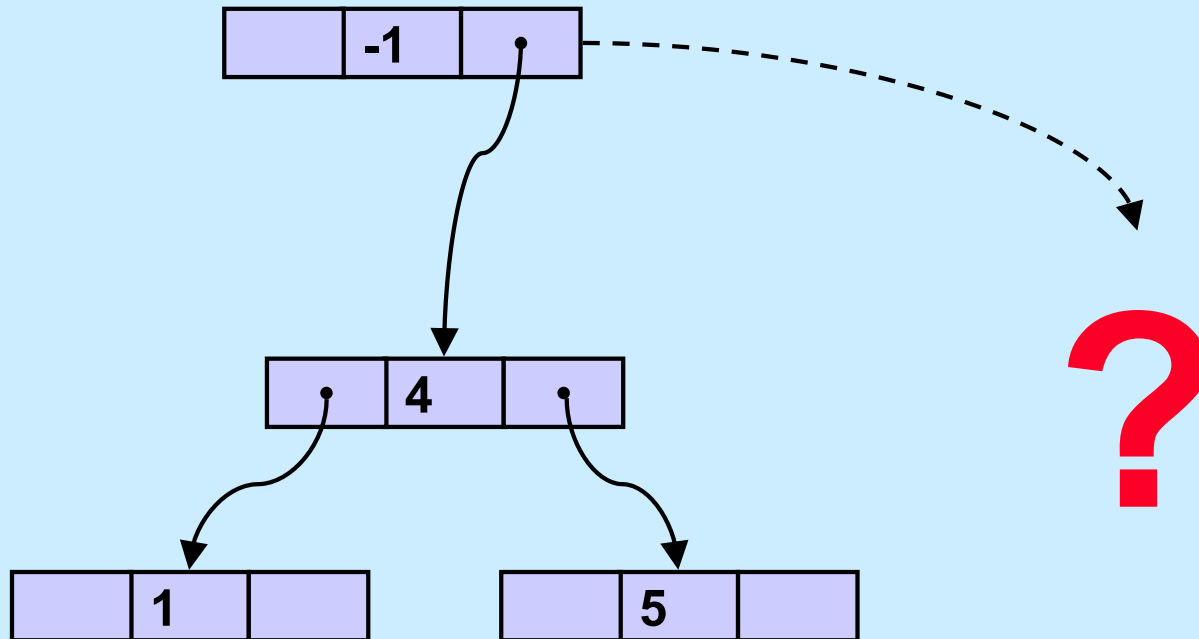
Binary Search Tree with Fine-Grained Synchronization I



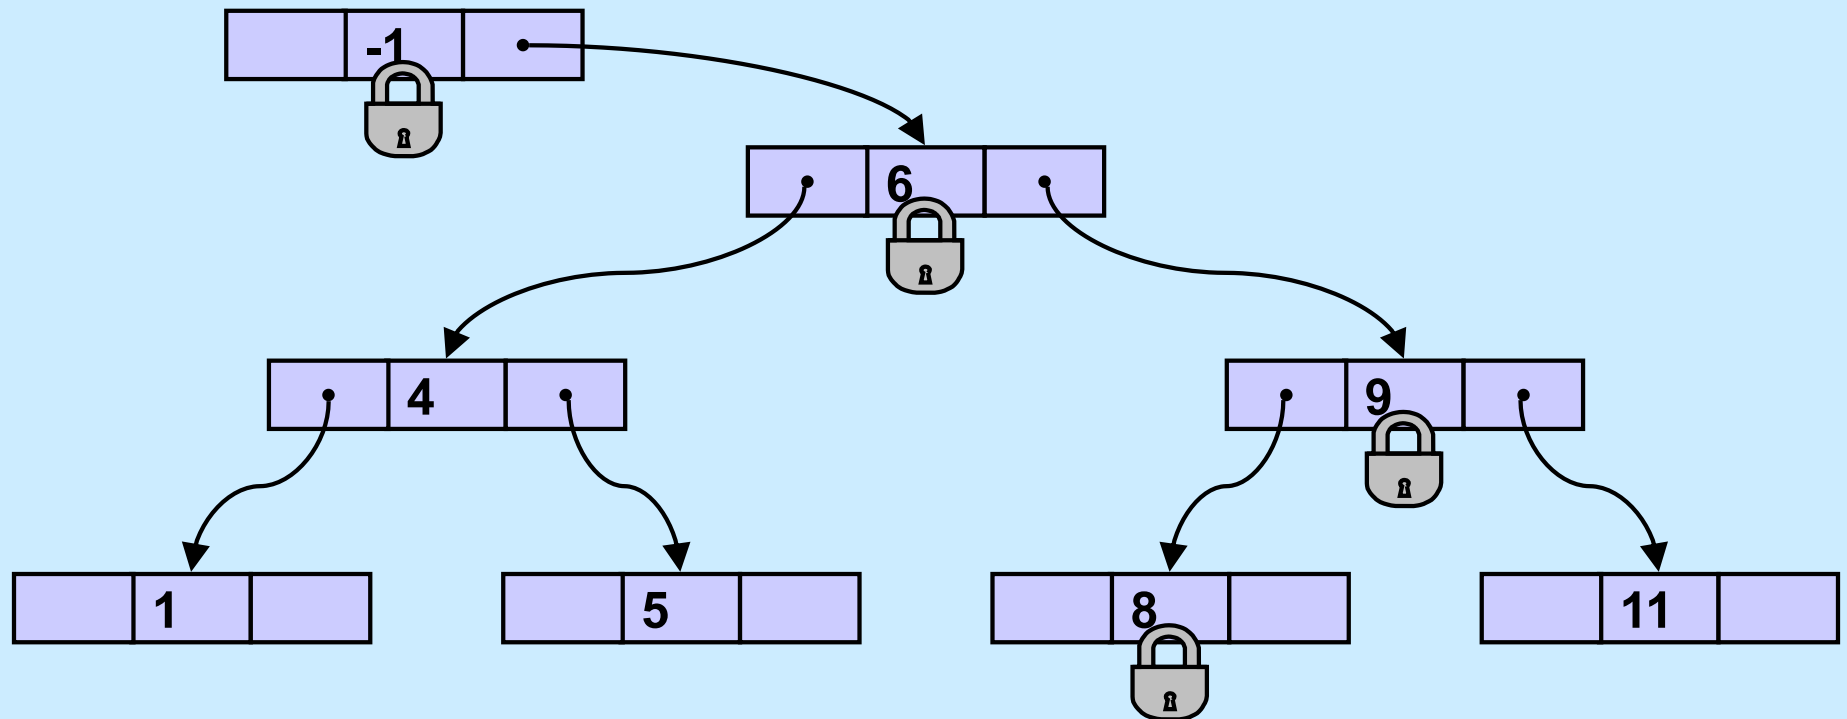
Binary Search Tree with Fine-Grained Synchronization II



Binary Search Tree with Fine-Grained Synchronization III



Doing It Right ...



C Code: Fine-Grained Search I

```
enum locktype {l_read, l_write};

#define lock(lt, lk) ((lt) == l_read)?
    pthread_rwlock_rdlock(lk):
    pthread_rwlock_wrlock(lk)

Node *search(int key,
    Node *parent, Node **parentp,
    enum locktype lt) {
    // parent is locked on entry
    Node *next;
    Node *result;
    if (key < parent->key) {
        if ((next = parent->lchild)
            == 0) {
            result = 0;
        } else {
            lock(lt, &next->lock);
            if (key == next->key) {
                result = next;
            } else {
                pthread_rwlock_unlock(
                    &parent->lock);
                result = search(key,
                    next, parentp, lt);
            }
        }
    }
    return result;
}
```

C Code: Fine-Grained Search II

```
} else {
    if ((next = parent->rchild)
        == 0) {
        result = 0;
    } else {
        lock(lt, &next->lock);
        if (key == next->key) {
            result = next;
        } else {
            pthread_rwlock_unlock(
                &parent->lock);
            result = search(key,
                next, parentpp, lt);
            return result;
        }
    }
}

if (parenttp != 0) {
    // parent remains locked
    *parenttp = parent;
} else
    pthread_rwlock_unlock(
        &parent->lock);
return result;
}
```

C Code: Add with Fine-Grained Synchronization I

```
int add(int key) {  
    Node *parent, *target, *newnode;  
    pthread_rwlock_wrlock(&head->lock);  
    if ((target = search(key, &head, &parent,  
        l_write)) != 0) {  
        pthread_rwlock_unlock(&target->lock);  
        pthread_rwlock_unlock(&parent->lock);  
        return 0;  
    }  
}
```

C Code: Add with Fine-Grained Synchronization II

```
newnode = malloc(sizeof(Node));
newnode->key = key;
newnode->lchild = newnode->rchild = 0;
pthread_rwlock_init(&newnode->lock, 0);
if (name < parent->name)
    parent->lchild = newnode;
else
    parent->rchild = newnode;
pthread_rwlock_unlock(&parent->lock);
return 1;
}
```