

CS 33

Hardware: The Memory Hierarchy

Many of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Random-Access Memory (RAM)

- **Key features**
 - **RAM** is traditionally packaged as a chip
 - basic storage unit is normally a **cell** (one bit per cell)
 - multiple RAM chips form a memory
- **Static RAM (SRAM)**
 - each cell stores a bit with a four- or six-transistor circuit
 - retains value indefinitely, as long as it is kept powered
 - relatively insensitive to electrical noise (EMI), radiation, etc.
 - faster and more expensive than DRAM
- **Dynamic RAM (DRAM)**
 - each cell stores bit with a capacitor; transistor is used for access
 - value must be refreshed every 10-100 ms
 - more sensitive to disturbances (EMI, radiation,...) than SRAM
 - slower and cheaper than SRAM

Supplied by CMU.

SRAM vs DRAM Summary

	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 or 6	1X	No	Maybe	100x	Cache memories
DRAM	1	10X	Yes	Yes	1X	Main memories, frame buffers

- **EDC = error detection and correction**
 - to cope with noise, etc.

Supplied by CMU.

What's Inside A Disk Drive?





Image courtesy of Seagate Technology

Disk Geometry


- Disks consist of **platters**, each with two **surfaces**
- Each surface consists of concentric rings called **tracks**
- Each track consists of **sectors** separated by **gaps**



Supplied by CMU.

Disk Geometry (Multiple-Platter View)

- Aligned tracks form a cylinder



Supplied by CMU.

Disk Capacity

- **Capacity**: maximum number of bits that can be stored
 - capacity expressed in units of gigabytes (GB), where $1 \text{ GB} = 2^{30} \text{ Bytes} \approx 10^9 \text{ Bytes}$
- Capacity is determined by these technology factors:
 - **recording density** (bits/in): number of bits that can be squeezed into a 1 inch segment of a track
 - **track density** (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment
 - **areal density** (bits/in²): product of recording and track density
- Modern disks partition tracks into disjoint subsets called **recording zones**
 - each track in a zone has the same number of sectors, determined by the circumference of innermost track
 - each zone has a different number of sectors/track

Supplied by CMU.

Computing Disk Capacity

**Capacity = (# bytes/sector) x (avg. # sectors/track) x
(# tracks/surface) x (# surfaces/platter) x
(# platters/disk)**

Example:


- 512 bytes/sector
- 600 sectors/track (on average)
- 40,000 tracks/surface
- 2 surfaces/platter
- 5 platters/disk

**Capacity = 512 x 600 x 40000 x 2 x 5
= 122,880,000,000
= 113.88 GB**

Supplied by CMU.


Note that 1GB = 2^{30} bytes.

Disk Operation (Single-Platter View)




Supplied by CMU.

Disk Operation (Multi-Platter View)



Supplied by CMU.

Disk Structure: Top View of Single Platter




Surface organized into tracks

Tracks divided into sectors

Supplied by CMU.


Disk Access



Head in position above a track

Supplied by CMU.


Disk Access



Rotation is counter-clockwise

Supplied by CMU.


Disk Access – Read



About to read blue sector

Supplied by CMU.


Disk Access – Read



After **BLUE**
read

After reading blue sector


Disk Access – Read



After **BLUE**
read


Red request scheduled next

Disk Access – Seek




Seek to red's track

Disk Access – Rotational Latency



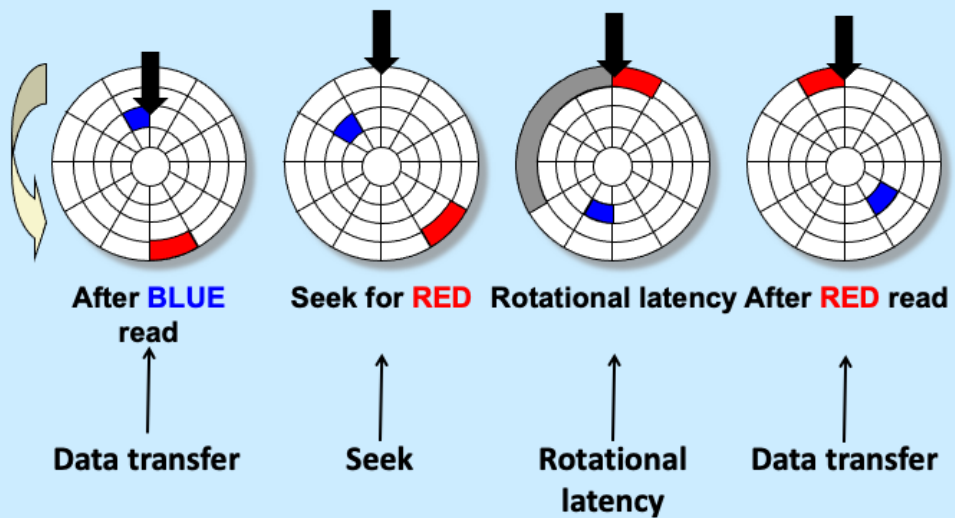
Disk Access – Read



Complete read of red

Supplied by CMU.

Disk Access – Service Time Components



Supplied by CMU.

Disk Access Time

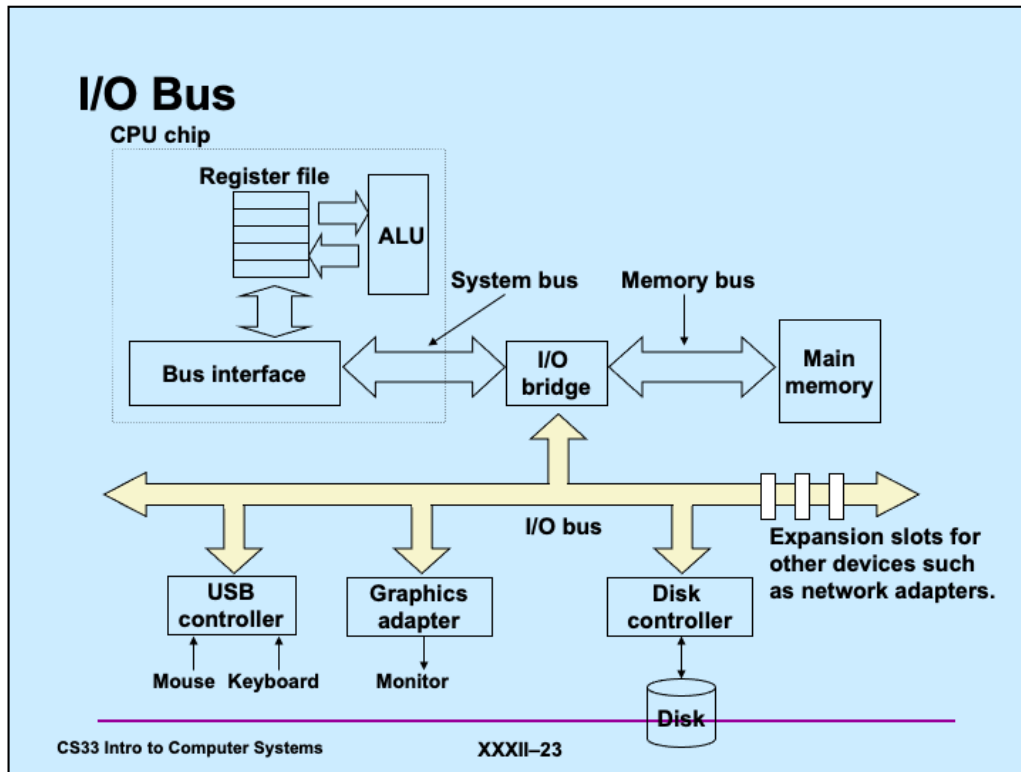
- Average time to access some target sector approximated by :
 - $T_{\text{access}} = T_{\text{avg seek}} + T_{\text{avg rotation}} + T_{\text{avg transfer}}$
- **Seek time** ($T_{\text{avg seek}}$)
 - time to position heads over cylinder containing target sector
 - typical $T_{\text{avg seek}}$ is 3–9 ms
- **Rotational latency** ($T_{\text{avg rotation}}$)
 - time waiting for first bit of target sector to pass under r/w head
 - typical rotation speed $R = 7200$ RPM
 - $T_{\text{avg rotation}} = \frac{1}{2} \times \frac{1}{R} \times 60 \text{ sec/1 min}$
- **Transfer time** ($T_{\text{avg transfer}}$)
 - time to read the bits in the target sector
 - $T_{\text{avg transfer}} = \frac{1}{R} \times \frac{1}{(\text{avg \# sectors/track})} \times 60 \text{ secs/1 min}$

Supplied by CMU.

Disk Access Time Example

- **Given:**
 - rotational rate = 7,200 RPM
 - average seek time = 9 ms
 - avg # sectors/track = 600
- **Derived:**
 - Tavg rotation = $1/2 \times (60 \text{ secs}/7200 \text{ RPM}) \times 1000 \text{ ms/sec} = 4 \text{ ms}$
 - Tavg transfer = $60/7200 \text{ RPM} \times 1/600 \text{ sects/track} \times 1000 \text{ ms/sec} = 0.014 \text{ ms}$
 - Taccess = 9 ms + 4 ms + 0.014 ms
- **Important points:**
 - access time dominated by seek time and rotational latency
 - first bit in a sector is the most expensive, the rest are free
 - SRAM access time is about 4 ns/doubleword, DRAM about 60 ns
 - » disk is about 40,000 times slower than SRAM
 - » 2,500 times slower than DRAM

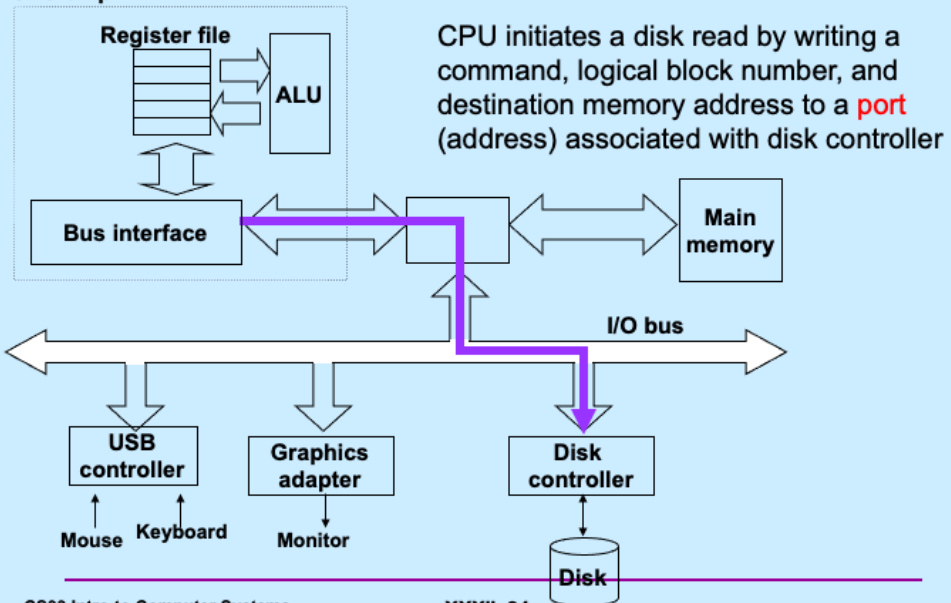
Supplied by CMU.



Supplied by CMU.

Reading a Disk Sector (1)

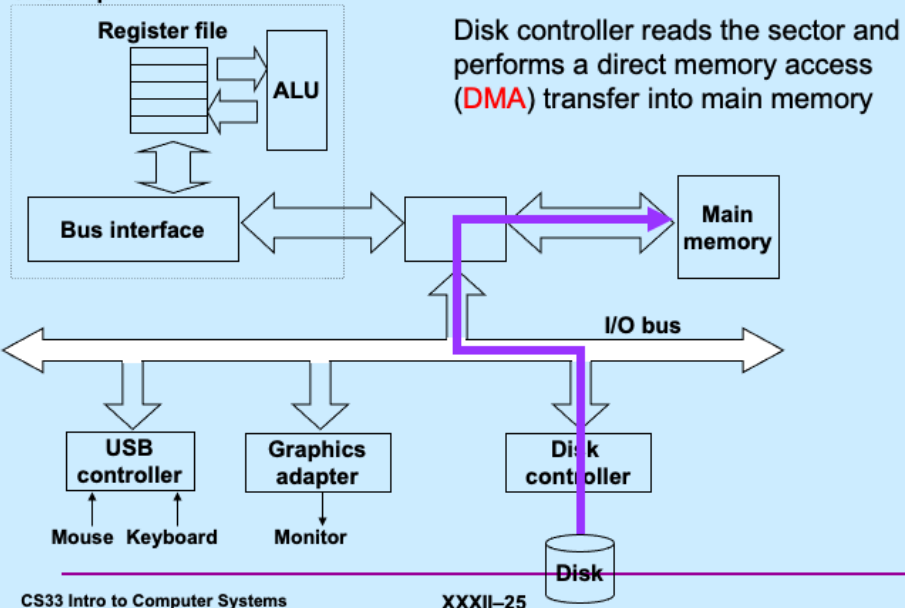
CPU chip



Supplied by CMU.

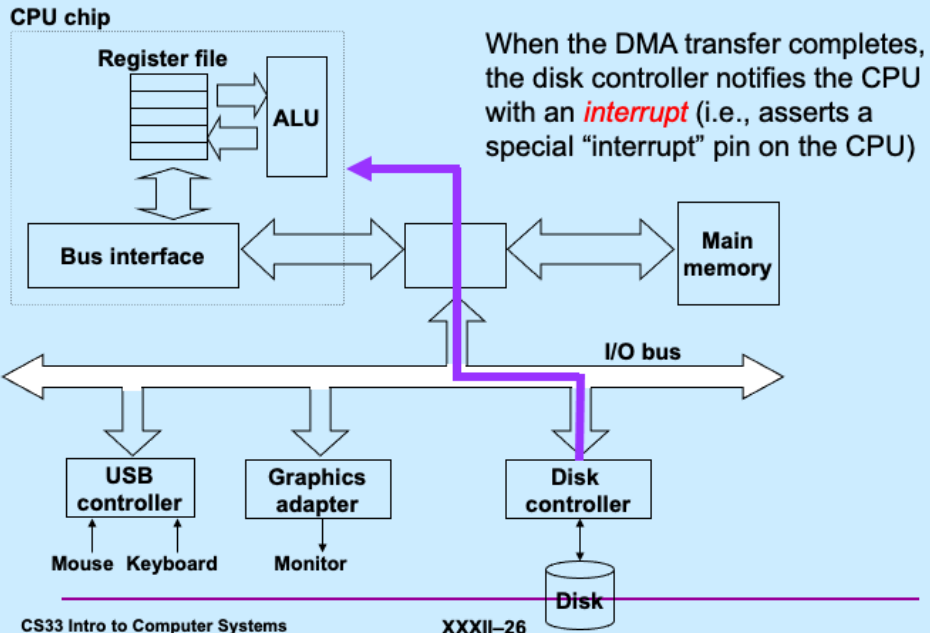
Reading a Disk Sector (2)

CPU chip



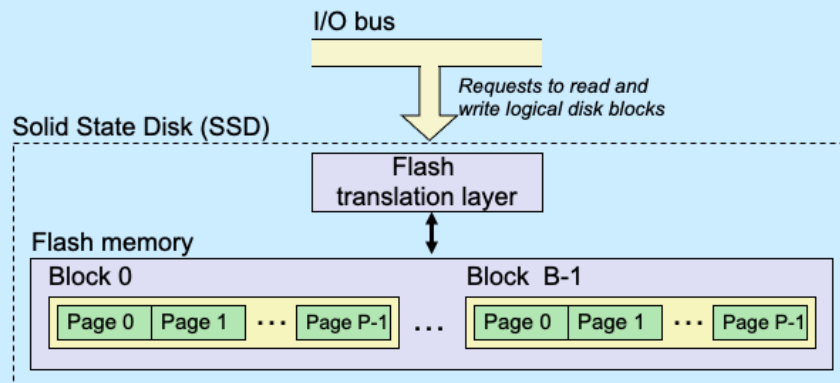
Supplied by CMU.

Reading a Disk Sector (3)



Supplied by CMU.

Solid-State Disks (SSDs)



- **Pages:** 512KB to 4KB; **blocks:** 32 to 128 pages
- **Data read/written** in units of pages
- **Page** can be written only after its block has been erased
- **A block** wears out after 100,000 repeated writes

Supplied by CMU.

SSD Performance Characteristics

Sequential read tput	250 MB/s	Sequential write tput	170 MB/s
Random read tput	140 MB/s	Random write tput	14 MB/s
Random read access	30 us	Random write access	300 us

- **Why are random writes so slow?**
 - erasing a block is slow (around 1 ms)
 - modifying a page triggers a copy of all useful pages in the block
 - » find a used block (new block) and erase it
 - » write the page into the new block
 - » copy other pages from old block to the new block

Supplied by CMU.


SSD Tradeoffs vs Rotating Disks

- **Advantages**
 - no moving parts → faster, less power, more rugged
- **Disadvantages**
 - have the potential to wear out
 - » mitigated by “wear-leveling logic” in flash translation layer
 - » e.g. Intel X25 guarantees 1 petabyte (10^{15} bytes) of random writes before they wear out
 - in 2010, about 100 times more expensive per byte
 - in 2017, about 6 times more expensive per byte
- **Applications**
 - smart phones, laptops
 - Apple “Fusion” drives

Supplied by CMU.

Reading a File on a Rotating Disk

- **Suppose the data of a file are stored on consecutive disk sectors on one track**
 - **this is the best possible scenario for reading data quickly**
 - » **single seek required**
 - » **single rotational delay**
 - » **all sectors read in a single scan**




Quiz 1

We have two files on the same (rotating) disk. The first file's data resides in consecutive sectors on one track, the second in consecutive sectors on another track. It takes a total of t seconds to read all of the first file then all of the second file.

Now suppose the files are read concurrently, perhaps a sector of the first, then a sector of the second, then the first, then the second, etc. Compared to reading them sequentially, this will take

- a) a lot less time
- b) around the same amount of time
- c) a lot more time



Quiz 2

We have two files on the same solid-state disk. Each file's data resides in consecutive blocks. It takes a total of t seconds to read all of the first file then all of the second file.

Now suppose the files are read concurrently, perhaps a block of the first, then a block of the second, then the first, then the second, etc. Compared to reading them sequentially, this will take

- a) a lot less time
- b) around the same amount of time
- c) a lot more time

Storage Trends

SRAM

Metric	1985	1990	1995	2000	2005	2010	2015	2015:1985
\$/MB	2,900	320	256	100	75	60	25	116
access (ns)	150	35	15	3	2	1.5	1.3	115

DRAM

Metric	1985	1990	1995	2000	2005	2010	2015	2015:1985
\$/MB	880	100	30	1	0.1	0.06	0.02	44,000
access (ns)	200	100	70	60	50	40	20	10
typical size (MB)	0.256	4	16	64	2,000	8,000	16,000	62,500

Disk

Metric	1985	1990	1995	2000	2005	2010	2015	2015:1985
\$/GB	100,000	8,000	300	10	5	.3	0.03	3,333,333
access (ms)	75	28	10	8	5	3	3	25
typical size (GB)	.01	.16	1	20	160	1,500	3,000	300,000

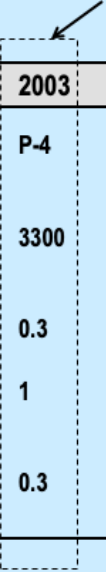
Supplied by CMU.

Current (2020) prices for SRAM vary a fair amount. As of 11/22, it can be had for around \$3/MB.

Current DRAM prices are around \$.00075/MB.

CPU Clock Rates

Inflection point in computer history
when designers hit the "Power Wall"




	1985	1990	1995	2000	2003	2005	2015	2015:1985
CPU	286	386	Pentium	P-III	P-4	Core 2	Core i7	---
Clock rate (MHz)	6	20	150	600	3300	2000	3000	500
Cycle time (ns)	166	50	6	1.6	0.3	0.50	0.33	500
Cores	1	1	1	1	1	2	4	4
Effective cycle time (ns)	166	50	6	1.6	0.3	0.25	0.08	2075

Supplied by CMU.

The CPU-Memory Gap

The gap widens between DRAM, disk, and CPU speeds

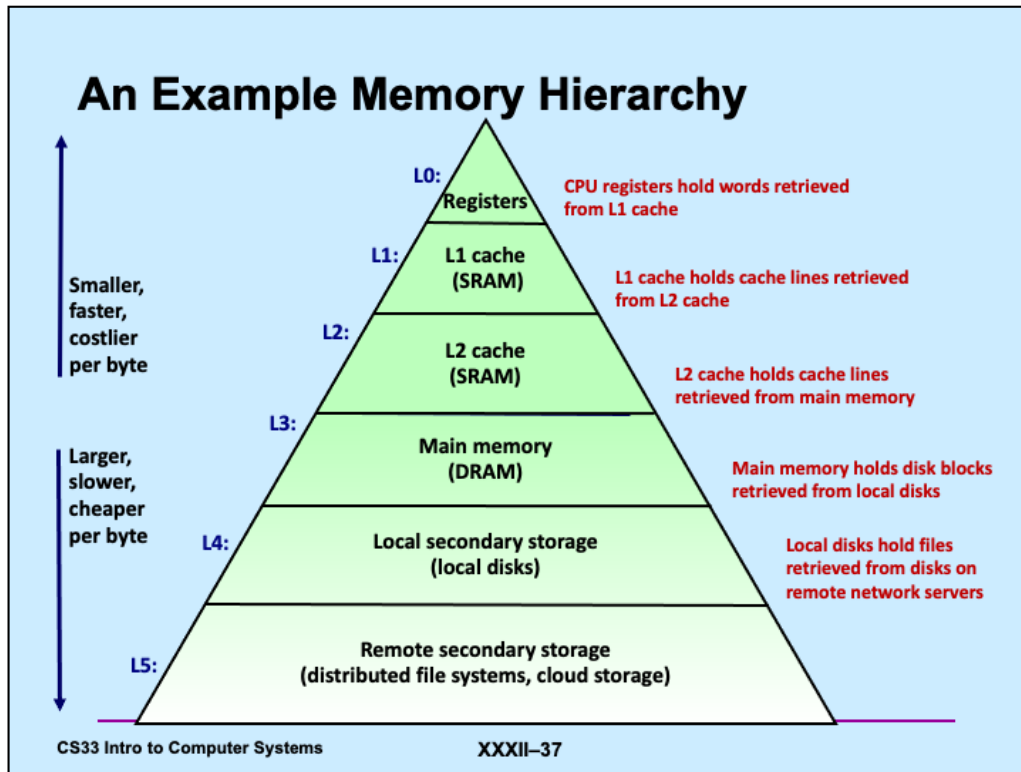


Supplied by CMU.

Memory Hierarchies

- **Some fundamental and enduring properties of hardware and software:**
 - fast storage technologies cost more per byte, have less capacity, and require more power (heat!)
 - the gap between CPU and main memory speed is widening
 - well written programs tend to exhibit good locality
- **These fundamental properties complement each other beautifully**
- **They suggest an approach for organizing memory and storage systems known as a **memory hierarchy****

Supplied by CMU.



Supplied by CMU.

Putting Things Into Perspective ...

- **Reading from:**
 - ... the L1 cache is like grabbing a piece of paper from your desk (3 seconds)
 - ... the L2 cache is picking up a book from a nearby shelf (14 seconds)
 - ... main system memory is taking a 4-minute walk down the hall to talk to a friend
 - ... a hard drive is like leaving the building to roam the earth for one year and three months

This analogy is from <http://duartes.org/gustavo/blog/post/what-your-computer-does-while-you-wait> (definitely worth reading!).

Disks Are Important

- **Cheap**
 - cost/byte much less than SSDs
- **(fairly) Reliable**
 - data written to a disk is likely to be there next year
- **Sometimes fast**
 - data in consecutive sectors on a track can be read quickly
- **Sometimes slow**
 - data in randomly scattered sectors takes a long time to read

Abstraction to the Rescue

- Programs don't deal with sectors, tracks, and cylinders
- Programs deal with *files*
 - maze.c rather than an ordered collection of sectors
 - OS provides the implementation

Implementation Problems

- **Speed**
 - **use the hierarchy**
 - » **copy files into RAM, copy back when done**
 - **optimize layout**
 - » **put sectors of a file in consecutive locations**
 - **use parallelism**
 - » **spread file over multiple disks**
 - » **read multiple sectors at once**

Implementation Problems

- **Reliability**
 - **computer crashes**
 - » what you thought was safely written to the file never made it to the disk — it's still in RAM, which is lost
 - » worse yet, some parts made it back to disk, some didn't
 - you don't know which is which
 - on-disk data structures might be totally trashed
 - **disk crashes**
 - » you had backed it up ... yesterday
 - **you screw up**
 - » you accidentally delete the entire directory containing your malloc solution

Implementation Problems

- **Reliability solutions**
 - **computer crashes**
 - » transaction-oriented file systems
 - » on-disk data structures always in well defined states
 - **disk crashes**
 - » files stored redundantly on multiple disks
 - **you screw up**
 - » file system automatically keeps “snapshots” of previous versions of files

You'll Soon Finish CS 33 ...

- You might

- celebrate



- take another systems course

- » 32

- » 138

- » 166

- » 167



- become a 33 TA



Systems Courses Next Semester

- **CS 32 (Intro to Software Engineering)**
 - you've mastered low-level systems programming
 - now do things at a higher level
 - learn software-engineering techniques using Java, XML, etc.
- **CS 138 (Distributed Systems)**
 - you now know how things work on one computer
 - what if you've got lots of computers?
 - some may have crashed, others may have been taken over by your worst (and smartest) enemy
- **CS 166 (Computer Systems Security)**
 - liked buffer?
 - you'll really like 166
- **CS 167/169 (Operating Systems)**
 - still mystified about what the OS does?
 - write your own!

The End

Well, not quite ...

Database is due on 12/10

No office hours Wednesday through Sunday

A few of us will be monitoring Piazza

Office hours resume next week

Happy coding and happy thanksgiving!